

ORACLE  
OPEN  
WORLD

#MySQL #oow16

# MySQL Performance Tuning 101

Jesper Wisborg Krogh

ORACLE®

Copyright © 2016, Oracle and/or its affiliates. All rights reserved.



## Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

# MySQL Performance Tuning 101

1. MySQL Performance Tuning Best Practices
2. Think
3. Monitoring
4. Create the Initial MySQL Configuration File
5. Buffer and Caches
6. Data consistency versus performance
7. Query Tuning
8. Consider the whole stack
9. Summary



# MySQL Performance Tuning Best Practices

ORACLE®

Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |

# MySQL Performance Tuning Best Practices

- Be wary about “best practices”
  - No two systems are the same
  - What was true earlier may no longer hold

# MySQL Performance Tuning Best Practices

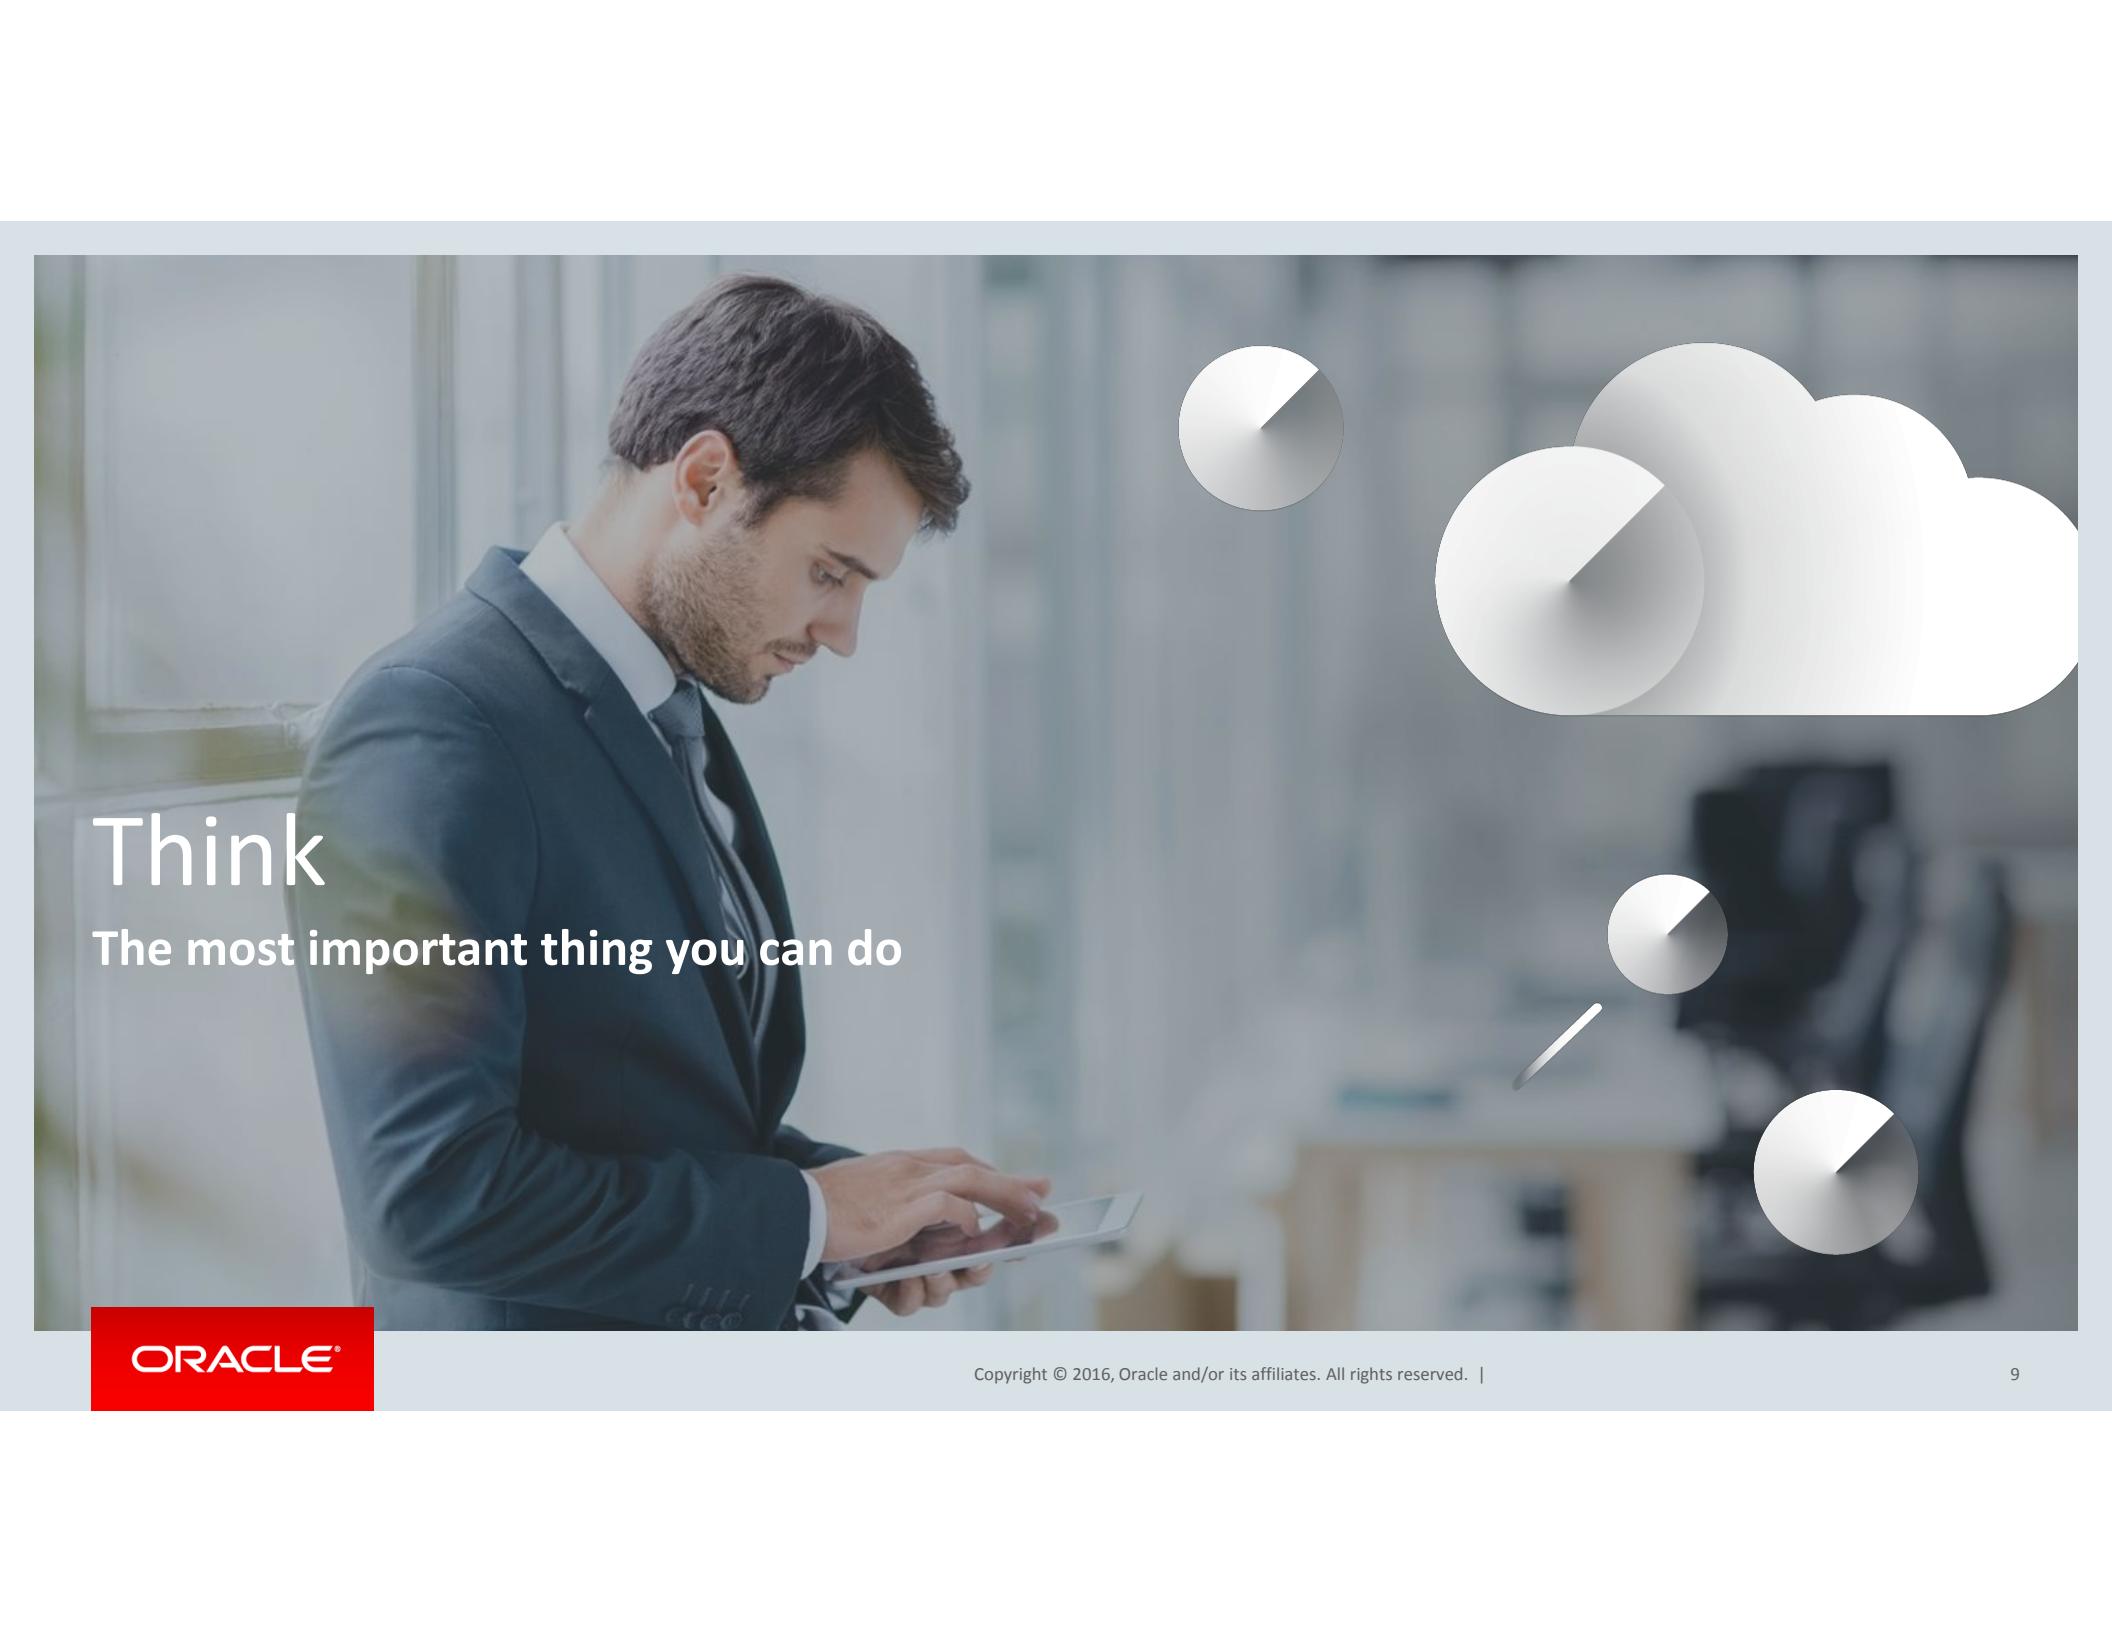
- Be wary about “best practices”
  - No two systems are the same
  - What was true earlier may no longer hold
- That said some guidelines can be given

# MySQL Performance Tuning Best Practices

- Think – consider what you are doing!
- Monitor your system
- Ensure you test your changes before deploying to production
  - The testing must reflect your production usage
- Make incremental changes
  - One change at a time
  - Relatively small changes

# MySQL Performance Tuning Best Practices

- Be mindful of your requirements
  - Some options give the choice between performance or data safety – what do you need?
- Often the default value is the best value
- Ensure all tables have a PRIMARY KEY
- InnoDB organizes the data according to the PRIMARY KEY:
  - The PRIMARY KEY is included in all secondary indexes in order to be able to locate the actual row => smaller PRIMARY KEY gives smaller secondary indexes.
  - A mostly sequential PRIMARY KEY is in general recommended to avoid inserting rows between existing rows.



**Think**  
The most important thing you can do

ORACLE®

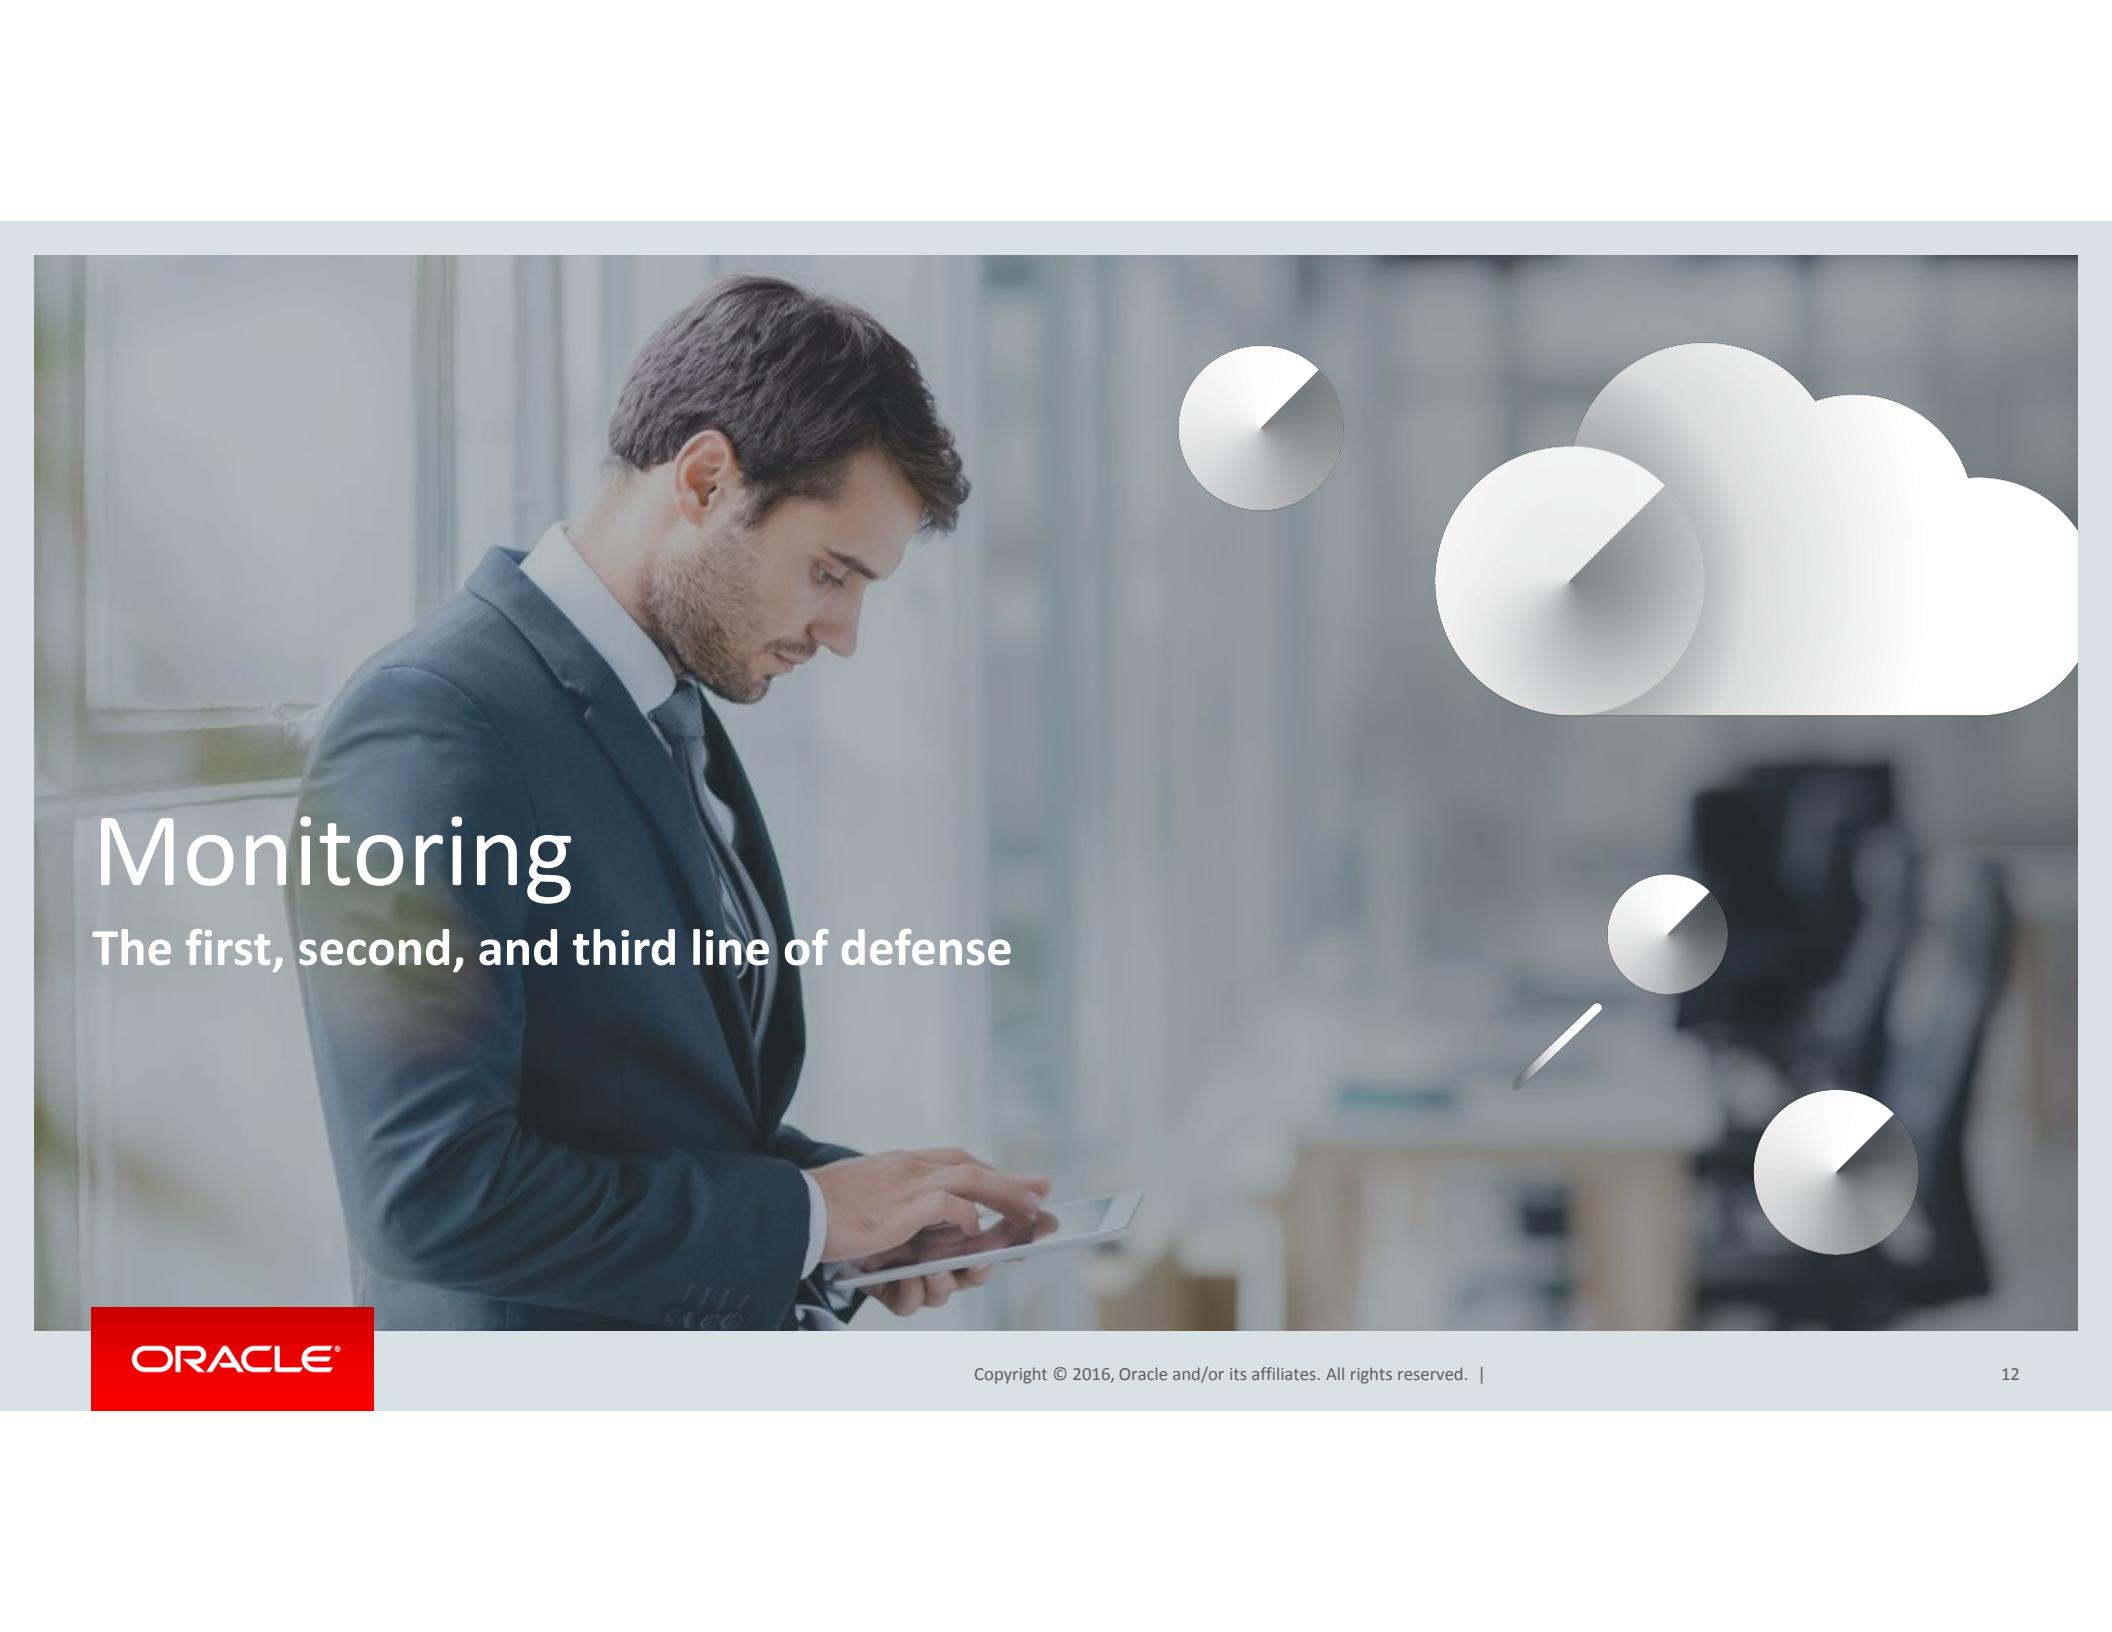
Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |

# Think

- Thinking is your best defense
- Analyze your performance problem:
  - Make sure you understand what the issue is:
    - No: performance is too slow
    - Yes:
      - The query takes 10 seconds, but it is used for interactive use so must complete in 0.1 second or less
      - The server must handle 200k queries per second
  - Determine the cause
    - Don't jump to conclusions
    - Consider the whole stack
    - Justify why you think you have found the cause

# Think

- Implement the solution
  - List all the solutions you can think of
    - Make sure you think outside the box
  - Explain why the solutions will work
  - Implement an action plan:
    - Test the action plan and update if necessary
    - Ensures you implement the solution the same on the test and production systems
    - If a regression occurs, you have a record of the steps you made



# Monitoring

## The first, second, and third line of defense

ORACLE®

Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |

12

# Monitoring

- Gives you a base line
- Tells you what has happened – useful to investigate performance issues
- Allows you to proactively handle potential issues

# Monitoring MySQL Enterprise Monitor

ORACLE® MySQL Enterprise Monitor

Dashboards ▾ Events Query Analyzer Reports & Graphs ▾ Refresh: Off ▾ admin ? ▾

All group - Graphs for last 1 hour (AEST) Edit

Database Statistics

Database Availability

Day: 100%  
Week: 99.871%  
Month: 99.871%

Graphs

Connections - All MySQL Instances

Statements... Database Activity - All MySQL Instances

Copyright © 2005, 2016, Oracle and/or its affiliates. All rights reserved.

Current Problem MySQL Instances

ID	Status	Emergency	Critical	Warning
No data available in table				

Showing 0 to 0 of 0 entries Show / hide columns

Current Problem Hosts

ID	Status	Emergency	Critical	Warning
ol7	Up	0	1	0

Showing 1 to 1 of 1 entries Show / hide columns

Current Emergency & Critical Events

Show 5 entries

Subject	Topic	Time	Actions
ol7, /dev/mapper/ol...	Filesystem /dev/m...	less than a minute ago	X

Showing 1 to 1 of 1 entries First Previous 1 Next Last

3.3.0.1097 - ol7 (10.0.2.15) - Sep 13, 2016 2:48:41 pm AEST (Up Since: 2 hours, 39 minutes ago) - About

The screenshot displays the MySQL Enterprise Monitor web interface. At the top, there's a navigation bar with links for Dashboards, Events, Query Analyzer, and Reports & Graphs. On the right side of the header, there are icons for switching between multiple windows, logging in as 'admin', and help. Below the header, a title bar says 'All group - Graphs for last 1 hour (AEST) Edit'. The main area is divided into several sections: 'Database Statistics' (with a 'Database Availability' chart), 'Graphs' (containing two stacked area charts for 'Connections - All MySQL Instances' and 'Database Activity - All MySQL Instances'), and three 'Current Problem' tables for MySQL Instances, Hosts, and Events. The 'Events' table shows one emergency event from 'ol7' about a file system issue. At the bottom, there are copyright notices for Oracle and a footer with a link to 'About'.

# Monitoring

- Many options available:
  - MySQL Enterprise Monitor
  - MySQL Plugin for Oracle Enterprise Monitor
  - Cacti
  - Zabbix
  - Nagios
  - And many more
- Make sure you configure alerts so you react appropriately to all events based on the severity level!

# Monitoring

## “Real time” monitoring

- On Linux perf is a great tool for live monitoring and can record over a time period
  - [https://perf.wiki.kernel.org/index.php/Main\\_Page](https://perf.wiki.kernel.org/index.php/Main_Page)
  - <http://www.brendangregg.com/perf.html>
- MySQL Enterprise Monitor has some reports to get snapshot data
- Workbench has performance reports based on the sys schema

# Monitoring MySQL Workbench Performance Reports

The screenshot shows the MySQL Workbench Performance Reports interface for a local instance (3306). The left sidebar lists various reports under categories like Hot Spots for I/O, High Cost SQL Statements, and Database Schema Statistics. The 'Top I/O by Event Category' report is selected and displayed in the main pane. The title is 'Top I/O by Event Category' and the subtitle is 'Show the highest IO Data usage by event categories'. A table lists the top 10 IO Event Types with their corresponding statistics.

IO Event Type	Total IOs (#)	Total Time (s)	Min Time (us)	Avg Time (ms)	Max Time (s)
innodb/innodb_data_file	14467	23.07	0.00	1.59	2.35
sql/FRM	9006	13.29	0.00	1.48	0.72
innodb/innodb_log_file	21	3.64	0.81	173.57	1.84
mysys/cnf	38	1.81	0.09	47.71	1.81
myisam/kfile	114	1.47	0.02	12.92	0.45
myisam/dfile	2056	0.94	0.02	0.46	0.31
csv/metadata	59	0.39	0.02	6.55	0.23
sql/ERRMSG	5	0.36	29.71	72.94	0.29
sql/file_parser	36	0.28	0.00	7.80	0.15
sql/dbopt	80	0.03	0.02	0.33	0.02

# Monitoring

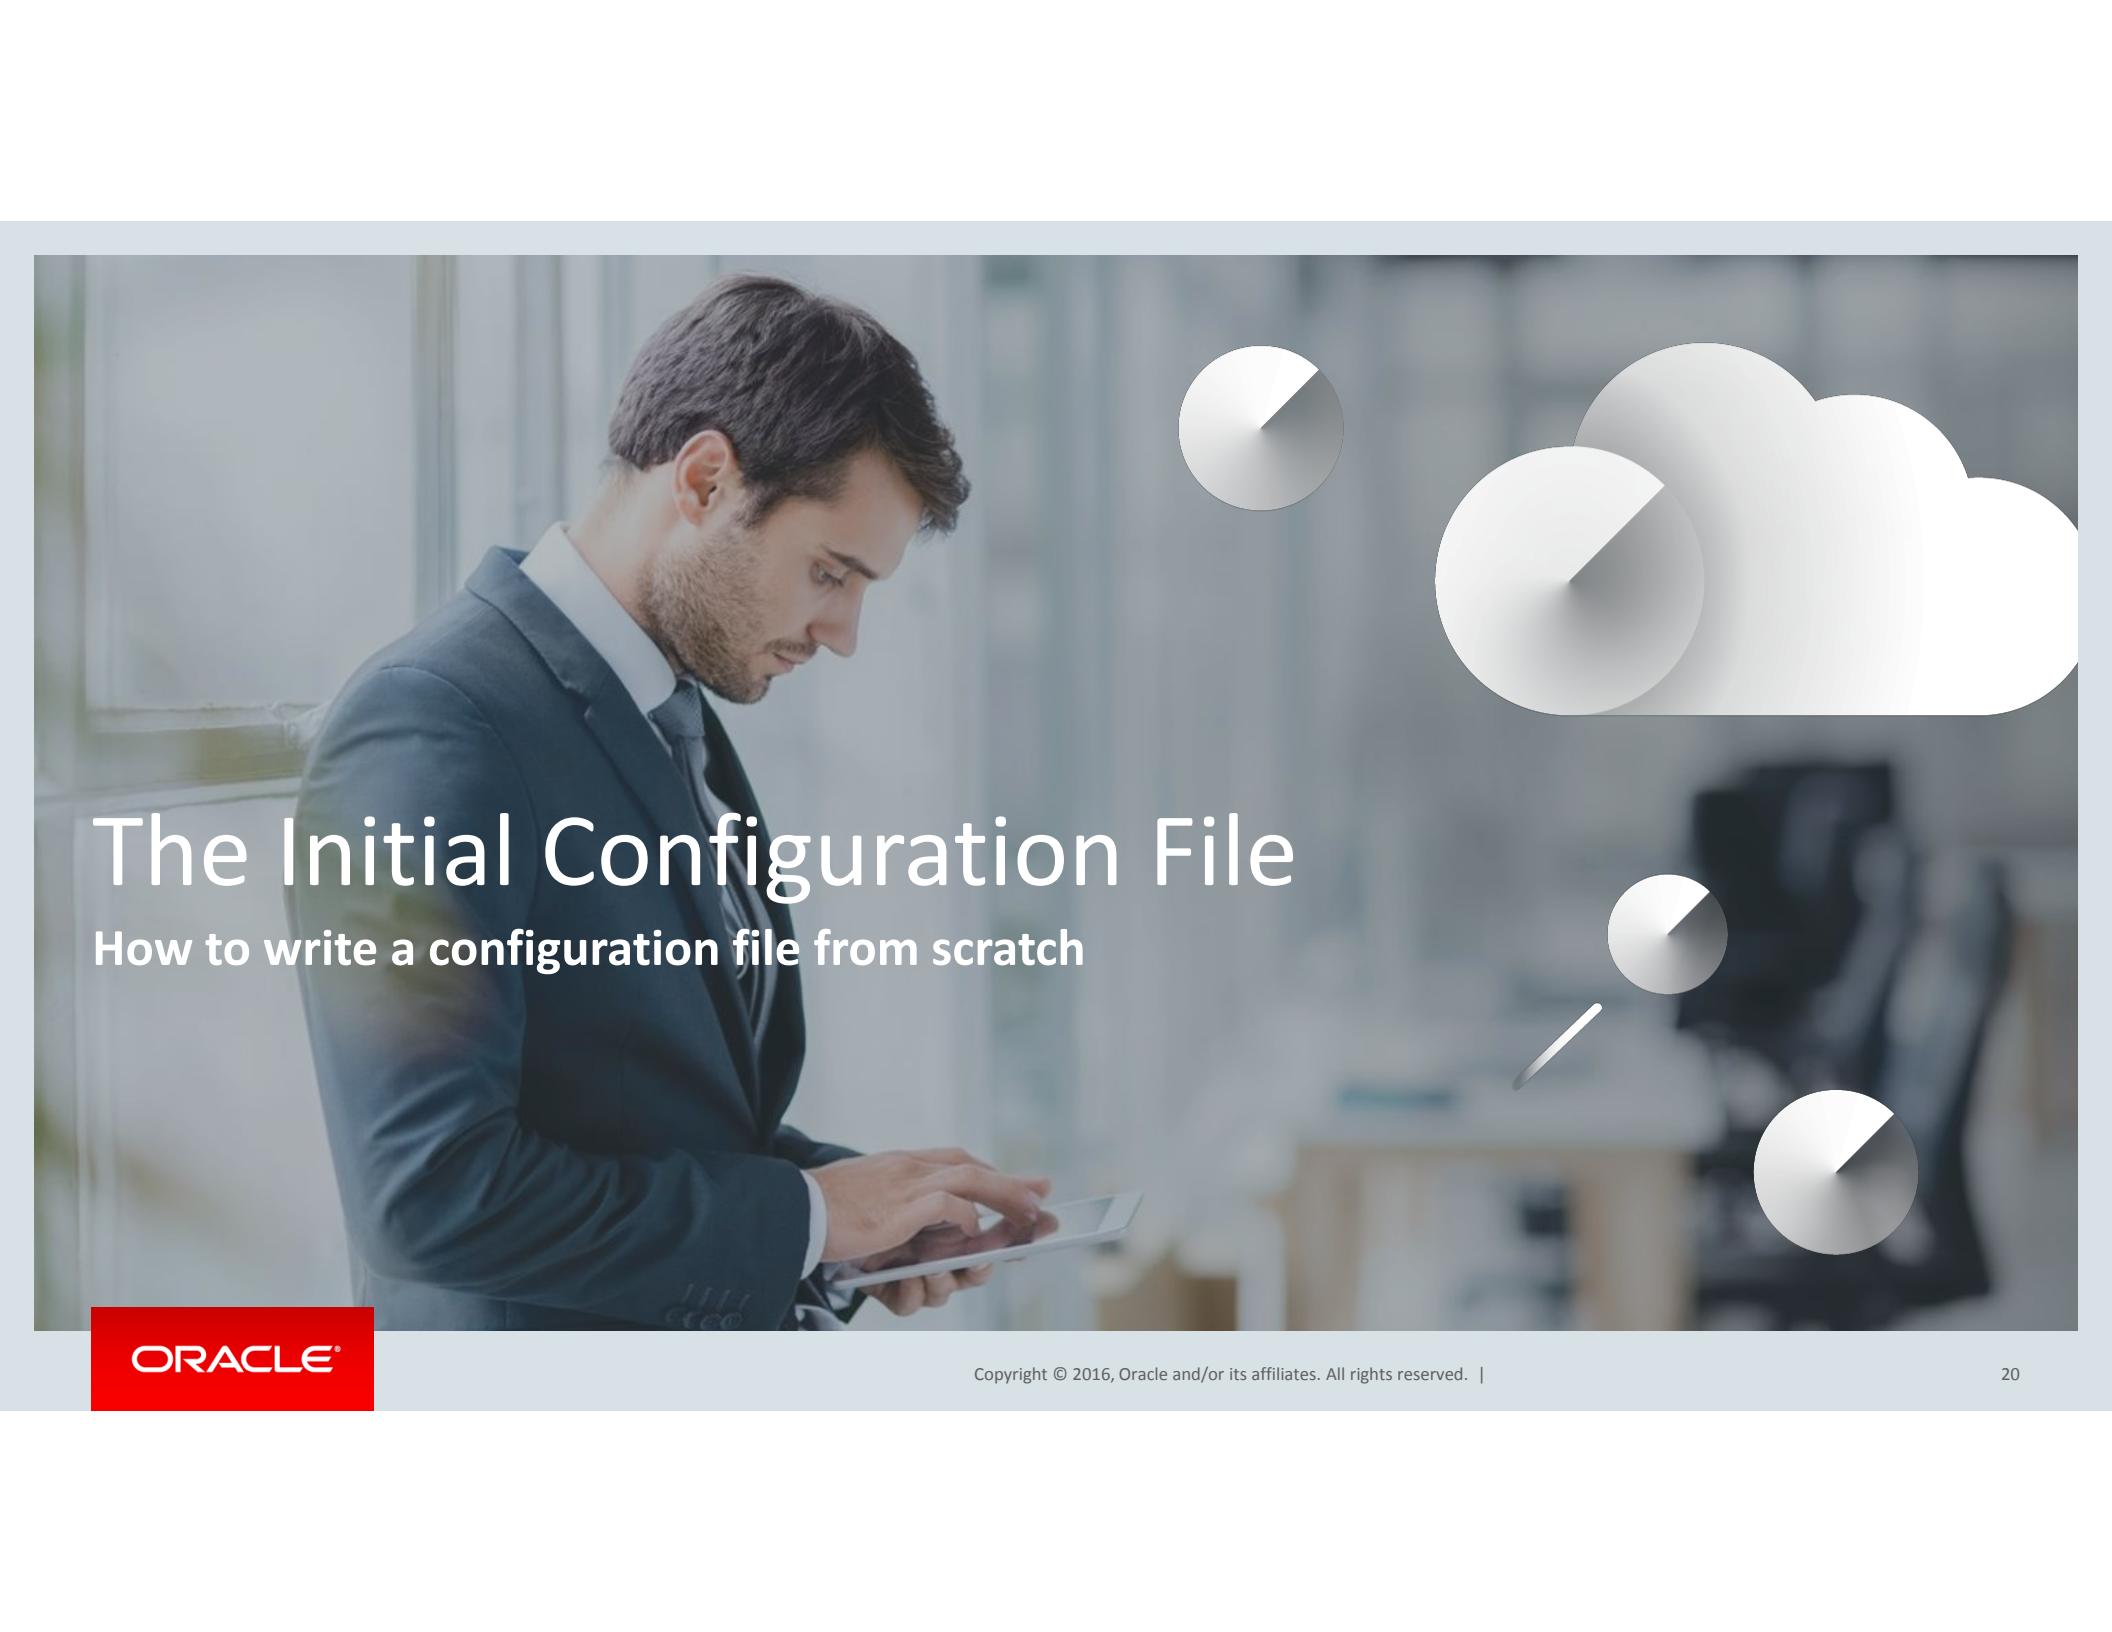
## What was that sys schema?

- The sys schema is a collection of views, functions, and procedures aiming to make the Information Schema and Performance Schema simpler to use.
- Included by default in MySQL 5.7+
- Also available from <https://github.com/mysql/mysql-sys> for MySQL 5.6
- Formerly known as ps\_helper by Mark Leith
- For a deep dive into the sys schema, see **Mark Leith's talk Thursday at 1:15pm in Park Central - City**
- <https://dev.mysql.com/doc/refman/5.7/en/sys-schema.html>

# Monitoring

## Monitoring has overhead

- While monitoring is good, there is such as thing as too much monitoring
- All monitoring has overhead – some more than others
- Particularly be careful with queries like:
  - SHOW PROCESSLIST
  - SHOW ENGINE INNODB STATUS
- If performed too often they can cause an outage on their own
- In 5.6+: Use performance\_schema.threads instead of SHOW PROCESSLIST

A professional man in a dark blue suit and tie is looking down at a white tablet device he is holding in his hands. He is positioned on the left side of the frame, facing right. The background is a blurred office environment. Floating around him are several white pie charts of different sizes, some with visible segments and others more solid. The overall theme suggests data analysis or configuration management.

# The Initial Configuration File

**How to write a configuration file from scratch**

# Initial Configuration File

1. Start with an empty configuration file
2. Set paths, port, etc.
3. Enable additional monitoring
4. Set capacity settings
5. Don't do much more!

# Initial Configuration File

**Start with empty configuration file**

- Remember the configuration templates in MySQL 5.5 and earlier?
  - my-huge.cnf/my-huge.ini
  - my-innodb-heavy-4G.cnf/my-innodb-heavy-4G.ini (20kB!)
  - my-large.cnf/my-large.ini
  - my-medium.cnf/my-medium.cnf
  - my-small.cnf/my-small.cnf
- In MySQL 5.6+:
  - my-default.cnf (1.1kB)

# Initial Configuration File

## my-default.cnf

```
[mysqld]

# Remove leading # and set to the amount of RAM for the most important data
# cache in MySQL. Start at 70% of total RAM for dedicated server, else 10%.
# innodb_buffer_pool_size = 128M

# Remove leading # to turn on a very important data integrity option: logging
# changes to the binary log between backups.
# log_bin

# These are commonly set, remove the # and set as required.
# basedir = .....
# datadir = .....
# port = .....
# server_id = .....
# socket = .....

# Remove leading # to set options mainly useful for reporting servers.
# The server defaults are faster for transactions and fast SELECTs.
# Adjust sizes as needed, experiment to find the optimal values.
# join_buffer_size = 128M
# sort_buffer_size = 2M
# read_rnd_buffer_size = 2M

sql_mode=NO_ENGINE_SUBSTITUTION,STRICT_TRANS_TABLES
```

# Initial Configuration File

**Start with empty configuration file**

- Lots of work has gone into improving the defaults in MySQL 5.6 and 5.7
- The work continues
- So it is no longer necessary with the huge templates

# Initial Configuration File

## Paths

- Some important paths:
  - datadir – where the data lives by default – default location for:
    - Data and indexes for non-InnoDB tables
    - InnoDB file-per-table, general, system, and undo log tablespaces
  - innodb\_data\_home\_dir – default location for the InnoDB system tablespace (ibdata1)
  - innodb\_log\_group\_home\_dir – path to the InnoDB redo log files
  - innodb\_undo\_directory – path to the InnoDB undo tablespaces
  - log\_bin – dual functionality: enable binary logging and set path/file name prefix
  - log\_error – makes sure you know where it is

# Initial Configuration File

## Paths

- It can be an advantage to have the paths pointing to separate disk systems
  - Reduces contention
  - I/O often becomes a bottleneck
  - Can place the hot files on faster disks
    - For example with `innodb_flush_log_at_trx_commit = 1` and high transaction commit rate, it can be necessary to place the InnoDB redo log on SSDs to support the high rate of flushes.
- Note: File-per-table tablespaces and general tablespaces can located outside datadir when they are created

# Initial Configuration File

## Enable additional monitoring

- If you use InnoDB, enable all the INNODB\_METRICS counters:
  - `innodb_monitor_enable = '%'`
    - The overhead has turned out to be small so worth the extra details
- Ensure the Performance Schema is enabled
  - Does have overhead, but provides very useful information for performance tuning
  - Consider enabling additional consumers and instruments as required (for example for transactions in MySQL 5.7)
  - The Performance Schema can also be configured dynamically at runtime

# Initial Configuration File

## Capacity settings

- Most important options for performance:
  - innodb\_buffer\_pool\_size for InnoDB workloads
  - key\_buffer\_size for MyISAM workloads
- Other important InnoDB options:
  - innodb\_buffer\_pool\_instances
  - innodb\_log\_file\_size / innodb\_log\_files\_in\_group
  - innodb\_thread\_concurrency
  - innodb\_undo\_tablespaces
  - innodb\_file\_per\_table (enabled by default in 5.6+)

# Initial Configuration File

## Capacity settings

- Other capacity options:
  - max\_connections
  - table\_definition\_cache
  - table\_open\_cache
  - table\_open\_cache\_instances

# Initial Configuration File

## Capacity settings - `innodb_buffer_pool_size`

- MySQL Reference Manual (about `innodb_buffer_pool_size`):  
**“On a dedicated database server, you might set the buffer pool size to 80% of the machine's physical memory size.”**
  - What if you have 1TB of memory? Do you still want to reserve 20% for other uses?
- Instead:
  - How much memory does the host have?
  - Subtract memory required by OS and other processes
  - Subtract memory required by MySQL other then the InnoDB buffer pool
  - Choose minimum of this and the size of the “working data set”

## innodb\_buffer\_pool\_size – Side Note

- In MySQL 5.7, innodb\_buffer\_pool\_size can be changed dynamically:

```
mysql> SELECT (POOL_SIZE*@@global.innodb_page_size/1024/1024) AS 'PoolSize (MB)'  
       FROM information_schema.INNODB_BUFFER_POOL_STATS\G  
*****  
***** 1. row *****  
PoolSize (MB): 127.98437500  
  
mysql> SELECT VARIABLE_VALUE  
       FROM performance_schema.global_status  
      WHERE VARIABLE_NAME = 'Innodb_buffer_pool_resize_status'\G  
*****  
***** 1. row *****  
VARIABLE_VALUE: Completed resizing buffer pool at 160912 18:49:50.  
  
mysql> SELECT (POOL_SIZE*@@global.innodb_page_size/1024/1024) AS 'PoolSize (MB)'  
       FROM information_schema.INNODB_BUFFER_POOL_STATS\G  
*****  
***** 1. row *****  
PoolSize (MB): 4095.98437500
```

# Initial Configuration File

## Capacity settings - `innodb_buffer_pool_instances`

- Specifies how many instances to split the buffer pool into - can reduce contention for concurrent workload
- Rule of thumb:
  - `innodb_buffer_pool_size` <= 1G: 1 instance
  - `innodb_buffer_pool_size` <= 16G: 8 instances
  - `innodb_buffer_pool_size` > 16G: 2G per instance
  - Max allowed value: 64
- But test with your workload!
- Requires restart

# Initial Configuration File

## Capacity settings – InnoDB redo log

- Total redo log size defined by two options:
  - innodb\_log\_file\_size
  - innodb\_log\_files\_in\_group
- Total size = innodb\_log\_file\_size \* innodb\_log\_files\_in\_group
  - Max supported total redo log size:
    - MySQL 5.5 and earlier: just below 4G
    - MySQL 5.6 and later: just below 512G
- Should be large enough to avoid “excessive” checkpointing

# Initial Configuration File

## Capacity settings – InnoDB redo log – Is it large enough?

- Get current Log sequence number (LSN) and last checkpoint:
  - SHOW ENGINE INNOD STATUS:

```
---  
LOG  
---  
Log sequence number 602763740  
Log flushed up to 602763740  
Pages flushed up to 584668961  
Last checkpoint at 555157885
```

- INNODB\_METRICS:

```
mysql> SELECT NAME, COUNT  
       FROM information_schema.INNODB_METRICS  
      WHERE NAME IN ('log_lsn_current', 'log_lsn_last_checkpoint');  
+-----+-----+  
| NAME           | COUNT |  
+-----+-----+  
| log_lsn_last_checkpoint | 555157885 |  
| log_lsn_current     | 602763740 |  
+-----+-----+
```

# Initial Configuration File

## Capacity settings – InnoDB redo log – Is it large enough?

- Get current Log sequence number (LSN) and last checkpoint:
  - Sys schema metrics view:

```
mysql> SELECT *
      FROM sys.metrics
     WHERE Variable_name IN ('log_lsn_current', 'log_lsn_last_checkpoint');
+-----+-----+-----+-----+
| Variable_name | Variable_value | Type          | Enabled |
+-----+-----+-----+-----+
| log_lsn_current | 602763740 | InnoDB Metrics - recovery | YES    |
| log_lsn_last_checkpoint | 555157885 | InnoDB Metrics - recovery | YES    |
+-----+-----+-----+-----+
```

# Initial Configuration File

## Capacity settings – InnoDB redo log – Is it large enough?

- Calculate the amount of used redo log:
  - Used log =  $\text{log\_lsn\_current} - \text{log\_lsn\_last\_checkpoint}$   
= 602763740 - 555157885  
= 47605855 (bytes)
- Compare to total size:
  - Used % =  $(\text{Used log} / \text{Total log}) * 100$   
=  $(47605855 / (\text{innodb\_log\_file\_size} * \text{innodb\_log\_files\_in\_group})) * 100$   
=  $(47605855 / 100663296) * 100$   
= 47.29 %

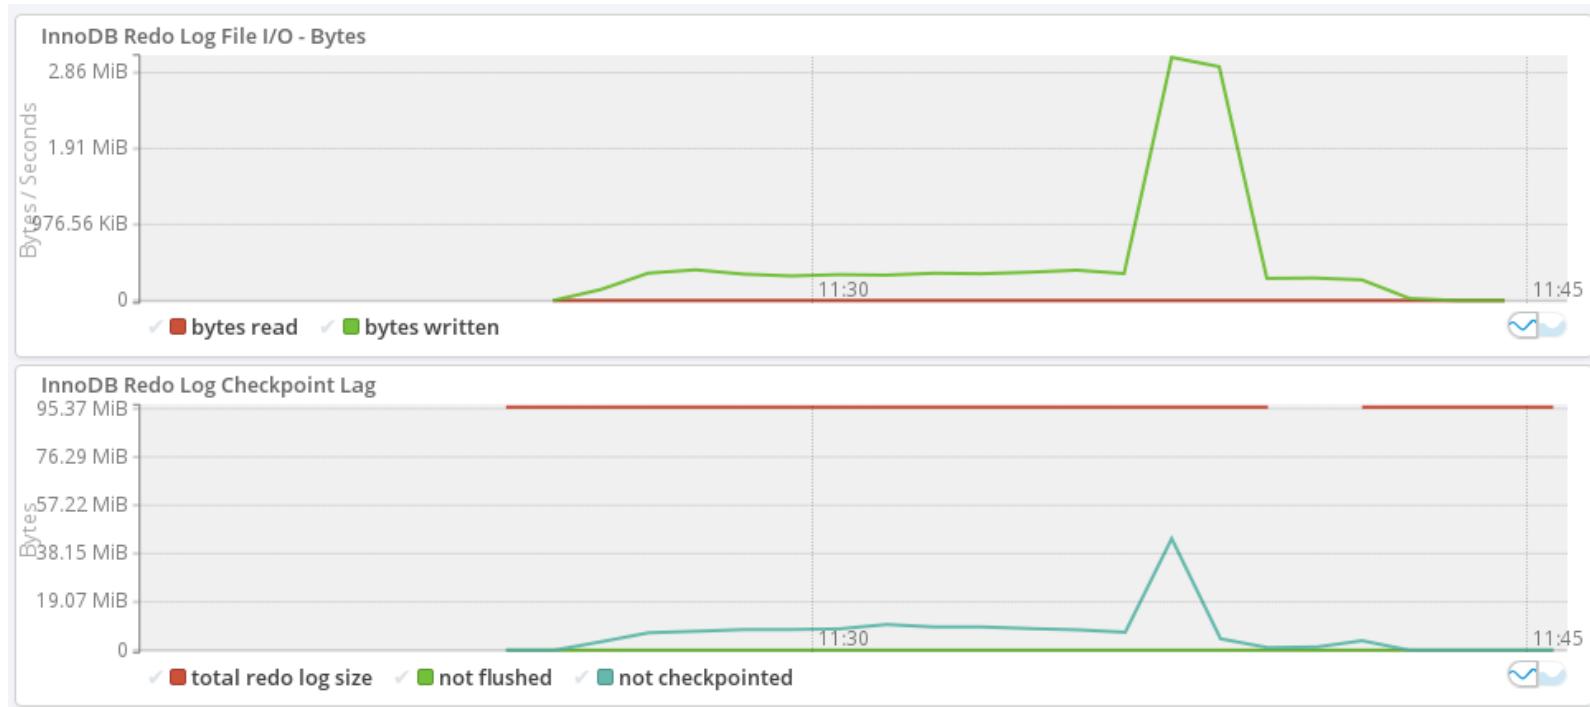
# Initial Configuration File

## Capacity settings – InnoDB redo log – Is it large enough?

- If the usage reaches 75% an asynchronous flush is triggered
  - I/O can be so intensive that all other work stalls
  - Main InnoDB thread may be in the state: **flushing buffer pool pages**
- Ensure you have enough head room to handle peak loads
  - For example aim at using at most 60% or 70% of the redo log
- Important that your monitoring solution monitors the redo log usage
- Improvements to the flushing algorithm and new options to control the I/O have been added in newer MySQL versions

# Initial Configuration File

## Capacity settings – InnoDB redo log – Is it large enough?



# Initial Configuration File

## Capacity settings - `innodb_thread_concurrency`

- Defines how many queries InnoDB will allow to execute concurrently
- 0 means unlimited and is often a good value if the number of vCPUs is the limiting factor
- Otherwise benchmark to find your optimal value – often good values are:
  - MySQL 5.5: Between 8 and 24
  - MySQL 5.6: Between 16 and 64
  - MySQL 5.7: Between 32 and 128
- Do not mix up with `thread_concurrency` (removed in 5.7!)

# Initial Configuration File

## InnoDB Undo Log

- `innodb_undo_tablespaces` can only be set before initializing the datadir!
- If set creates undo log tablespaces outside the system tablespace (`ibdata1`)
- Advantages:
  - Keeps the system tablespace smaller
  - Can keep the undo log on faster disks
  - In MySQL 5.7 the undo log tablespaces can be truncated
- Each undo tablespace file is 10M initially
- Maximum of 95 undo tablespaces in 5.7

# Initial Configuration File

## Other capacity options

- `max_connections` – be careful setting this too large as each connection requires memory
- `table_definition_cache` – ensure all tables can be in the cache. If you expect 4000 tables, set `table_definition_cache > 4000`
- `table_open_cache` – each table can be open more than once
- `table_open_cache_instances` - Typically a good value is 16

# Initial Configuration File

## Other capacity options

- max\_connections, table\_definition\_cache, and table\_open\_cache determines the default value for several Performance Schema options
  - Small instance:
    - max\_connections <= 151 && table\_definition\_cache <= 400 && table\_open\_cache <= 2000
  - Medium instance:
    - max\_connections <= 302 && table\_definition\_cache <= 800 && table\_open\_cache <= 4000
  - Otherwise large instance

# Initial Configuration File

## Other capacity options – MySQL 5.6 Performance Schema Memory Usage

- max\_connections = 151; table\_open\_cache = 200

table_definition_cache	Instance Size	Total Performance Schema Memory Usage
400	Small	52.6 MB
401	Medium	90.7 MB
800	Medium	98.4 MB
801	Large	400 MB

- MySQL 5.7 has more dynamic memory usage

# Oracle MySQL Cloud Service

**Simple, Automated, Integrated & Enterprise-Ready**

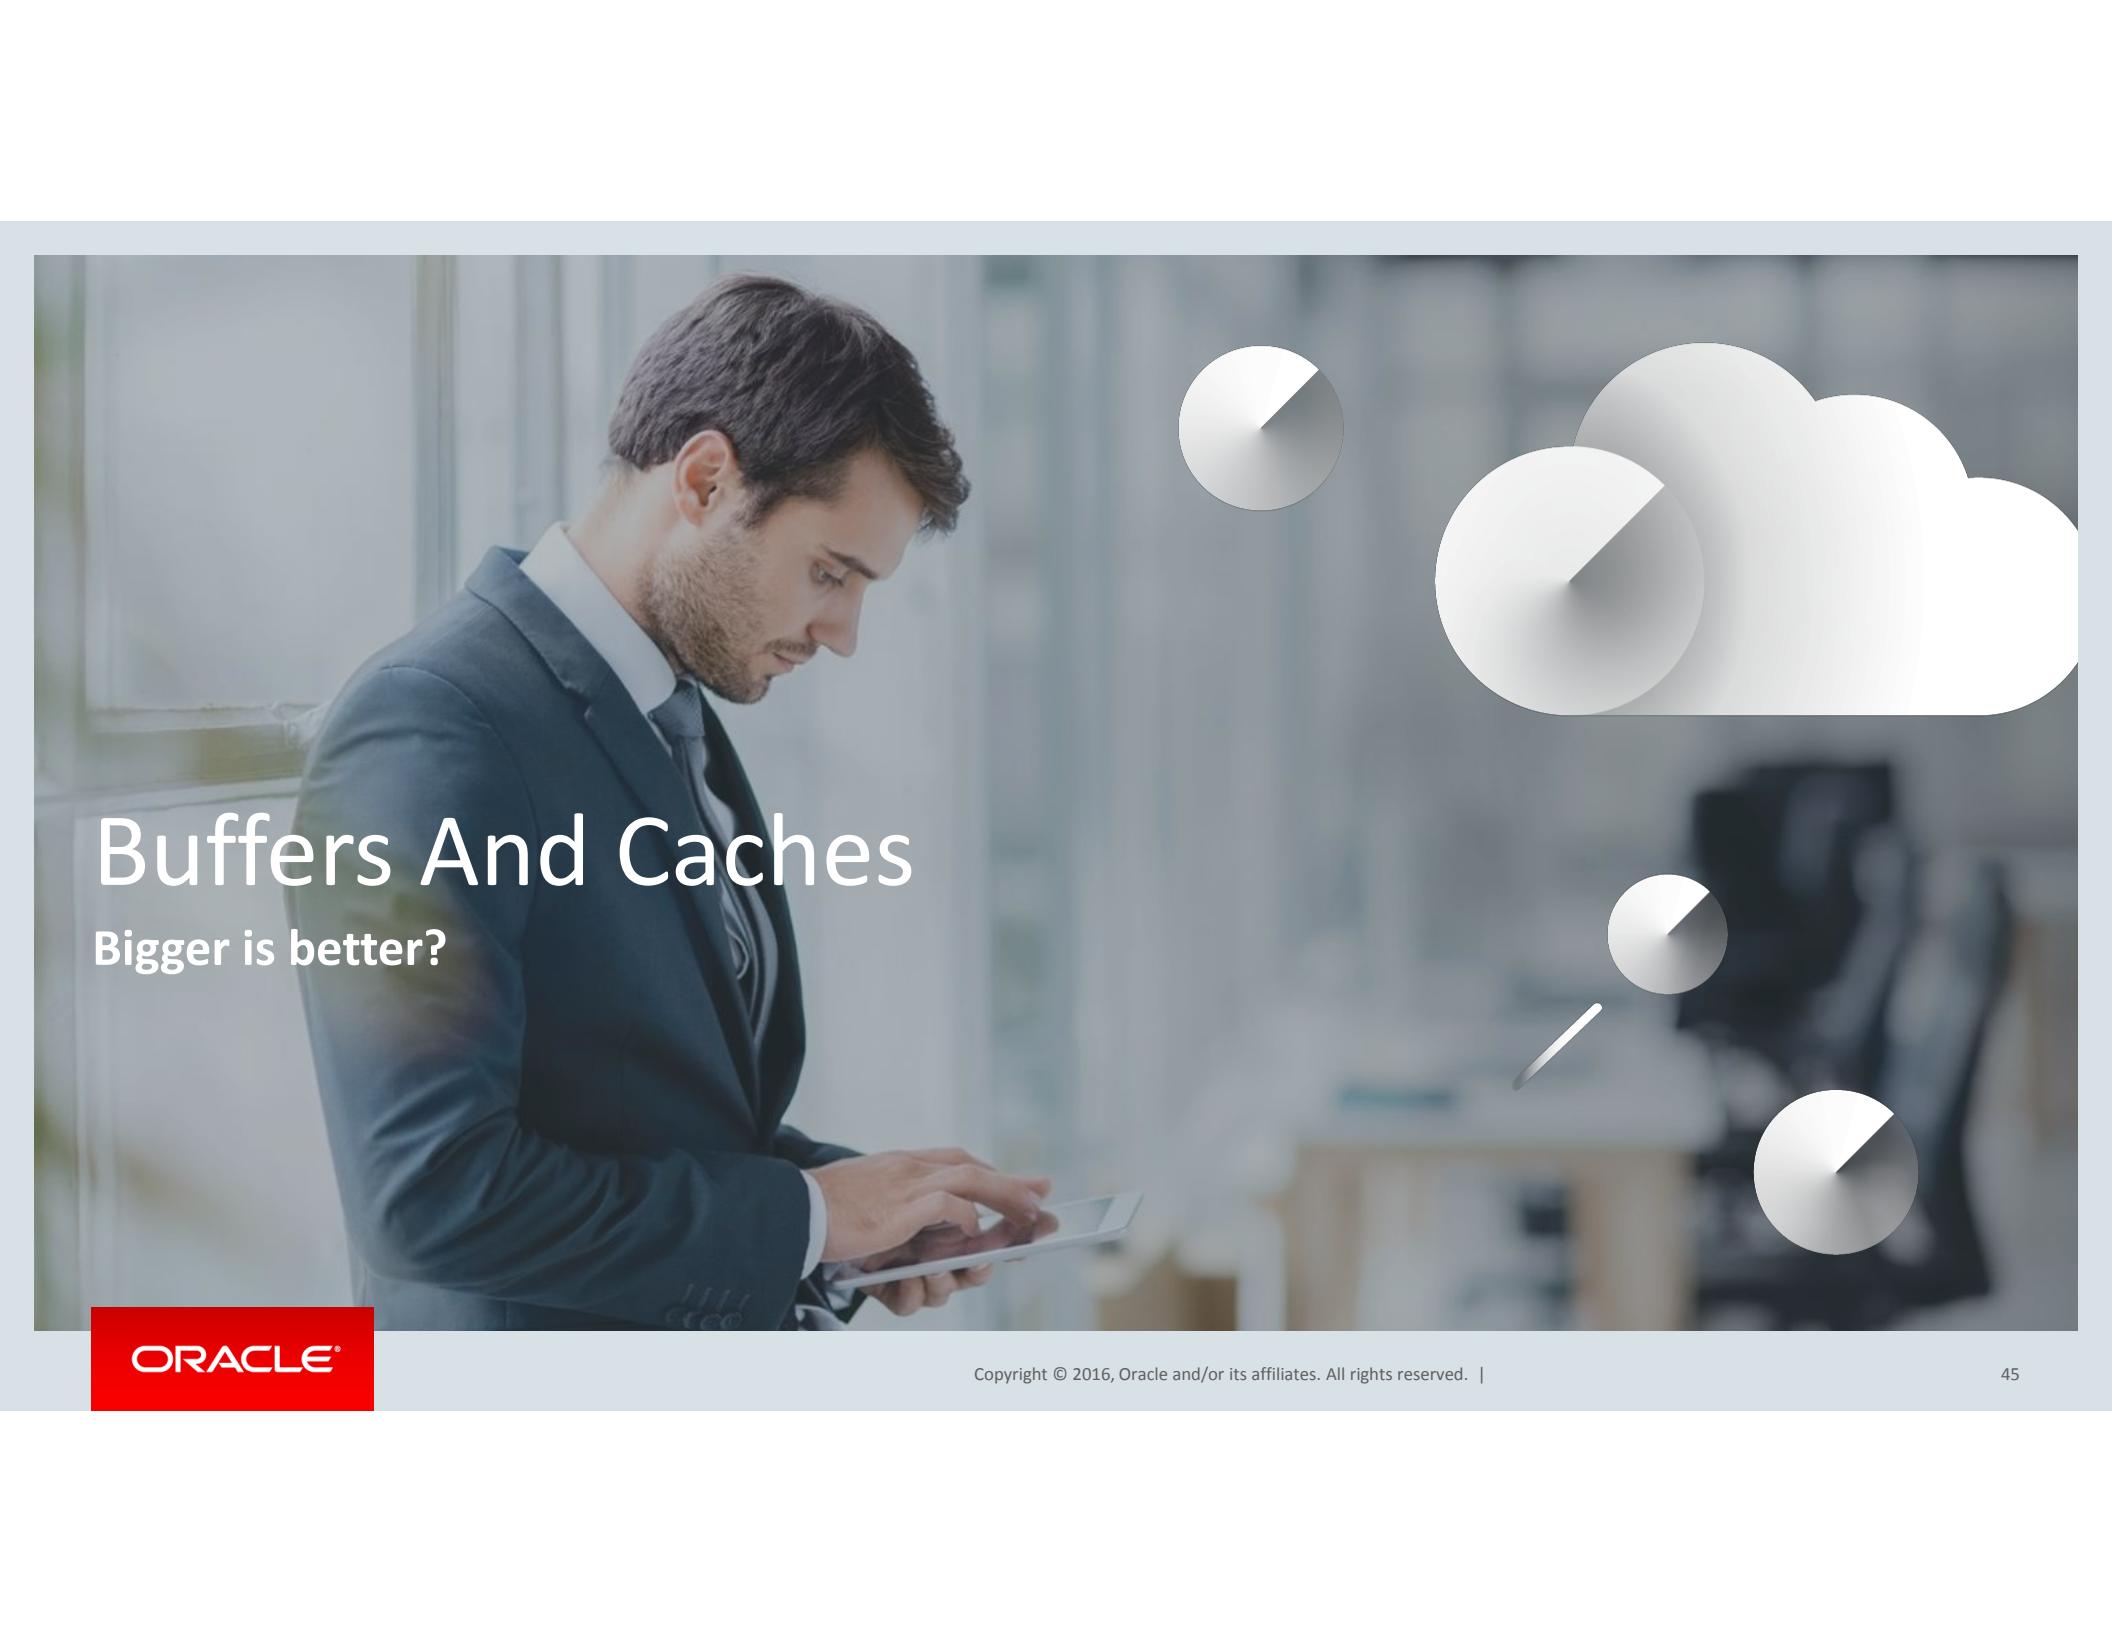
- The #1 Open Source Database in Oracle Cloud
- Only public cloud integrating MySQL Enterprise Edition
- Rapidly, securely and cost-effectively develop & deploy modern MySQL-based applications



[Cloud.oracle.com/mysql](http://Cloud.oracle.com/mysql): Learn more & Free Trial!

ORACLE®

Copyright © 2015, Oracle and/or its affiliates. All rights reserved. |



# Buffers And Caches

Bigger is better?



# Buffers And Caches

## Several buffers and caches available

- Query Cache - `query_cache_size`
- Per query buffers:
  - `join_buffer_size`
  - `read_buffer_size`
  - `read_rnd_buffer_size`
  - `sort_buffer_size`
- And more

# Buffers And Caches

**Several buffers and caches available**

- Surely enabling caches is a good thing?
- Surely the larger caches and buffers the better?

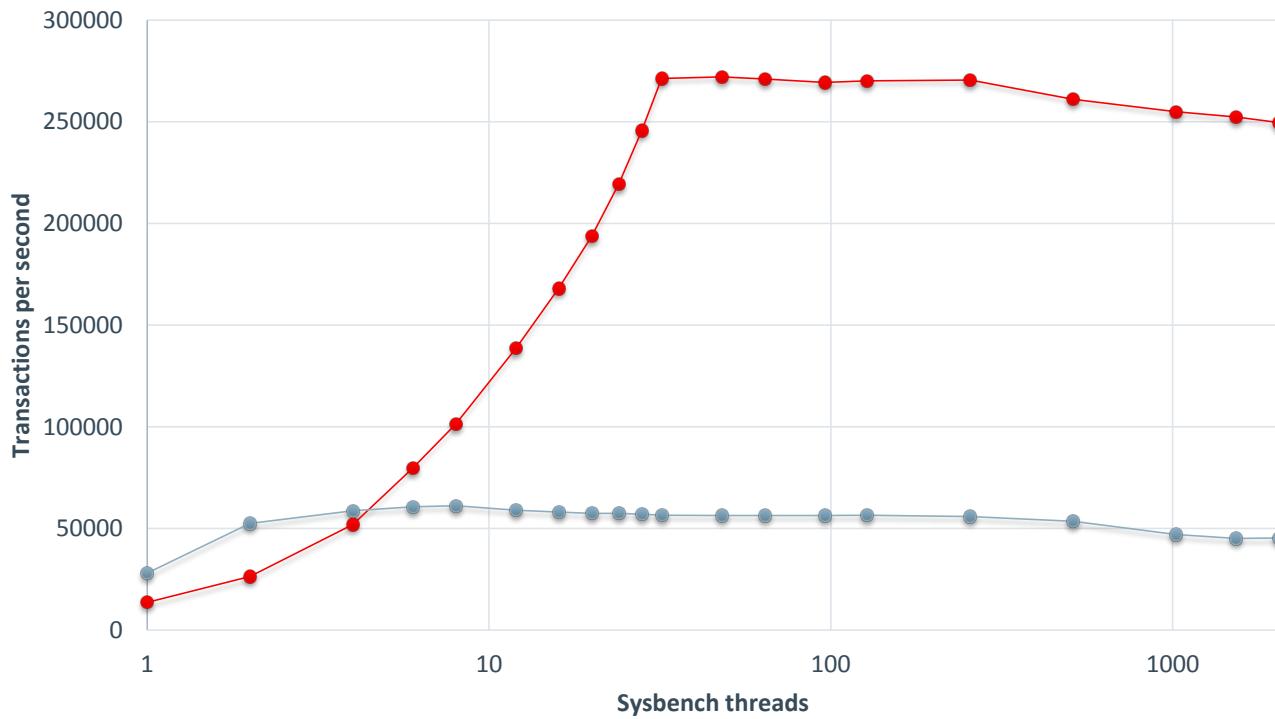
# Buffers And Caches

Several buffers and caches available

- Surely enabling caches is a good thing?
- Surely the larger caches and buffers the better?
- **No!**

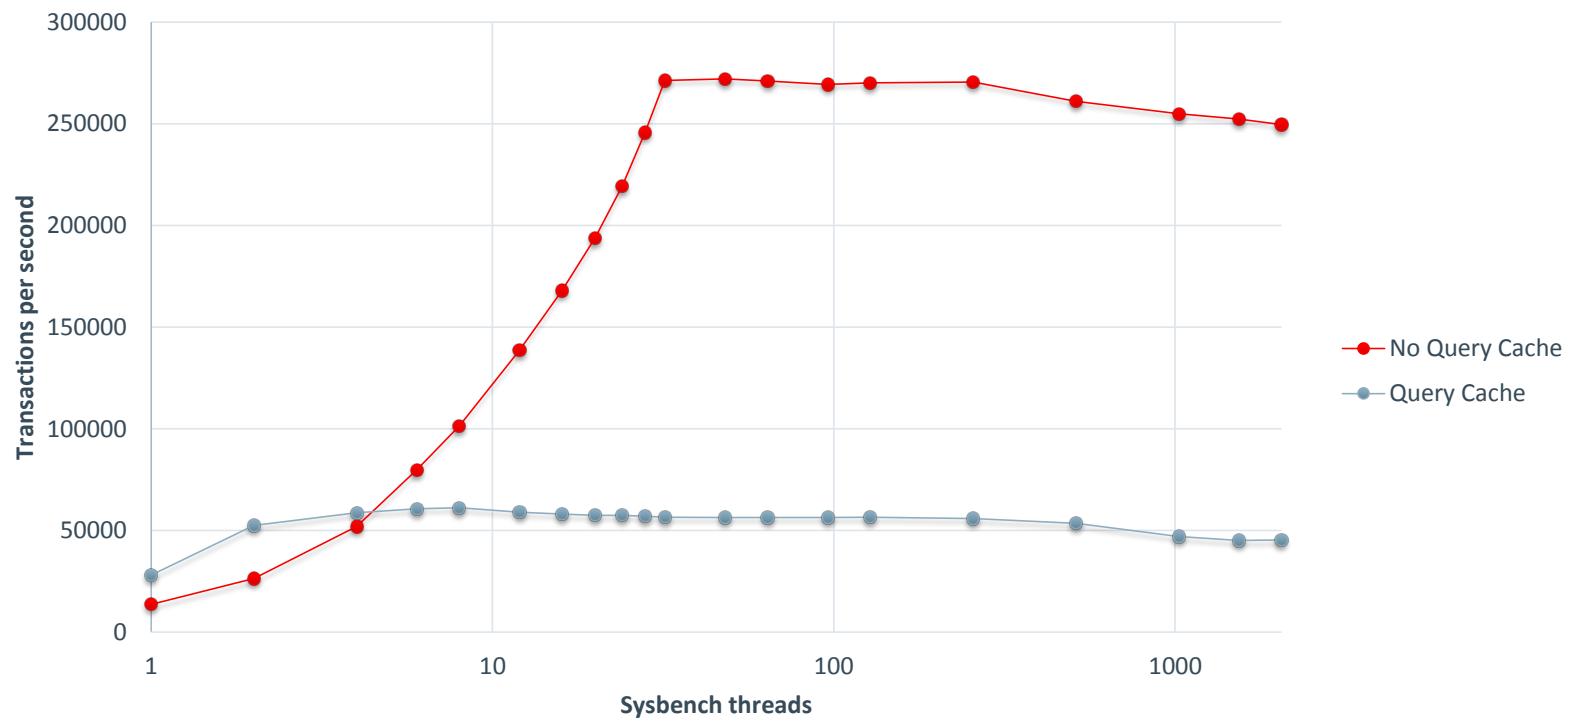
# Buffers And Caches

## Query Cache



# Buffers And Caches

## Query Cache



# Buffers And Caches

## Query Cache

- It is rarely good to enable the Query Cache
  - Guarded by a single mutex
- Most workloads is best left with the Query Cache disabled (the default):
  - `query_cache_type = 0`
- If you think your workload benefits from the Query Cache, test it
  - The more writes, the less benefit
  - The more data fitting into the buffer pool, the less benefit
  - The more complex queries and the larger scans, the more benefit
- Often other caching solutions are a better option.

# Buffers And Caches

## `join_buffer_size`

- In MySQL 5.5 and earlier: only used for plain index scans, range index scans, and joins that do not use indexes
  - No reason to have it larger than the size of each matching row
- In MySQL 5.6+ also used for Batched Key Access (BKA)
  - Queries using BKA can benefit from a larger join buffer
- Is the minimum size allocated!
- Usually a small global (default for new connections) value is best: 32k-256k
- Increase for a session as needed

# Buffers And Caches

## sort\_buffer\_size

- Like for join\_buffer\_size a small global value is usually best: 32k-256k
- Can monitor using the global status variable Sort\_merge\_passes:

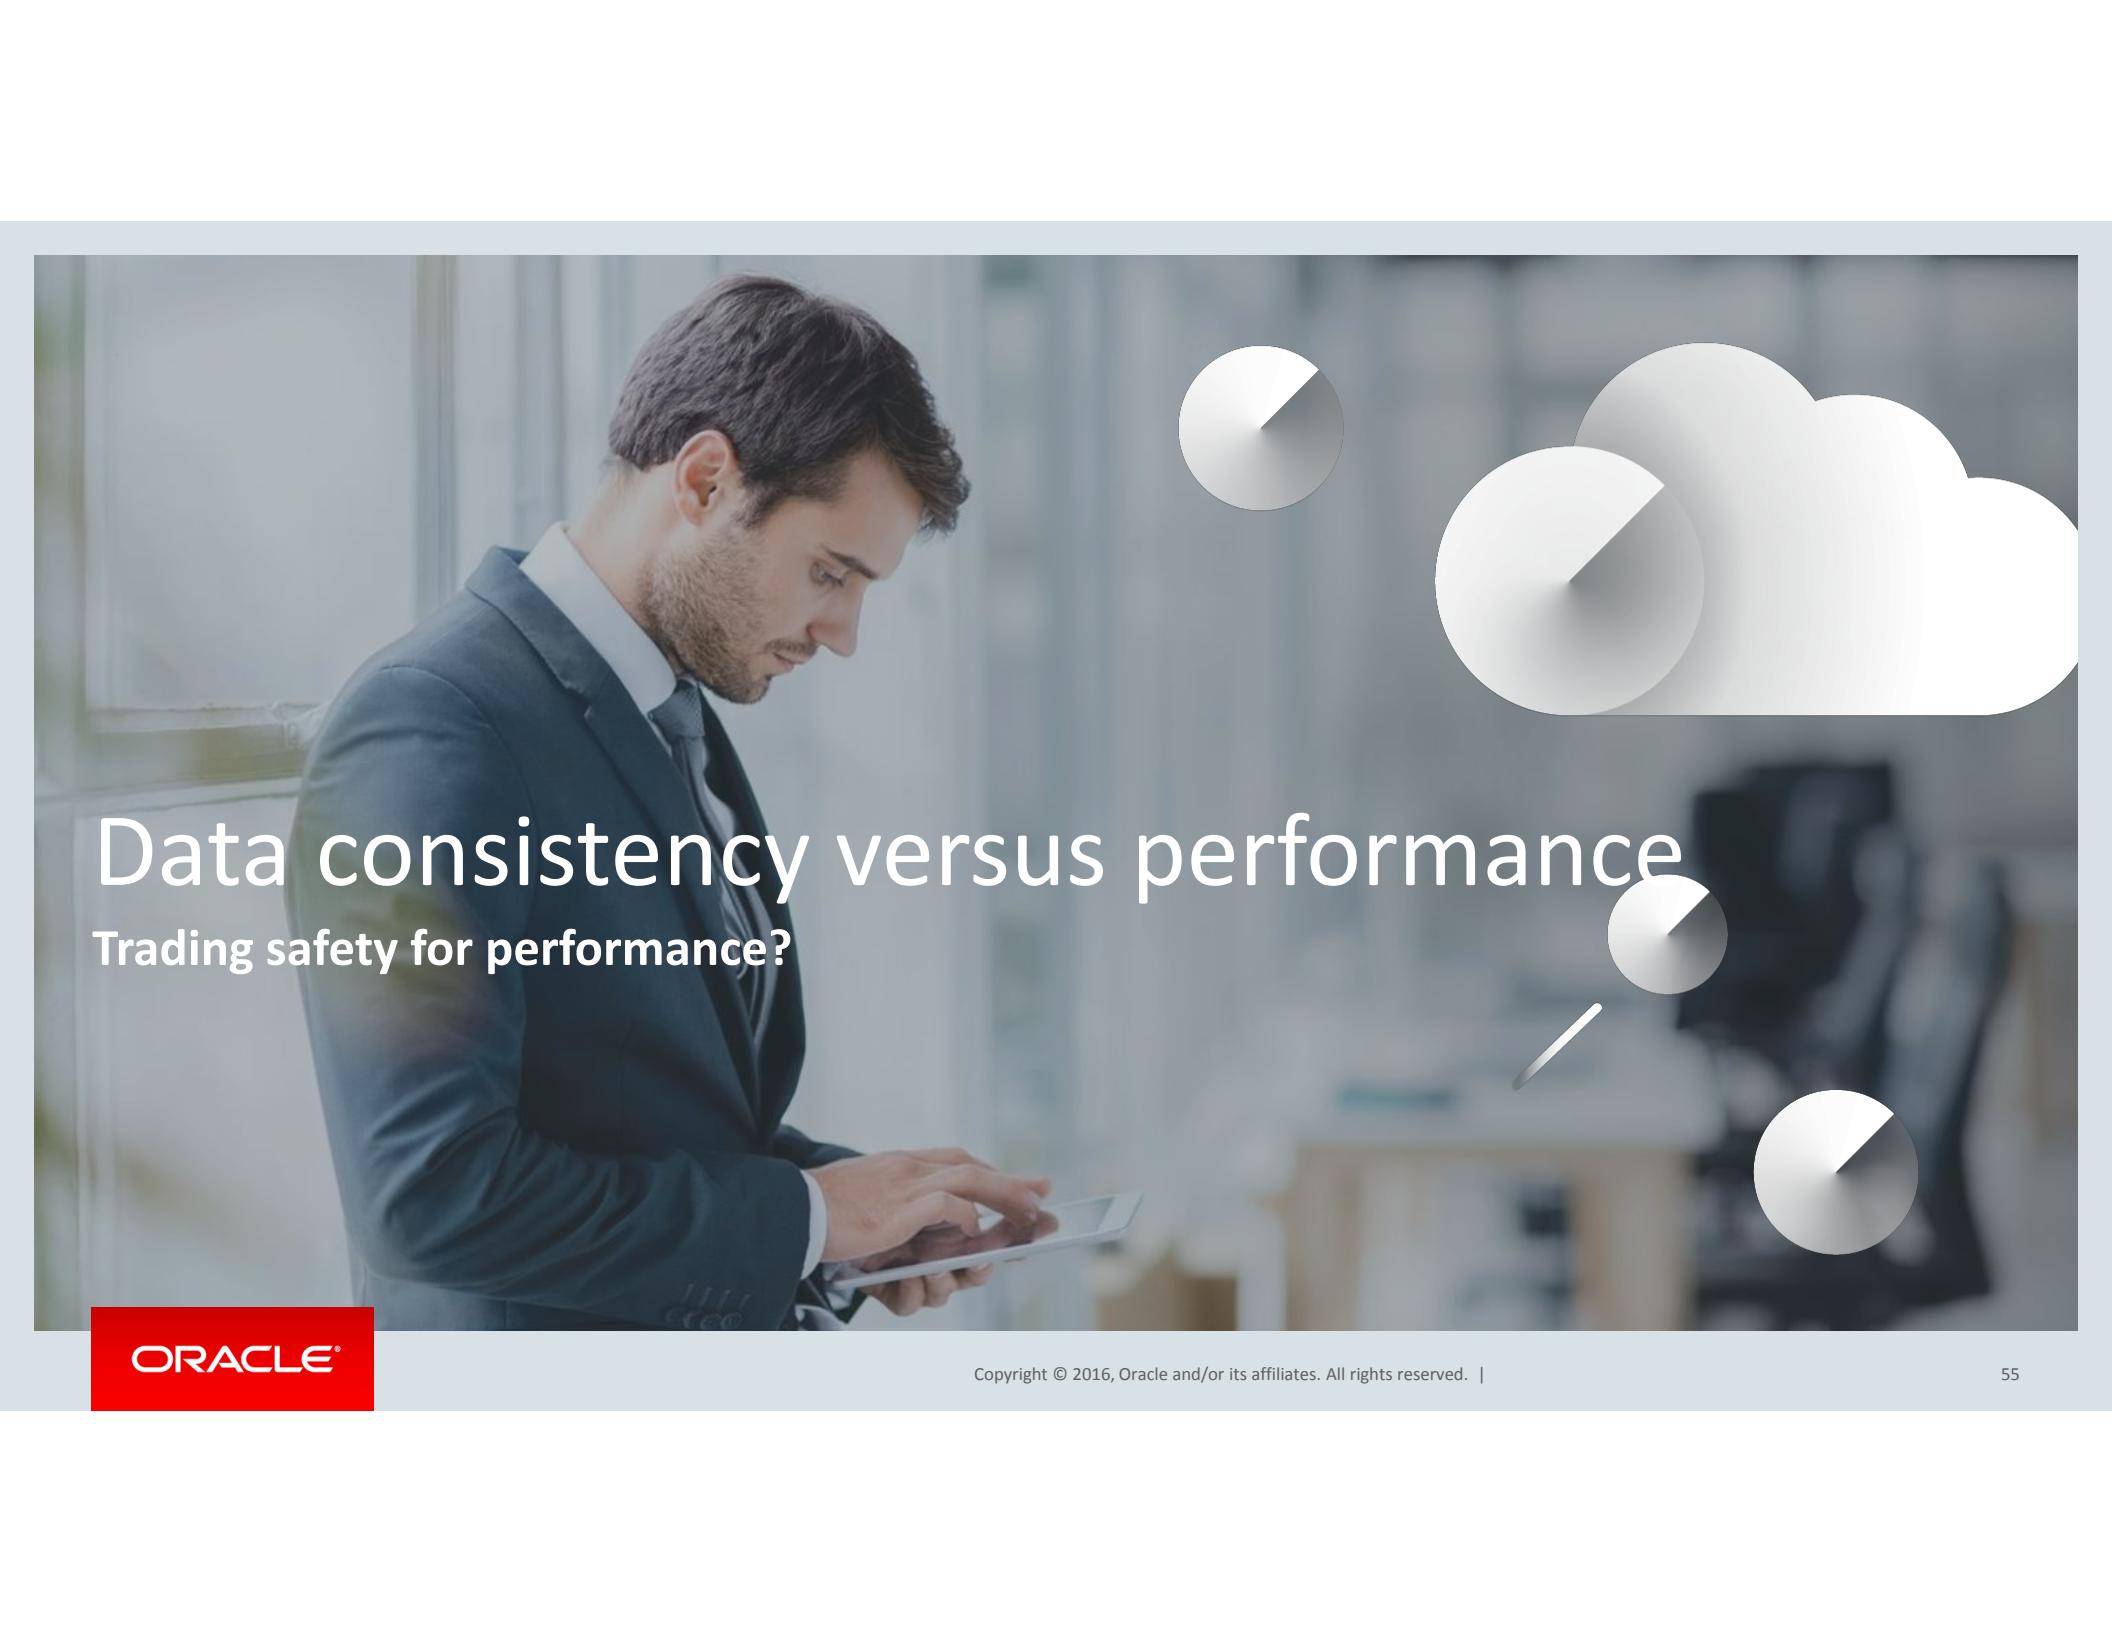
```
mysql> SHOW GLOBAL STATUS LIKE 'Sort_merge_passes';
+-----+-----+
| Variable_name      | Value |
+-----+-----+
| Sort_merge_passes | 0     |
+-----+-----+
1 row in set (0.00 sec)
```

- Aim at a few Sort\_merge\_passes per second on a busy server
- Increase for a session as needed

# Buffers And Caches

## Why is it important with small global values?

- Some buffers are allocated in full each time they are used
- One query may use several buffers, so large buffers can cause significant memory usage
- Memory allocations are relatively expensive
- For example Linux glibc malloc changes memory allocation algorithm when crossing threshold (typical 256k or 512k).
  - Algorithm for larger allocations can be 40 times slower than for smaller allocation

A professional man in a dark blue suit and tie is looking down at a white tablet device he is holding in his hands. He has short brown hair and a light beard. The background is a blurred office environment. Four white pie charts of varying sizes are floating in the upper right area of the slide.

# Data consistency versus performance

Trading safety for performance?

ORACLE®

Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |

55

# Data Consistency Versus Performance

- The two options most commonly considered:
  - innodb\_flush\_log\_at\_trx\_commit
  - sync\_binlog
- There are others as well

# Data Consistency Versus Performance

## innodb\_flush\_log\_at\_trx\_commit

- Possible values in order of data safety:
  - 1: is theoretically the slowest, but with fast SSD it may be around as fast as 2 and 0
  - 2: flushed every second (`innodb_flush_log_at_timeout` in 5.6+)
    - Transactions may be lost if the OS crashes
  - 0: MySQL never fsyncs
    - Transactions may be lost if the OS crashes
- Defaults to 1 – flush the redo log after each commit
  - Required for D in ACID
  - For this reason it is the recommended value

# Data Consistency Versus Performance

## innodb\_flush\_log\_at\_trx\_commit

- What if `innodb_flush_log_at_trx_commit = 1` is too slow and you want the D in ACID?
  - Make sure the redo logs are on separate disks
    - Do not be tempted to have the individual redo logs on different disks
  - Consider SSD particularly if you have a high commit rate
  - Battery-backed disk cache also makes flushes cheaper

# Data Consistency Versus Performance

## sync\_binlog

- Specifies how many transactions between each flush of the binary log
  - 0: when rotating the binary log or when the OS decides
  - 1: at every transaction commit – this is the safest
  - N: every N commits
- Default value:
  - MySQL 5.6 and earlier: 0
  - MySQL 5.7 and later: 1
- MySQL 5.6 and later support group commit for InnoDB giving less overhead of sync\_binlog = 1

# Data Consistency Versus Performance

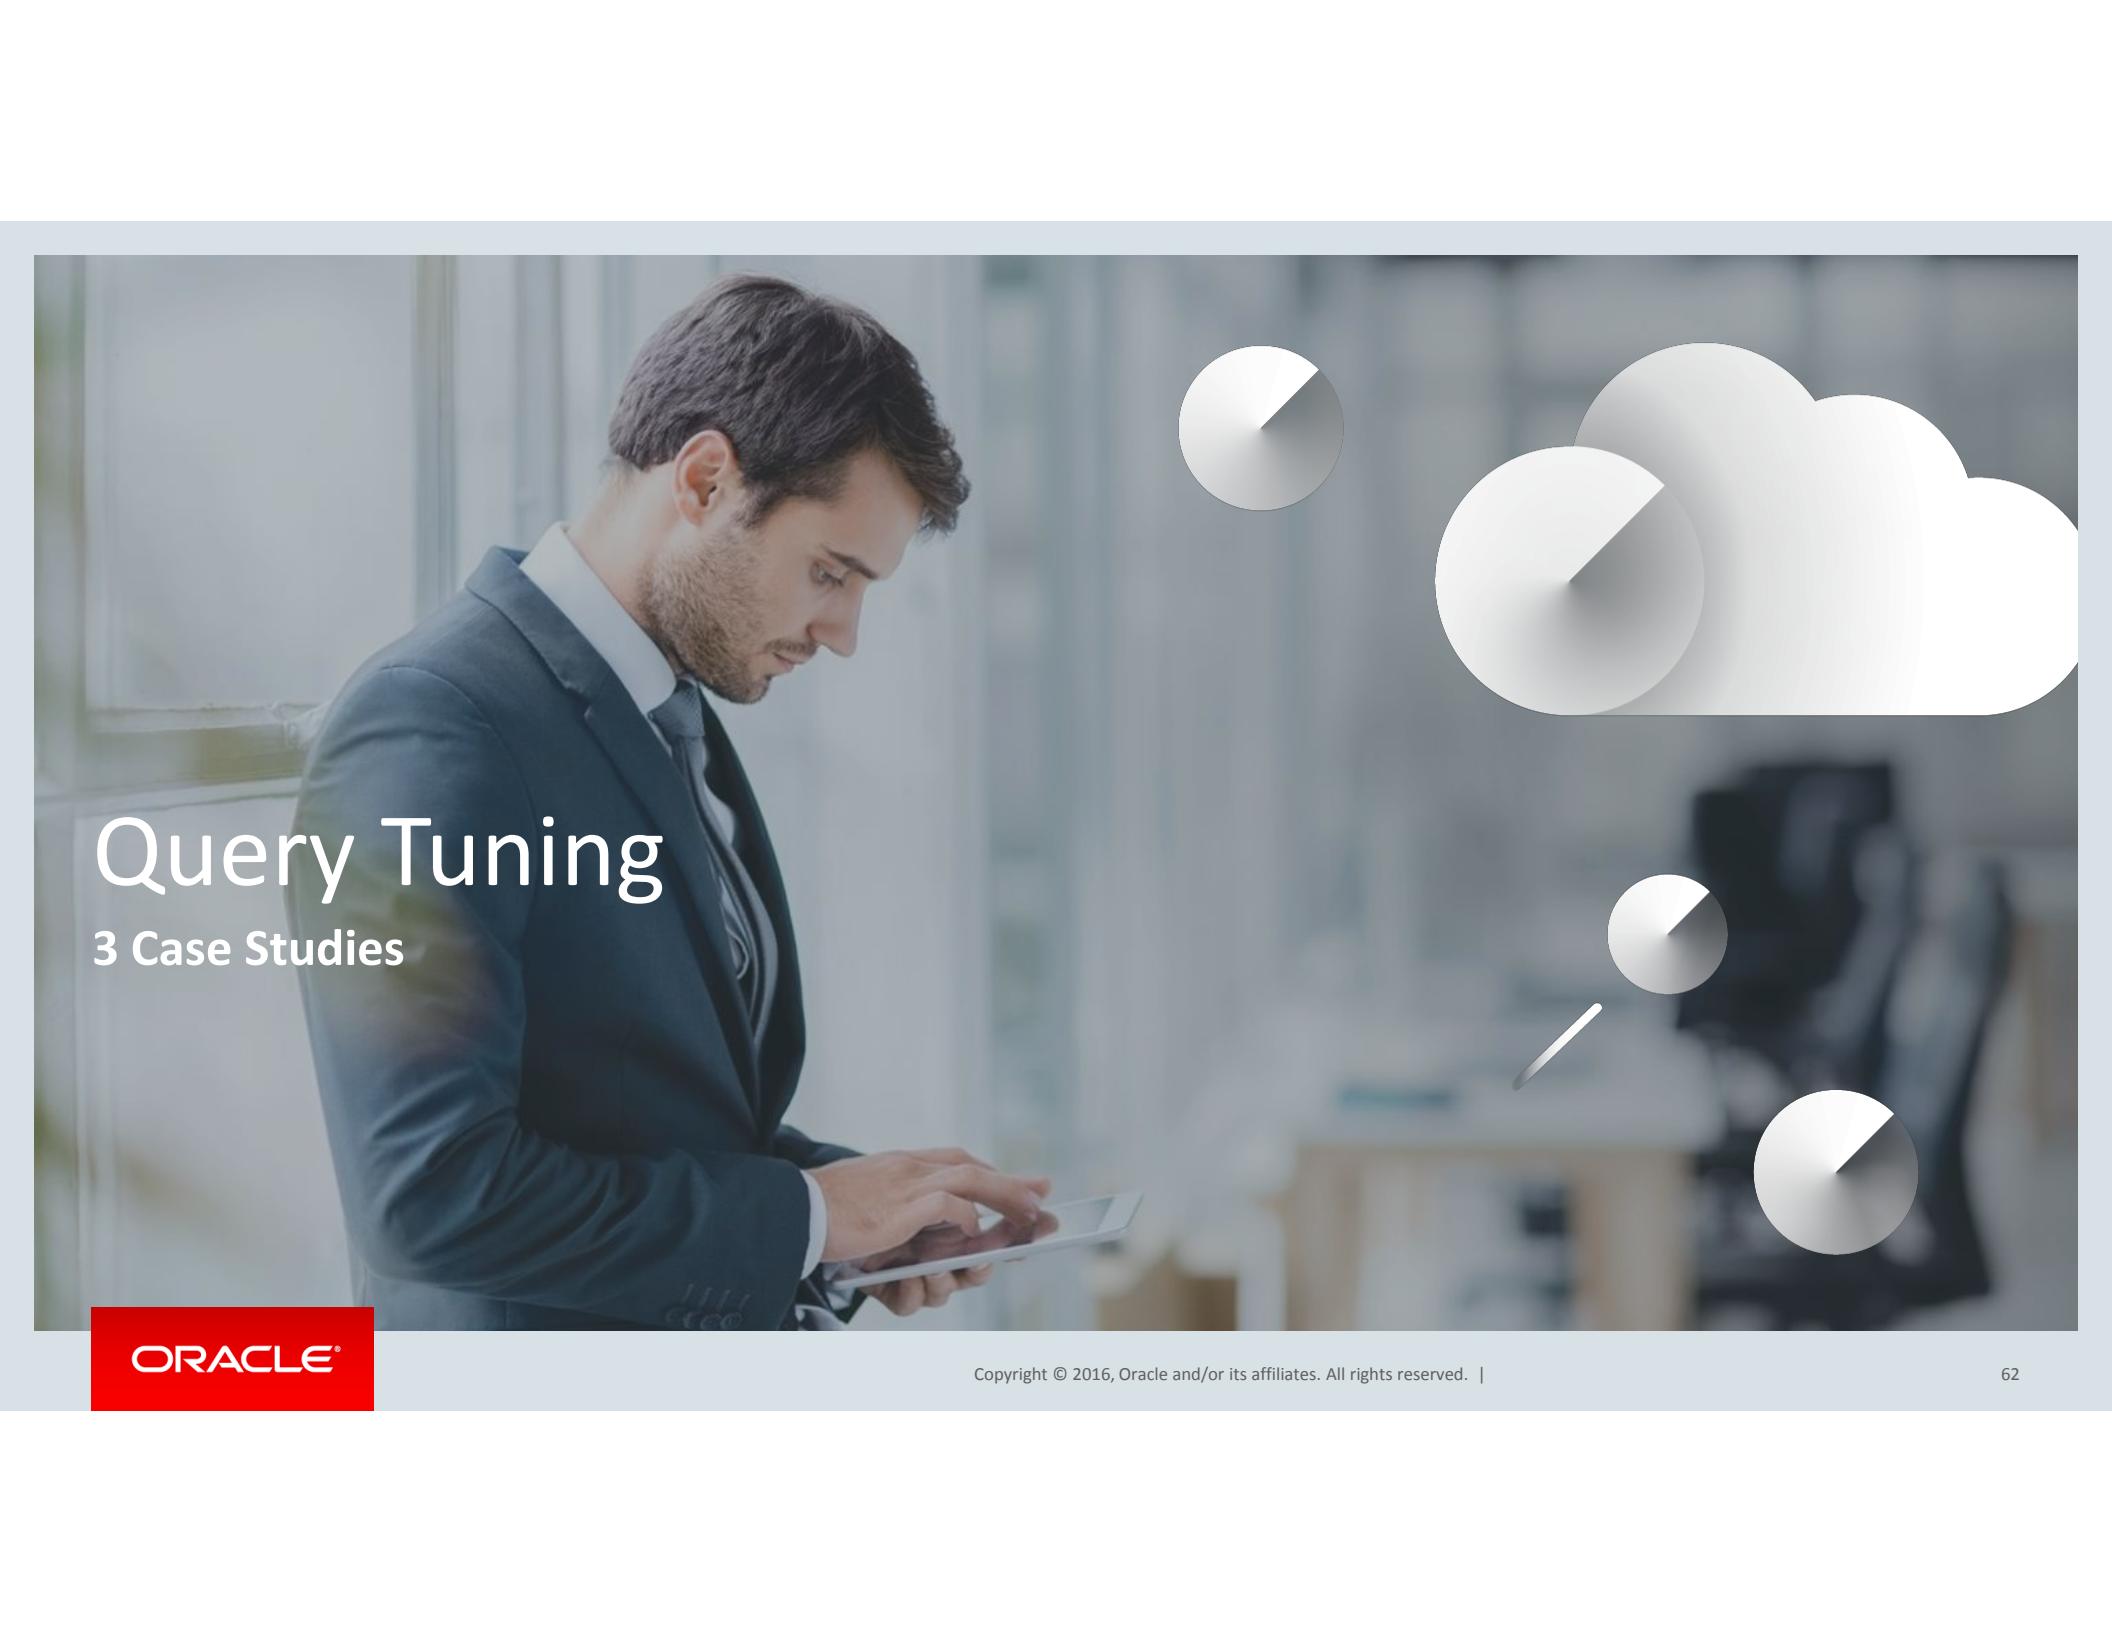
## sync\_binlog

- sync\_binlog != 1 means that slaves most likely will have to be rebuilt if the master crashes
- But sync\_binlog = 0 must be the best performance wise, right?

# Data Consistency Versus Performance

## sync\_binlog

- `sync_binlog != 1` means that slaves most likely will have to be rebuilt if the master crashes
- But `sync_binlog = 0` must be the best performance wise, right?
  - By default: `max_binlog_size = 1G`
  - 1G is not a lot of memory nowadays
  - So the OS may buffer the entire binary log
  - When rotating the binary log, up to 1G will be flushed to disk
    - Stalls all other commits until the binary log rotation is done
- So in short: `sync_binlog = 0` may give the best throughput, but `sync_binlog = 1` may give the most predictable performance



# Query Tuning

## 3 Case Studies

ORACLE®

Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |

62

# Query Tuning

1. Determine the query to optimize
2. Optimize the query

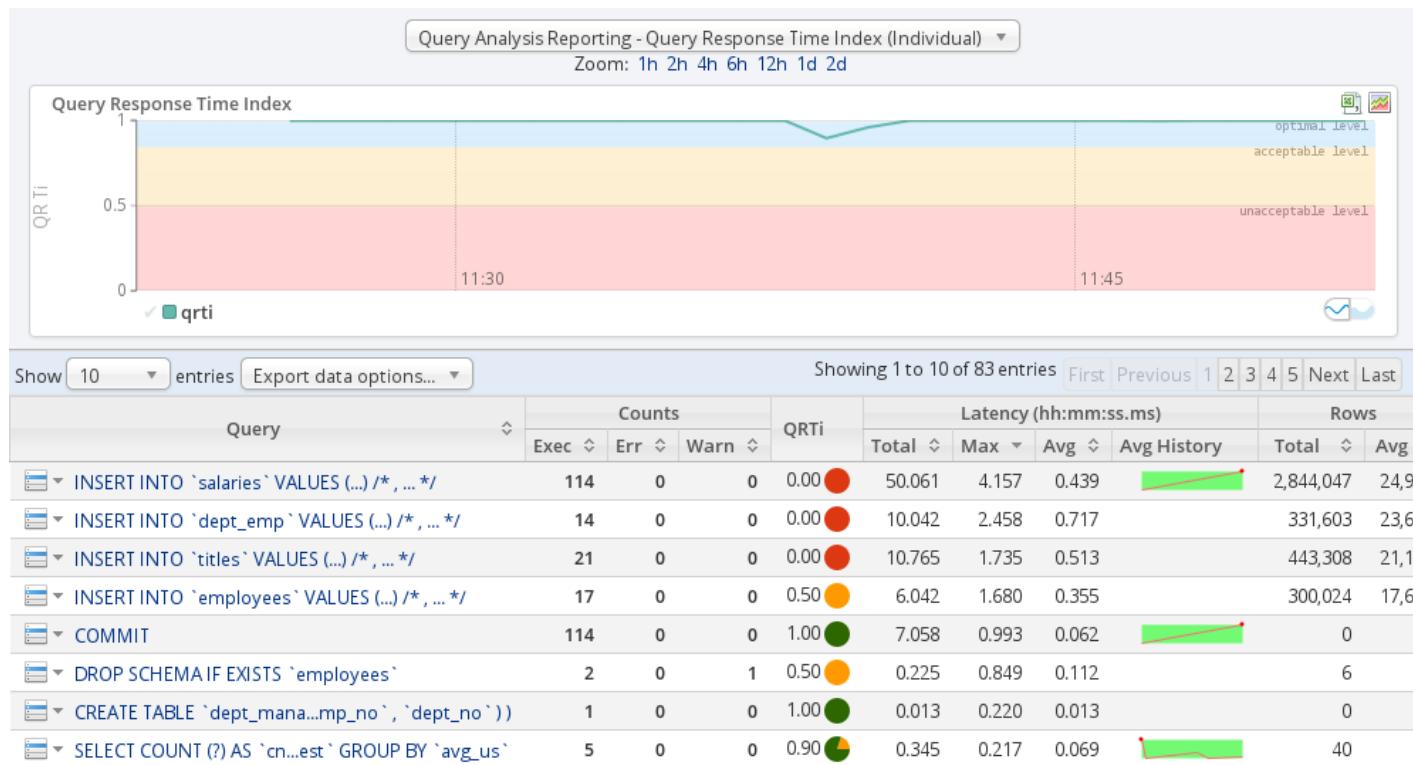
# Query Tuning

## Finding queries to optimize

- Application metrics
- Manually executing query and determining it is too slow or uses too many resources
- Slow Query Log
- MySQL Enterprise Monitor Query Analyzer
- `performance_schema.events_statements_summary_by_digest`
- `sys.statements_with_runtimes_in_95th_percentile`

# Query Tuning

## Finding queries to optimize - MySQL Enterprise Monitor Query Analyzer



# Query Tuning

## Finding queries to optimize – Performance Schema Digest Summary

```
mysql> SELECT LEFT(DIGEST_TEXT, 64) AS DIGEST_TEXT, COUNT_STAR, SUM_TIMER_WAIT, MAX_TIMER_WAIT
    FROM performance_schema.events_statements_summary_by_digest
   ORDER BY MAX_TIMER_WAIT DESC
     LIMIT 10;
+-----+-----+-----+-----+
| DIGEST_TEXT | COUNT_STAR | SUM_TIMER_WAIT | MAX_TIMER_WAIT |
+-----+-----+-----+-----+
| INSERT INTO `salaries` VALUES (...) /* , ... */ |      342 | 159811231808000 | 4156961573000 |
| INSERT INTO `dept_emp` VALUES (...) /* , ... */ |       42 | 31561264335000 | 2458392698000 |
| INSERT INTO `titles` VALUES (...) /* , ... */ |       63 | 35738435708000 | 1735350241000 |
| INSERT INTO `employees` VALUES (...) /* , ... */ |       51 | 18004605187000 | 1679817477000 |
| INSERT INTO `sbtest` ( `k` , `c` , `pad` ) VALUES (...) /* , ... */ |      10 | 5241286782000 | 1247361451000 |
| COMMIT |      342 | 31984662051000 | 992714081000 |
| DROP SCHEMA IF EXISTS `employees` |        6 | 1252459420000 | 848771265000 |
| CREATE TABLE `sbtest` ( `id` INTEGER UNSIGNED NOT NULL AUTO_INCREMENT |       1 | 565468324000 | 565468324000 |
| CREATE TABLE `dept_manager` ( `dept_no` CHARACTER (?) NOT NULL , |       3 | 355874700000 | 220491035000 |
| SELECT COUNT (?) AS `cnt` , `round`( ( `performance_schema` . ` |       6 | 386062170000 | 217206520000 |
+-----+-----+-----+-----+
10 rows in set (0.00 sec)
```

# Query Tuning

## Finding queries to optimize – 95<sup>th</sup> percentile

```
mysql> SELECT query, exec_count, total_latency, max_latency
      FROM sys.statements_with_runtimes_in_95th_percentile;
+-----+-----+-----+-----+
| query | exec_count | total_latency | max_latency |
+-----+-----+-----+-----+
| INSERT INTO `dept_emp` VALUES (...) /* , ... */ | 42 | 31.56 s | 2.46 s |
| INSERT INTO `titles` VALUES (...) /* , ... */ | 63 | 35.74 s | 1.74 s |
| CREATE TABLE `sbtest` ( `id` I ... ( `id` ) ) ENGINE = `innodb` | 1 | 565.47 ms | 565.47 ms |
| INSERT INTO `sbtest` ( `k` , `... d` ) VALUES (...) /* , ... */ | 10 | 5.24 s | 1.25 s |
| INSERT INTO `salaries` VALUES (...) /* , ... */ | 342 | 2.66 m | 4.16 s |
| INSERT INTO `employees` VALUES (...) /* , ... */ | 51 | 18.00 s | 1.68 s |
| DROP SCHEMA IF EXISTS `employees` | 6 | 1.25 s | 848.77 ms |
| LOAD DATA LOCAL INFILE ? INTO TABLE `City` CHARACTER SET `utf8` | 1 | 167.90 ms | 167.90 ms |
+-----+-----+-----+-----+
8 rows in set (0.01 sec)
```

# Query Tuning

## Optimize the query

- EXPLAIN is your best friend:
  - Shows information about the query plan
  - Tells where indexes can be used and not used
- In MySQL 5.6 and earlier use EXPLAIN EXTENDED/EXPLAIN PARTITIONS
- In MySQL 5.6+ use FORMAT=JSON for more information

# Query Tuning

## Optimize the query – Example 1

```
CREATE TABLE `City` (
  `ID` int(11) NOT NULL AUTO_INCREMENT,
  `Name` char(35) NOT NULL DEFAULT '',
  `CountryCode` char(3) NOT NULL DEFAULT '',
  `District` char(20) NOT NULL DEFAULT '',
  `Population` int(11) NOT NULL DEFAULT '0',
  PRIMARY KEY (`ID`)
) ENGINE=InnoDB AUTO_INCREMENT=4080 DEFAULT CHARSET=latin1

CREATE TABLE `Country` (
  `Code` char(3) NOT NULL DEFAULT '',
  `Name` char(52) NOT NULL DEFAULT '',
  `Continent` enum('Asia','Europe','North America',...) NOT NULL DEFAULT 'Asia',
  `Region` char(26) NOT NULL DEFAULT '',
  ...
  PRIMARY KEY (`Code`)
) ENGINE=InnoDB DEFAULT CHARSET=latin1
1 row in set (0.00 sec)
```

# Query Tuning

## Optimize the query – Example 1

```
mysql> SELECT sys.ps_thread_id(NULL); -- Gets the Performance Schema THREAD_ID for the current connection
+-----+
| sys.ps_thread_id(NULL) |
+-----+
|          11419 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT ci.Name, co.Name AS Country, ci.Population
      FROM Country co
            STRAIGHT_JOIN City ci ON ci.CountryCode = co.Code
      WHERE co.Continent = 'North America'
      ORDER BY ci.Population DESC
      LIMIT 5;
+-----+-----+-----+
| Name          | Country        | Population |
+-----+-----+-----+
| Ciudad de México | Mexico          |     8591309 |
| New York       | United States   |     8008278 |
| Los Angeles    | United States   |     3694820 |
| Chicago         | United States   |     2896016 |
| La Habana       | Cuba            |     2256000 |
+-----+-----+-----+
5 rows in set (0.05 sec)
```

# Query Tuning

## Optimize the query – Example 1

```
mysql> SELECT * FROM performance_schema.events_statements_current WHERE THREAD_ID = 11419\G
/* Some columns removed */
      THREAD_ID: 11419
    TIMER_WAIT: 53361516000
        LOCK_TIME: 392000000
          SQL_TEXT: SELECT ci.Name, co.Name AS Country, ... DESC LIMIT 5
            DIGEST: d27cbfe06a40d745a5c498f7f5824529
        DIGEST_TEXT: SELECT `ci` . `Name` , `co` . `Name` AS `Country` , ... DESC LIMIT ?
            ERRORS: 0
            WARNINGS: 0
            ROWS_SENT: 5
            ROWS_EXAMINED: 4904
CREATED_TMP_DISK_TABLES: 0
  CREATED_TMP_TABLES: 1
    SELECT_FULL_JOIN: 1
SELECT_FULL_RANGE_JOIN: 0
    SELECT_RANGE: 0
SELECT_RANGE_CHECK: 0
    SELECT_SCAN: 1
  SORT_MERGE_PASSES: 0
    SORT_RANGE: 0
    SORT_ROWS: 5
    SORT_SCAN: 1
    NO_INDEX_USED: 1
NO_GOOD_INDEX_USED: 0
```

# Query Tuning

## Optimize the query – Example 1

```
mysql> SELECT * FROM sys.session WHERE thd_id = 11419\G
/* Some columns removed */
***** 1. row ****
    thd_id: 11419
    conn_id: 11394
        user: root@localhost
  lock_latency: 392.00 us
rows_examined: 4904
  rows_sent: 5
rows_affected: 0
  tmp_tables: 1
tmp_disk_tables: 0
    full_scan: YES
last_statement: SELECT ci.Name, co.Name AS Cou ... BY ci.Population DESC LIMIT 5
last_statement_latency: 53.36 ms
  current_memory: 72.13 KiB
      last_wait: NULL
last_wait_latency: NULL
      source: NULL
    trx_latency: NULL
    trx_state: NULL
  trx_autocommit: NULL
          pid: 31971
program_name: mysql
```

# Query Tuning

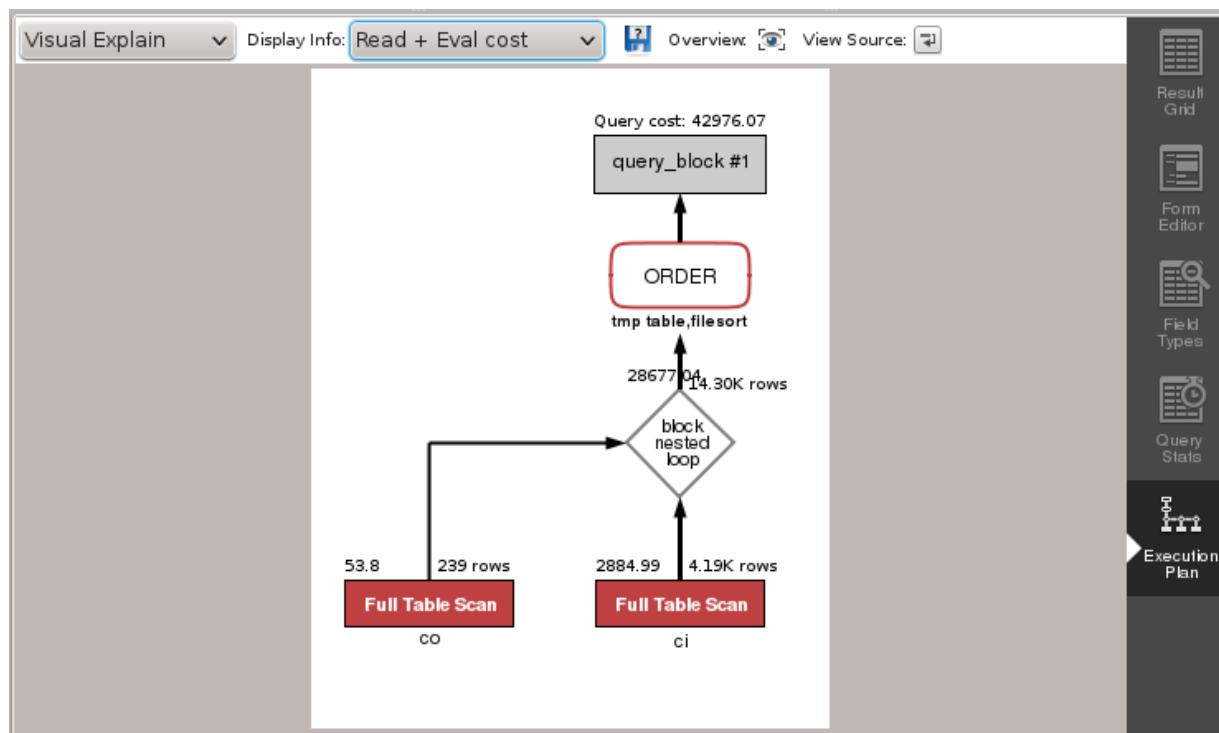
## Optimize the query – Example 1

```
mysql> EXPLAIN SELECT ci.Name, co.Name AS Country, ci.Population
   FROM Country co
        STRAIGHT_JOIN City ci ON ci.CountryCode = co.Code
 WHERE co.Continent = 'North America'
 ORDER BY ci.Population DESC
 LIMIT 5;
+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | rows | Extra
+-----+-----+-----+-----+-----+-----+
|  1 | SIMPLE      | co    | ALL  | PRIMARY       | NULL | 239  | Using where; Using temporary; Using filesort
|  1 | SIMPLE      | ci    | ALL  | NULL          | NULL | 4188 | Using where; Using join buffer (Block Nested Loop)
+-----+-----+-----+-----+-----+-----+
2 rows in set, 1 warning (0.00 sec)

mysql> SHOW WARNINGS\G
***** 1. row *****
Level: Note
Code: 1003
Message: /* select#1 */ select `world`.`ci`.`Name` AS `Name`, `world`.`co`.`Name` AS `Country`, `world`.`ci`.`Population` AS `Population` from `world`.`Country` `co` straight_join `world`.`City` `ci` where ((`world`.`ci`.`CountryCode` = `world`.`co`.`Code`) and (`world`.`co`.`Continent` = 'North America')) order by `world`.`ci`.`Population` desc limit 5
1 row in set (0.00 sec)
```

# Query Tuning

## Optimize the query – Example 1



# Query Tuning

## Optimize the query – Example 1

- The EXPLAIN output shows:
  - No indexes used for joins (full table scans)
  - No index used for WHERE clause
  - No index used for ORDER BY
- First consider an index useful for the JOIN:
  - FROM Country co STRAIGHT\_JOIN City ci ON ci.CountryCode = co.Code
  - Add an index in the City.CountryCode column:
    - ALTER TABLE City ADD INDEX (`CountryCode`);

# Query Tuning

## Optimize the query – Example 1

- New EXPLAIN:

id	table	type	possible_keys	key	key_len	ref	rows	Extra
1	co	ALL	PRIMARY	NULL	NULL	NULL	239	Using where; ... (same) ...
1	ci	ref	CountryCode	CountryCode	3	world.co.Code	18	NULL

- Much better, but still no index used for the Country table

- Try add:
  - An index on the Country.Continent column for the WHERE clause
  - An index on the City.Population column for the ORDER BY
- ALTER TABLE Country ADD INDEX (Continent);
- ALTER TABLE City ADD INDEX (Population);

# Query Tuning

## Optimize the query – Example 1

- New EXPLAIN:

```
+-----+-----+-----+-----+-----+-----+-----+
| table | type  | possible_keys | key      | key_len | ref     | rows   | Extra
+-----+-----+-----+-----+-----+-----+-----+
| co    | ref   | PRIMARY,Continent | Continent | 1       | const   | 37    | Using temporary; Using filesort |
| ci    | ref   | CountryCode       | CountryCode | 3       | world.co.Code | 18    | NULL
+-----+-----+-----+-----+-----+-----+-----+
```

- Good for the Country table
- But no help for the ORDER BY
  - MySQL can only use one index for each table
  - Adding (CountryCode, Population) will not help as City table is the second table
    - ALTER TABLE City ADD INDEX (CountryCode, Population);

# Query Tuning

## Optimize the query – Example 1

- New EXPLAIN:

table	type	possible_keys	key	ref	rows	Extra
co	ref	PRIMARY,Continent	Continent	const	37	Using temporary; Using filesort
ci	ref	CountryCode,CountryCode_2	CountryCode	world.co.Code	18	NULL

- What if the JOIN order is reversed?

```
SELECT ci.Name, co.Name AS Country, ci.Population
  FROM City ci
    STRAIGHT_JOIN Country co ON ci.CountryCode = co.Code
 WHERE co.Continent = 'North America'
 ORDER BY ci.Population ASC
 LIMIT 5;
```

# Query Tuning

## Optimize the query – Example 1

- New EXPLAIN:

```
mysql> EXPLAIN SELECT ci.Name, co.Name AS Country, ci.Population
      FROM City ci
            STRAIGHT_JOIN Country co ON ci.CountryCode = co.Code
      WHERE co.Continent = 'North America'
      ORDER BY ci.Population ASC
      LIMIT 5;
+-----+-----+-----+-----+-----+-----+
| table | type  | possible_keys          | key       | key_len | ref           | rows | Extra        |
+-----+-----+-----+-----+-----+-----+
| ci    | ALL   | CountryCode,CountryCode_2 | NULL     | NULL    | NULL          | 4188 | Using filesort |
| co    | eq_ref| PRIMARY,Continent      | PRIMARY  | 3       | world.ci.CountryCode | 1    | Using where   |
+-----+-----+-----+-----+-----+-----+
```

- So eliminated the internal temporary table (but not related to the indexes)

# Query Tuning

## Optimize the query – Example 1

- Now the City table looks like:

```
CREATE TABLE `City` (
  `ID` int(11) NOT NULL AUTO_INCREMENT,
  `Name` char(35) NOT NULL DEFAULT '',
  `CountryCode` char(3) NOT NULL DEFAULT '',
  `District` char(20) NOT NULL DEFAULT '',
  `Population` int(11) NOT NULL DEFAULT '0',
  PRIMARY KEY (`ID`),
  KEY `CountryCode` (`CountryCode`),
  KEY `CountryCode_2` (`CountryCode`, `Population`),
  KEY `Population`(`Population`)
) ENGINE=InnoDB AUTO_INCREMENT=4080 DEFAULT CHARSET=latin1
```

- Note: two indexes starting with the CountryCode table
  - (CountryCode, Population) can be used instead of (CountryCode) as they have same left prefix.
  - Best to drop one of them

# Query Tuning

## Optimize the query – Example 1

- Side note – you can check for unused indexes:

```
SELECT OBJECT_SCHEMA,
       OBJECT_NAME,
       INDEX_NAME
  FROM performance_schema.table_io_waits_summary_by_index_usage
 WHERE INDEX_NAME IS NOT NULL
   AND COUNT_STAR = 0
   AND OBJECT_SCHEMA != 'mysql'
   AND INDEX_NAME != 'PRIMARY'
 ORDER BY OBJECT_SCHEMA, OBJECT_NAME;
```

- Or:

```
mysql> SELECT * FROM sys.schema_unused_indexes WHERE object_schema = 'world';
+-----+-----+-----+
| object_schema | object_name      | index_name      |
+-----+-----+-----+
| world        | City             | CountryCode    |
| world        | City             | CountryCode_2  |
| world        | City             | Population     |
| world        | CountryLanguage | CountryCode    |
+-----+-----+-----+
```

# Query Tuning

## Optimize the query – Example 2

- Employees sample database – find all employees that has been in more than one department:

```
SELECT *  
  FROM employees  
 WHERE emp_no IN (SELECT emp_no FROM dept_emp GROUP BY emp_no HAVING COUNT(*) > 1);
```

- MySQL 5.5: Killed it after 4 hours

# Query Tuning

## Optimize the query – Example 2

- Employees sample database – find all employees that has been in more than one department:

```
SELECT *
  FROM employees
 WHERE emp_no IN (SELECT emp_no FROM dept_emp GROUP BY emp_no HAVING COUNT(*) > 1);
```

- MySQL 5.5: Killed it after 4 hours
- EXPLAIN – note the dependent subquery:

```
+----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref | rows | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+
|  1 | PRIMARY     | employees | ALL | NULL          | NULL | NULL    | NULL | 300030 | Using where |
|  2 | DEPENDENT SUBQUERY | dept_emp | index | NULL          | emp_no | 4       | NULL |      1 | Using index |
+----+-----+-----+-----+-----+-----+-----+-----+
```

# Query Tuning

## Optimize the query – Example 2

- Potential rewrites:

- Use derived table: 0.15 second

```
SELECT employees.*  
  FROM (SELECT emp_no FROM dept_emp GROUP BY emp_no HAVING COUNT(*) > 1) tmp  
 INNER JOIN employees USING (emp_no);
```

- Regular JOIN: 0.4 second

```
SELECT e.emp_no, e.birth_date, e.first_name, e.last_name, e.gender, e.hire_date  
  FROM employees e  
    INNER JOIN dept_emp USING (emp_no)  
      GROUP BY emp_no  
        HAVING COUNT(*) > 1;
```

# Query Tuning

## Optimize the query – Example 2

- Potential rewrites:

- Use derived table: 0.15 second

```
SELECT employees.*  
  FROM (SELECT emp_no FROM dept_emp GROUP BY emp_no HAVING COUNT(*) > 1) tmp  
 INNER JOIN employees USING (emp_no);
```

- Regular JOIN: 0.4 second

```
SELECT e.emp_no, e.birth_date, e.first_name, e.last_name, e.gender, e.hire_date  
  FROM employees e  
    INNER JOIN dept_emp USING (emp_no)  
      GROUP BY emp_no  
        HAVING COUNT(*) > 1;
```

- Or upgrade to 5.6+ (can materialize the subquery)
  - Original query: 0.48 second

# Query Tuning

## Optimize the query – Example 3

- Schema:

```
CREATE TABLE `t1` (
  `id` int(11) DEFAULT NULL,
  `name` varchar(10) DEFAULT NULL
) ENGINE=InnoDB DEFAULT CHARSET=latin1

CREATE TABLE `t2` (
  `id` int(11) DEFAULT NULL,
  `name` varchar(10) DEFAULT NULL
) ENGINE=InnoDB DEFAULT CHARSET=latin1
```

- Query:

```
mysql> SELECT COUNT(*) FROM t1 INNER JOIN t2 USING (id);
+-----+
| COUNT(*) |
+-----+
|      50000 |
+-----+
1 row in set (2 min 3.36 sec)
```



# Query Tuning

## Optimize the query – Example 3

- EXPLAIN:

table	type	possible_keys	key	key_len	ref	rows	Extra
t1	ALL	NULL	NULL	NULL	NULL	50283	NULL
t2	ALL	NULL	NULL	NULL	NULL	50283	Using where; Using join buffer (Block Nested Loop)

- Obviously best choice is to add indexes, but assume that is not possible

# Query Tuning

## Optimize the query – Example 3

- In 5.6+ rewrite using subqueries
  - In 5.7 it is necessary to switch off the derived\_merge optimizer switch:

```
mysql> SET optimizer_switch='derived_merge=off';
Query OK, 0 rows affected (0.00 sec)
```

- The new query:

```
mysql> SELECT COUNT(*) FROM (SELECT * FROM t1) AS t1 INNER JOIN (SELECT * FROM t2) AS t2 USING (id);
+-----+
| COUNT(*) |
+-----+
|    50000 |
+-----+
1 row in set (0.12 sec)
```

- A factor 1000 speedup! Why?

# Query Tuning

## Optimize the query – Example 3

- New EXPLAIN:

+-----+   id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra	-----+
1   PRIMARY   <derived2>   ALL   NULL   NULL   NULL   NULL   50283   Using where										
1   PRIMARY   <derived3>   ref   <auto_key0>   <auto_key0>   5   t1.id   10   NULL										
3   DERIVED   t2   ALL   NULL   NULL   NULL   NULL   50283   NULL										
2   DERIVED   t1   ALL   NULL   NULL   NULL   NULL   50283   NULL										

- The optimizer added an auto key!

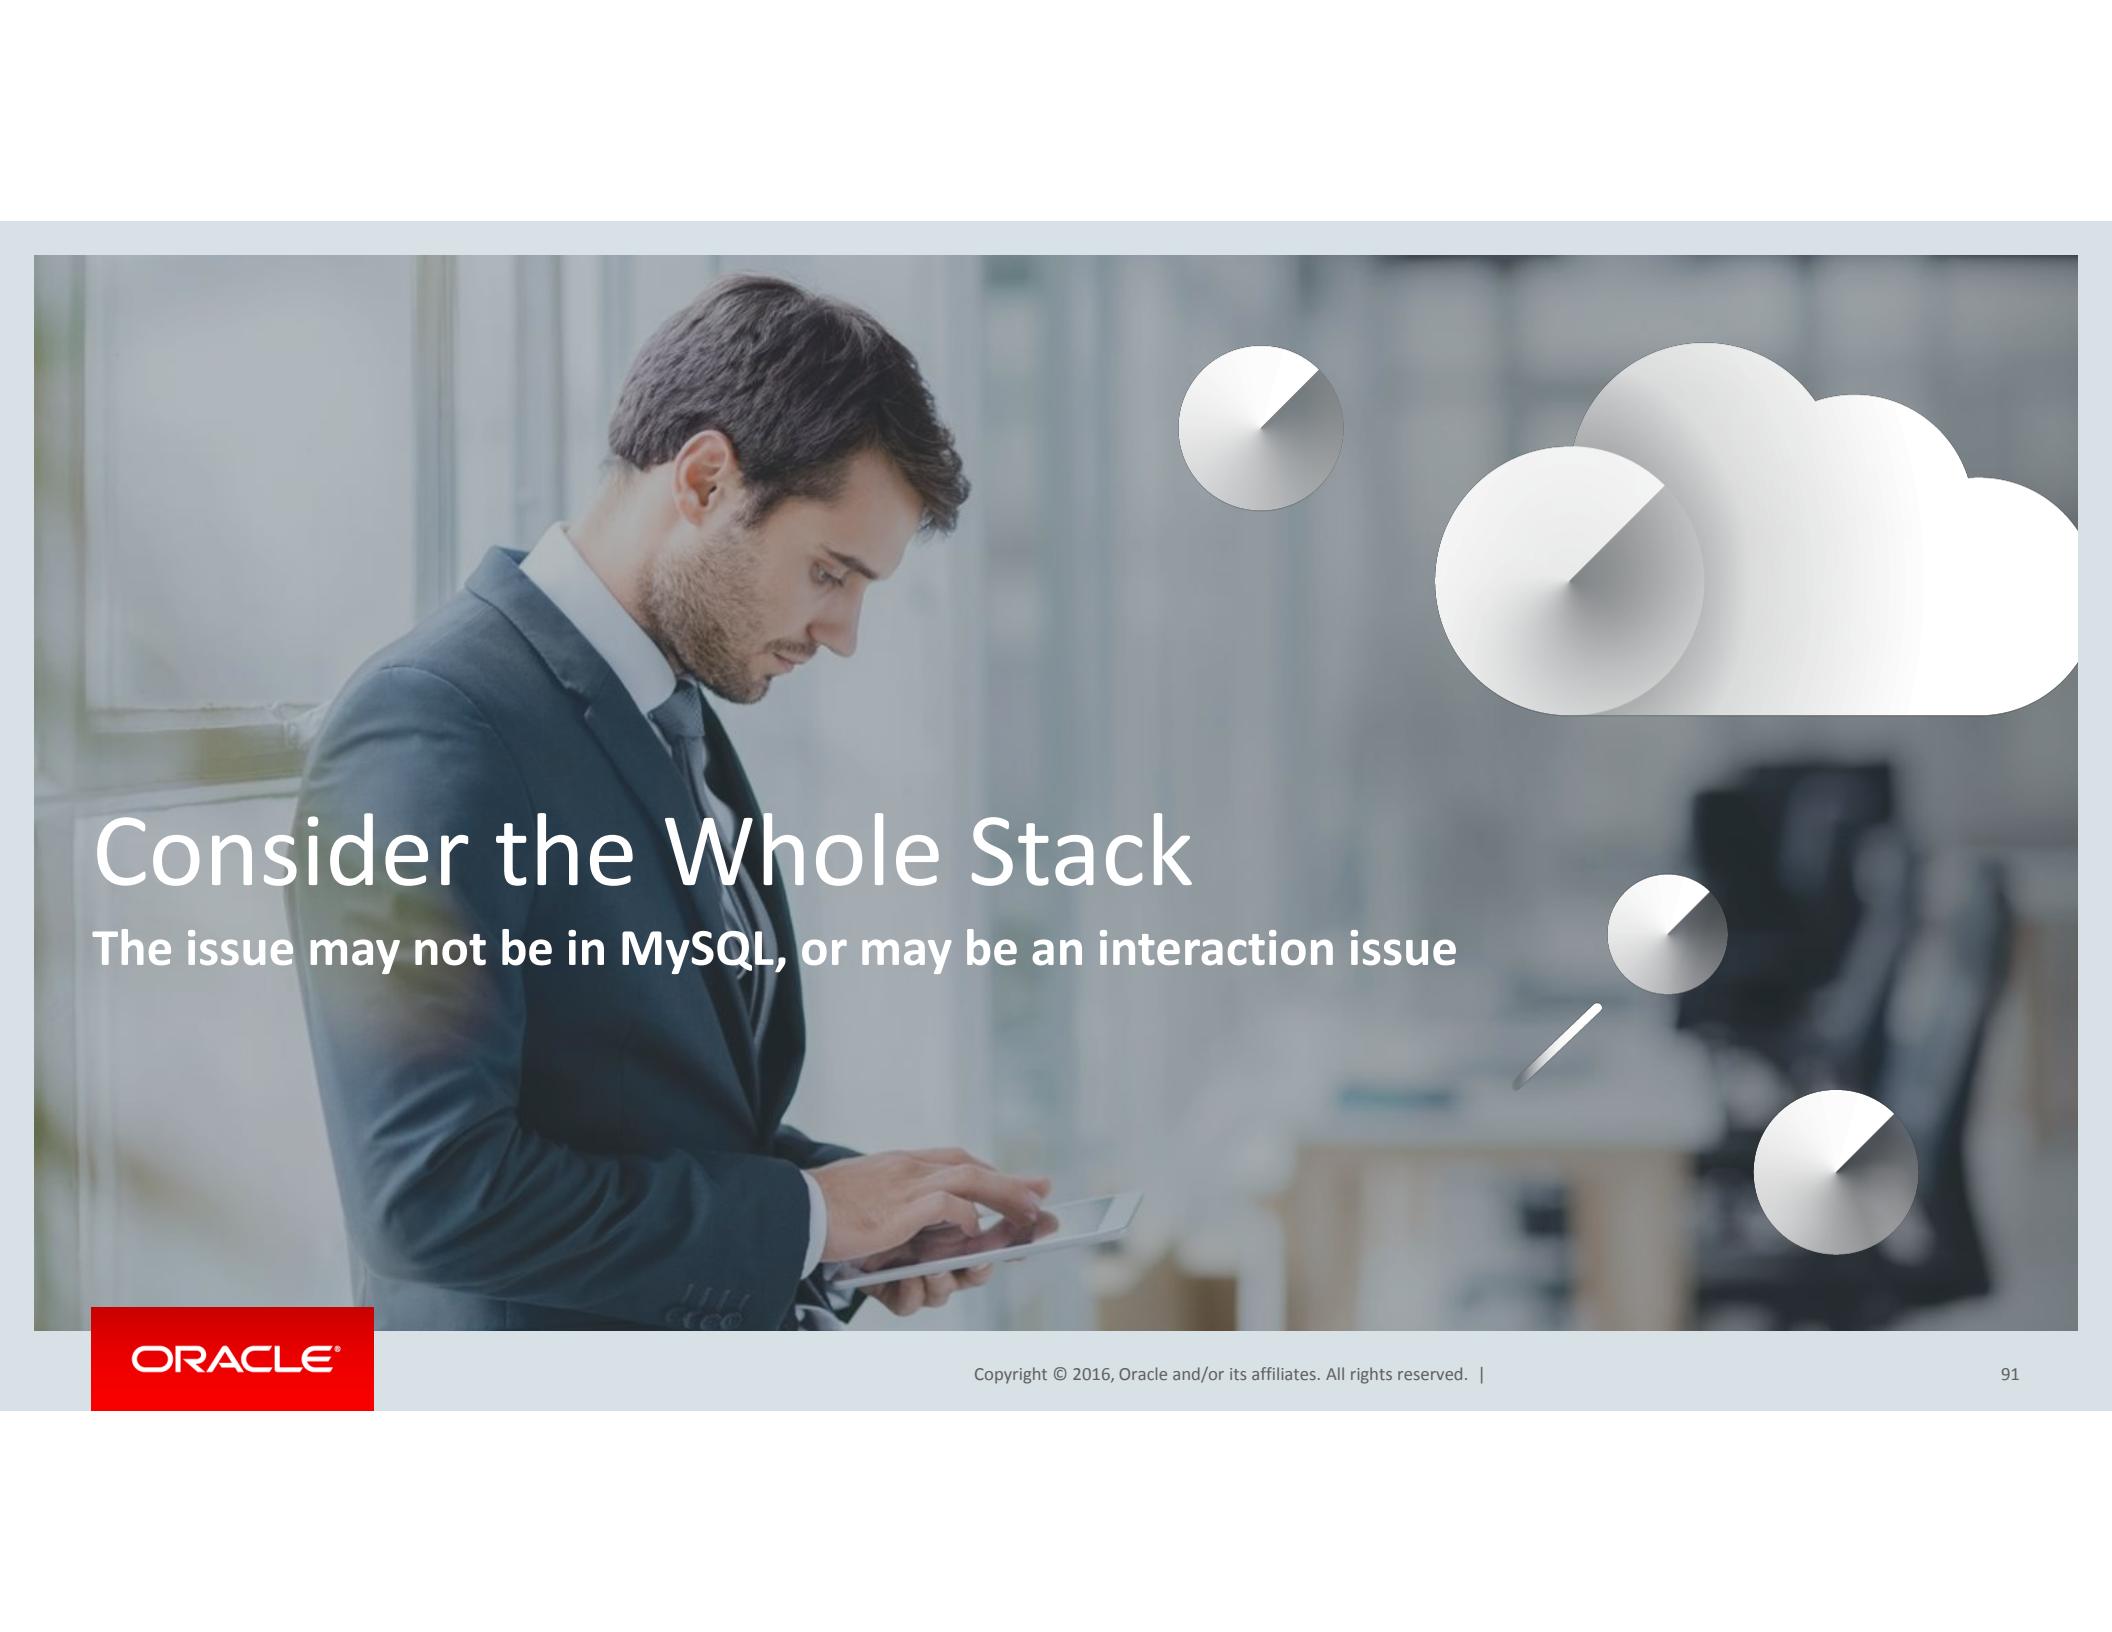
# Query Tuning

## Optimize the query – Example 3

- New EXPLAIN:

+-----+ <th>id</th> <th>select_type</th> <th>table</th> <th>type</th> <th>possible_keys</th> <th>key</th> <th>key_len</th> <th>ref</th> <th>rows</th> <th>Extra</th> <th>-----+</th>	id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra	-----+
1   PRIMARY   <derived2>   ALL   NULL   NULL   NULL   NULL   NULL   50283   Using where											
1   PRIMARY   <derived3>   ref   <auto_key0>   <auto_key0>   5   t1.id   10   NULL											
3   DERIVED   t2   ALL   NULL   NULL   NULL   NULL   NULL   50283   NULL											
2   DERIVED   t1   ALL   NULL   NULL   NULL   NULL   NULL   50283   NULL											

- The optimizer added an auto key!
- Far fetched? A real world use case is for the schema\_table\_statistics view in the sys schema.

A professional man in a dark suit and tie is looking down at a white tablet device he is holding in his hands. He is positioned on the left side of the frame, with a blurred office environment in the background. On the right side of the slide, there are several white, semi-transparent pie charts of different sizes and orientations, along with a large white cloud icon. The overall theme suggests a focus on data analysis and cloud computing.

# Consider the Whole Stack

The issue may not be in MySQL, or may be an interaction issue

# Consider the Whole Stack

- The issue may occur at any place in the stack, e.g.:
  - Application
  - Application host/hardware
  - Network between application and host with MySQL
  - MySQL host/hardware
  - MySQL
- This should also be considered when monitoring
- Consider the OS/hardware settings

# Consider the Whole Stack

## I/O settings on the operating system

- Test with various file systems, for example:
  - Ext4
  - XFS
  - ZFS
- The optimal file system varies with time
- Use a battery backed file system:
  - Allows you to disable the write barrier in Ext4 and XFS

# Consider the Whole Stack

## I/O settings on the operating system

- If the file system guarantees a full page write, the InnoDB doublewrite buffer can be disabled
- ZFS is the only main file system that makes this guarantee

# Consider the Whole Stack

## I/O settings on the operating system

- I/O Scheduler
  - Several Linux distributions use the CFQ scheduler by default
    - OK for reads
    - Serializes writes!
  - NOOP and deadline are usually better for MySQL workloads
    - Deadline is the default I/O scheduler for Oracle Linux

# Consider the Whole Stack

## I/O settings on the operating system

- I/O Scheduler

- Check the current scheduler:

```
shell# cat /sys/block/sda/queue/scheduler  
noop deadline [cfq]
```

- Update the scheduler dynamically:

```
shell# echo deadline > /sys/block/sda/queue/scheduler  
shell# cat /sys/block/sda/queue/scheduler  
noop [deadline] cfq
```

- To set at boot time, use the “elevator=deadline” boot option.

# Consider the Whole Stack

## Memory allocation library

- Linux glibc malloc can be a performance bottleneck
- Better to use alternative malloc library:
  - tcmalloc
  - jemalloc

# Consider the Whole Stack

## Memory allocation library

- Set alternative malloc library:
  - mysqld\_safe
    - [mysqld\_safe]  
malloc-lib = /usr/lib64/libjemalloc.so.1
  - systemd distributions:
    - Set LD\_PRELOAD in /etc/sysconfig/mysql

# Consider the Whole Stack

## Case study – Real world example

- MySQL Cluster, but illustrates well how you cannot focus on one part of the stack
- Data nodes often stalled
- Eventually data nodes were shut down when a stall took longer than 18 seconds by the watchdog to protect the rest of the cluster
- Always when trying to send data through TCP

# Consider the Whole Stack

## Case study – Real world example

- Turned out OS tools such as mpstat, iostat, and cat /proc/interrupts also stalled
- Longest stall was 30+ seconds
- Collected data using the Linux perf tool:
  - For every single stall, a memory compaction was ongoing at the OS level
- Memory very fragmented:

```
shell$ cat /proc/buddyinfo  
...  
Node 0, zone Normal 18232 103 32 1400 4 0 0 0 0 0 0
```

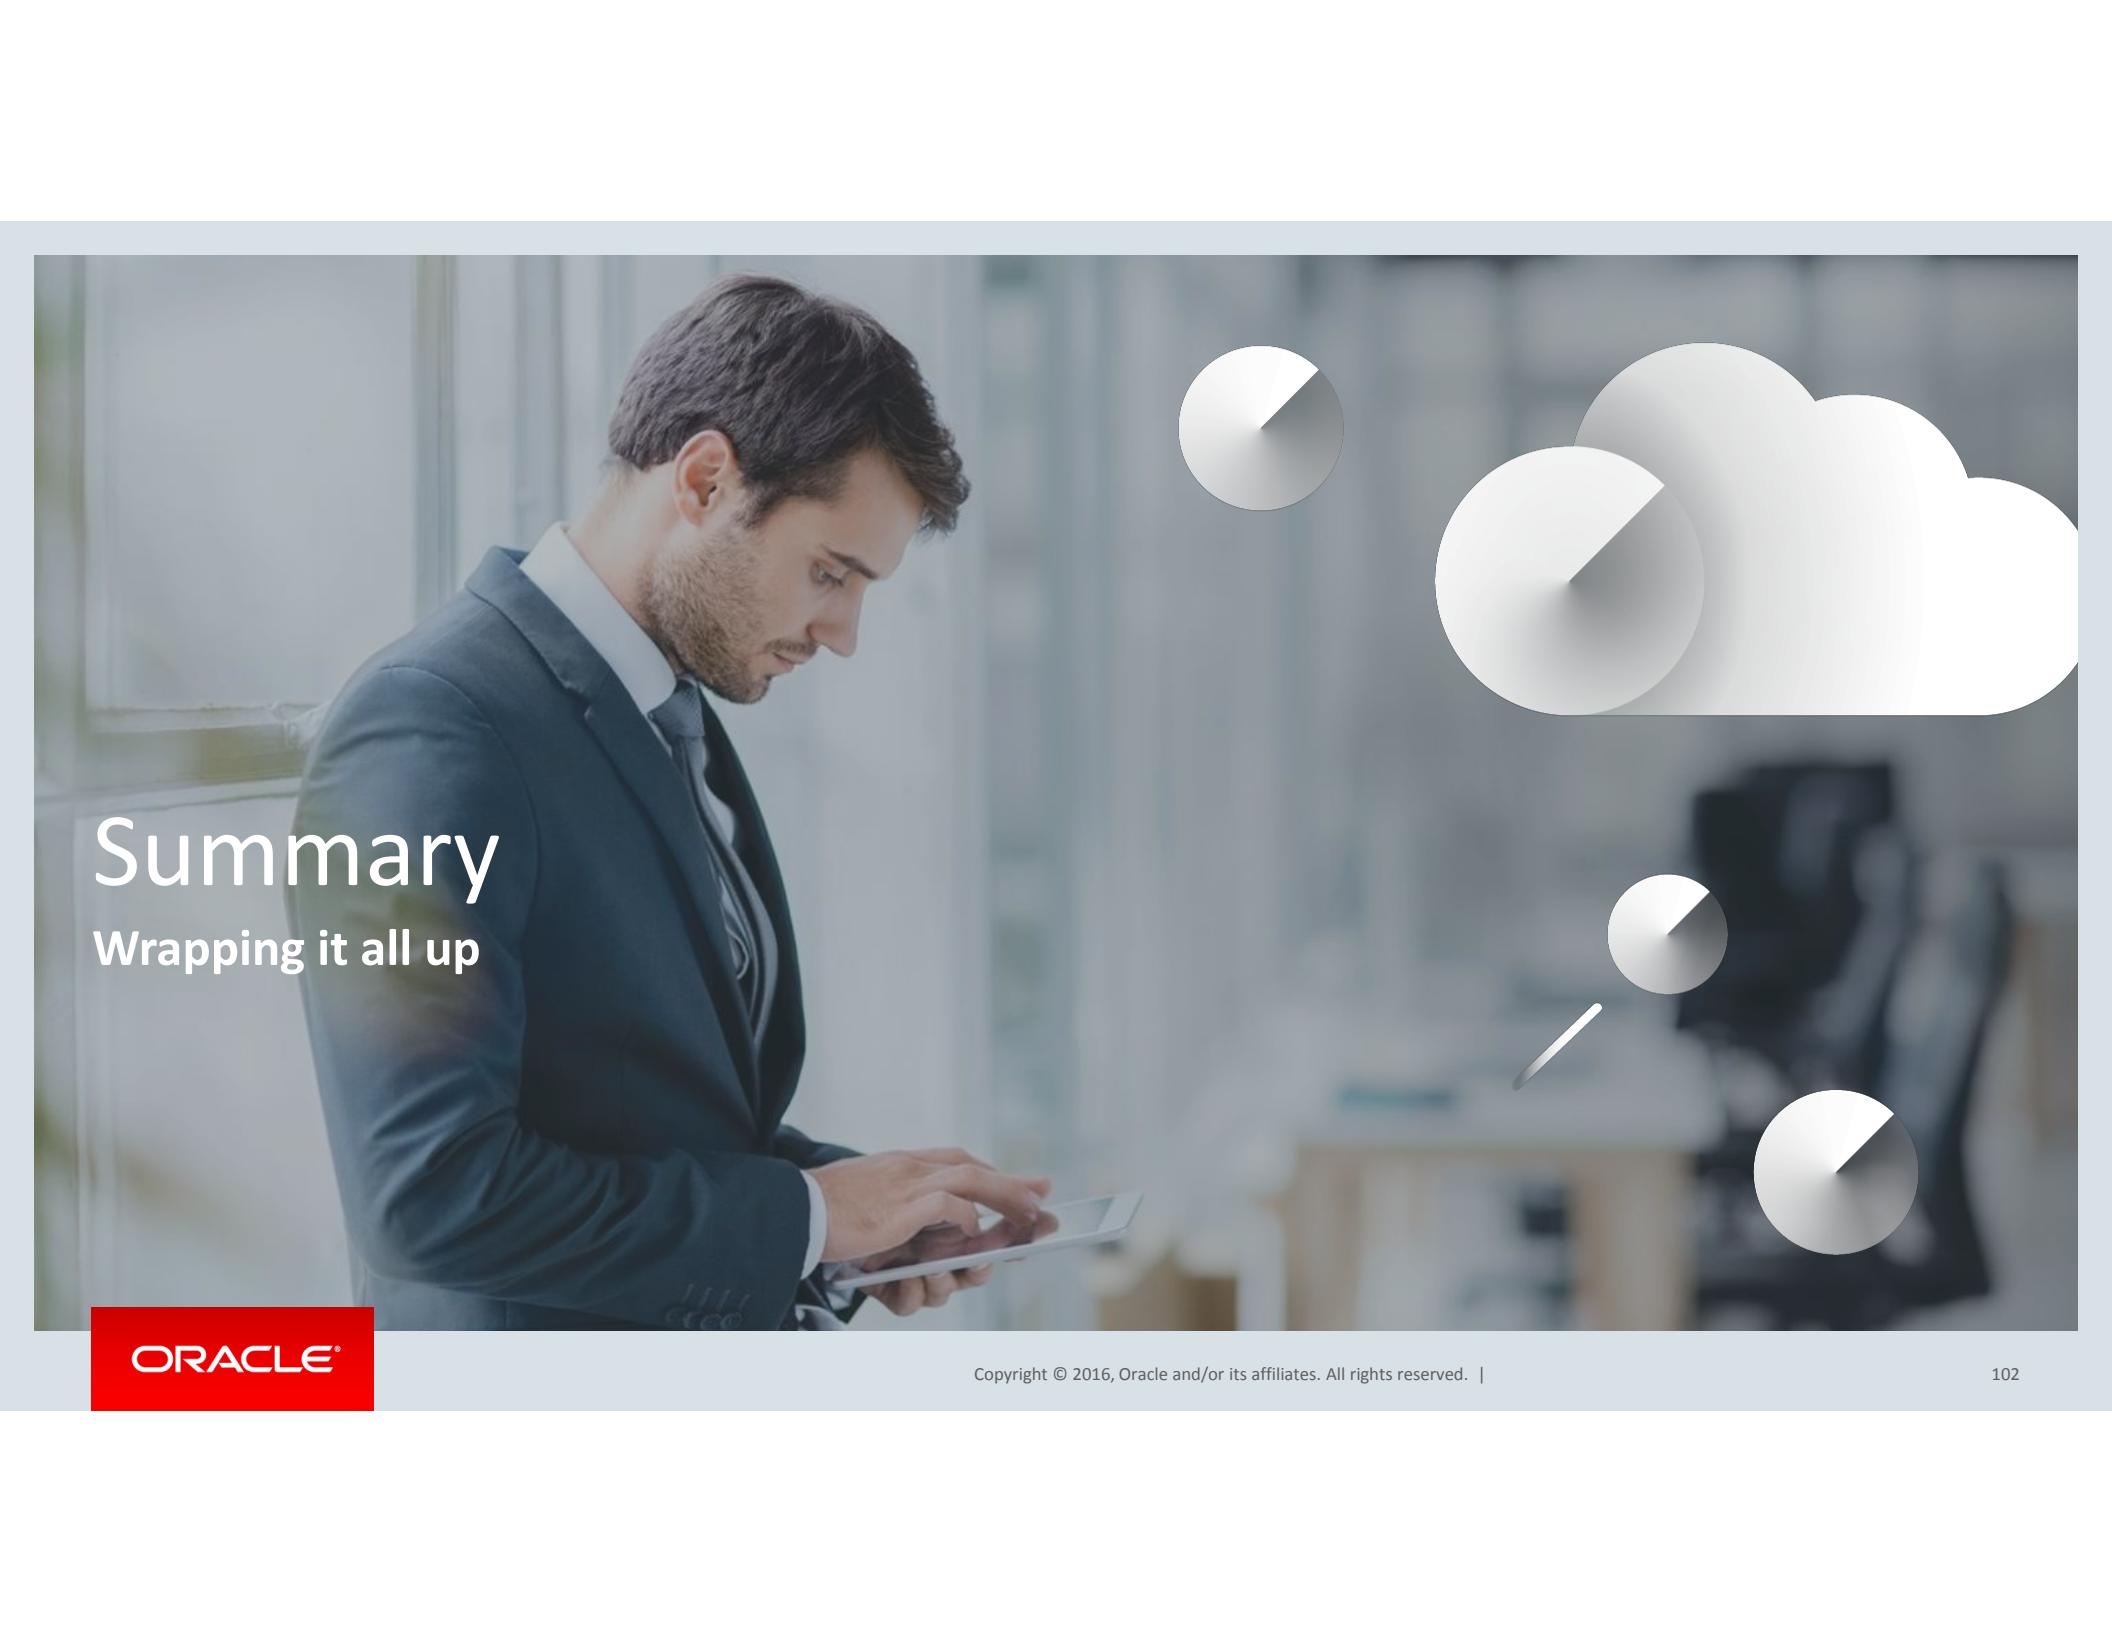
# Consider the Whole Stack

## Case study – Real world example

- MySQL Cluster allocates and touches memory up front, so should not cause fragmentation after startup
- Turned out fragmentation was triggered by the I/O cache at the OS level
- TCP required Direct Memory Access (DMA) which much use contiguous memory, so with fragmented memory triggered memory compaction
- Enabling Direct I/O resolved the issue

# Summary

## Wrapping it all up

A professional man in a dark blue suit and tie is looking down at a white tablet device he is holding in his hands. He is positioned on the left side of the frame, facing right. In the upper right area of the slide, there is a large, semi-transparent white cloud icon. Inside and around the cloud are several white pie charts of different sizes, some with thin lines extending from them. The background is a blurred office environment.

ORACLE®

Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |

102

# Summary

- MySQL performance tuning is similar to all other performance tuning:
  - Premature optimization is bad
  - Only change one thing at a time and don't make too large changes
  - Don't micro manage
  - One size does not fit all
  - Base your decisions on measurements
  - Understand what an option does before changing it
  - Understand your system (data)
  - Understand what you need
  - Consider the whole stack

# Some Related Talks

- Tuesday
  - 4:00 p.m. | Park Central – City
    - MySQL Monitoring: Finding Meaning in Beautiful Graphs [CON2385]
    - Peter Zaitsev, Percona Inc
  - 6:15 p.m. | Park Central – City
    - Meet the MySQL Engineering Team [BOF1967]
- Wednesday
  - 3:00 p.m. | Park Central – Stanford
    - MySQL Performance: Demystified Tuning and Best Practices [CON1993]
    - Dimitri Kravtchuk, Oracle
  - 3:00 p.m. | Park Central – City
    - MySQL Optimizer: What's New in 5.7 and Sneak Peek at 8.0 [CON6112]
    - Manyi Lu, Director Software Engineering, Oracle

## Some Related Talks

- Thursday
  - 9:30 a.m. | Park Central – City
    - MySQL Replication Troubleshooting for Oracle DBAs [CON1685]
    - Svetlana Smirnova, Percona
  - 9:45 a.m. | Hotel Nikko – Peninsula
    - Solving Performance Problems Using MySQL Enterprise Monitor [HOL5412]
    - Mark Matthews, Oracle
- Thursday (cont.)
  - 1:15 p.m. | Park Central – City
    - A MySQL Sys Schema Deep Dive [CON4589]
    - Mark Leith, Oracle
  - 2:30 p.m. | Park Central – Stanford
    - MySQL Security Best Practices [CON4680]
    - Georgi Kodinov and Mike Frank, Oracle

# Oracle MySQL Cloud Service

## Learn More @ OpenWorld & JavaOne



### Conference Presentation

- **CON4851**

#### *Introducing MySQL Cloud Service*

Park Central Hotel – 3rd Floor /  
Stanford Room

Sept 20<sup>th</sup> (Tues) @ 4PM

- **CON3810**

#### *Deep Dive: Oracle MySQL Cloud Service*

Park Central Hotel – 3<sup>rd</sup> Floor /  
Stanford

Sept 21<sup>st</sup> (Wed) @ 11AM

### Hands-on Lab

- **HOL4079**

#### *Hands on with Oracle MySQL Cloud Service*

Nikko Hotel - Peninsula

Sept 21<sup>st</sup> (Wed) @ 3:30PM

### Demo Pod

- *Database Showcase @ OpenWorld*

#### *Oracle MySQL Cloud Service*

Moscone Center South -  
Exhibition Hall

- *JavaHub @ JavaOne*

#### *IaaS & Databases: Compute VMs, Containers, MySQL, NoSQL, and Oracle Database in the Cloud*

Hilton Hotel – Grand Ballroom  
Exhibition Hall



# MySQL Community Reception @ Oracle OpenWorld

**Celebrate, Have Fun and Mingle with Oracle's MySQL Engineers & Your Peers**

- Tuesday, September 20 @ 7 pm
- Jillian's at Metreon: 175 Fourth Street, San Francisco  
At the corner of Howard and 4<sup>th</sup> st.; only 2-min walk from Moscone Center (same place as last year)

## Join us!

MySQL Community Reception  
at Oracle OpenWorld



Thank You

# Q&A

**ORACLE**<sup>®</sup>