

Performance Schema and Sys Schema

Mark Leith
Senior Software Development Manager
MySQL Enterprise Tools, Oracle



Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |

Program Agenda

- 1 ➤ This Introduction
- 2 ➤ Improvements to monitoring in MySQL 5.7
- 3 ➤ Improvements made so far in MySQL 8.0
- 4 ➤ Questions and feedback!

Performance Schema Improvements in MySQL 5.7

- 23 Worklogs completed
 - 14 Adding new instrumentation
 - 4 Performance / Scalability improvements
 - 5 Configuration / Refactoring improvements
- 35 new tables added (*5.7 has 87, 5.6 had 52, 5.5 had 17*)
- 419 new instruments (*5.7 has 972, 5.6 had 553, 5.5 had 222*)

Performance Schema Improvements in MySQL 8.0

- 4 Worklogs completed
 - 2 Adding new instrumentation
 - 1 Scalability Improvement
 - 1 Refactoring
- 7 new tables added (*8.0 has 96, 5.7 - 87, 5.6 - 52, 5.5 - 17*)
- 143 new instruments (*8.0 has 1152, 5.7 - 1009, 5.6 - 553, 5.5 - 222*)
- *So far ...*

Program Agenda

- 1 ➤ That Introduction
- 2 ➤ Improvements to monitoring in MySQL 5.7
- 3 ➤ Improvements made so far in MySQL 8.0
- 4 ➤ Questions and feedback!

"Understanding where MySQL can allocate memory can help us to find the cause in most cases. It is not as straightforward as it should be and I'm very hopeful future releases of MySQL, MariaDB or Drizzle bring improvements in this space allowing us to see directly for what purpose memory is allocated and so detect all kinds of memory usage problems easier."

Peter Zaitsev, CEO, Percona

WL#3249 - PERFORMANCE_SCHEMA, Instrument memory usage

<http://www.mysqlperformanceblog.com/2012/03/21/troubleshooting-mysql-memory-usage/>



WL#3249 - Instrument Memory Usage

- Added **212** different memory instrumentation types (all disabled by default)
- Records **current, high and low water marks** of allocations
- Does *not* track latency of memory allocation
- 5 new summary tables
 - `memory_summary_by_account_by_event_name`
 - `memory_summary_by_host_by_event_name`
 - `memory_summary_by_thread_by_event_name`
 - `memory_summary_by_user_by_event_name`
 - `memory_summary_global_by_event_name`

p_s.memory_summary_global_by_event_name

```
mysql> select * from performance_schema.memory_summary_global_by_event_name
-> order by current_number_of_bytes_used desc limit 1\G
***** 1. row *****
    EVENT_NAME: memory/innodb/buf_buf_pool
    COUNT_ALLOC: 1
    COUNT_FREE: 0
    SUM_NUMBER_OF_BYTES_ALLOC: 137428992
    SUM_NUMBER_OF_BYTES_FREE: 0
    LOW_COUNT_USED: 0
    CURRENT_COUNT_USED: 1
    HIGH_COUNT_USED: 1
    LOW_NUMBER_OF_BYTES_USED: 0
CURRENT_NUMBER_OF_BYTES_USED: 137428992
HIGH_NUMBER_OF_BYTES_USED: 137428992
```

InnoDB Buffer Pool

137.429 Megabytes

sys.memory_global_by_current_bytes view

```
mysql> select * from sys.memory_global_by_current_bytes\G
***** 1. row ****
  event_name: memory/innodb/buf_buf_pool
  current_count: 1
  current_alloc: 131.06 MiB
  current_avg_alloc: 131.06 MiB
    high_count: 1
    high_alloc: 131.06 MiB
    high_avg_alloc: 131.06 MiB
***** 2. row ****
  event_name: memory/innodb/log0log
  current_count: 9
  current_alloc: 16.01 MiB
  current_avg_alloc: 1.78 MiB
    high_count: 9
    high_alloc: 16.01 MiB
    high_avg_alloc: 1.78 MiB
  ...
  ...
```



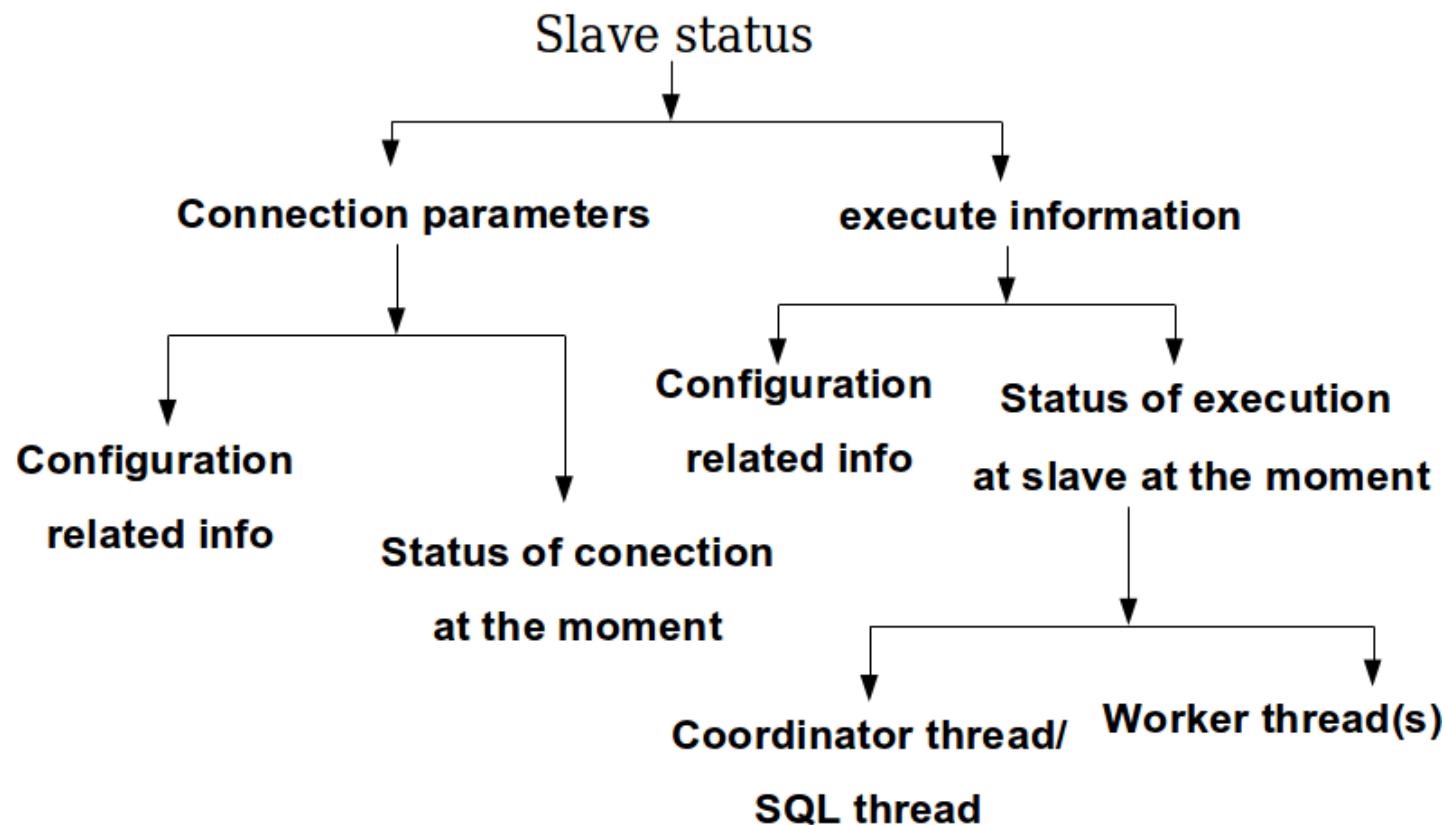
sys.user_summary view

```
mysql> select * from sys.user_summary\G
***** 1. row *****
      user: mark
statements: 3072
statement_latency: 1.77 s
statement_avg_latency: 575.29 us
table_scans: 7
file_ios: 20043
file_io_latency: 346.79 ms
current_connections: 1
total_connections: 1
unique_hosts: 1
current_memory: 515.81 KiB
total_memory_allocated: 30.69 MiB
```

WL#3656 - PERFORMANCE_SCHEMA table for SHOW SLAVE STATUS

- 6 new tables
 - `replication_connection_configuration`
 - `replication_connection_status`
 - `replication_execute_configuration`
 - `replication_execute_status`
 - `replication_execute_status_by_coordinator`
 - `replication_execute_status_by_worker`

SHOW SLAVE STATUS Split



Replication Connection Config

```
mysql> select * from replication_connection_configuration\G
***** 1. row *****
    CHANNEL_NAME:
        HOST: 127.0.0.1
        PORT: 13253
        USER: rsandbox
    NETWORK_INTERFACE:
        AUTO_POSITION: 1
        SSL_ALLOWED: NO
        SSL_CA_FILE:
        SSL_CA_PATH:
    SSL_CERTIFICATE:
        SSL_CIPHER:
        SSL_KEY:
    SSL_VERIFY_SERVER_CERTIFICATE: NO
        SSL_CRL_FILE:
        SSL_CRL_PATH:
    CONNECTION_RETRY_INTERVAL: 60
    CONNECTION_RETRY_COUNT: 86400
    HEARTBEAT_INTERVAL: 30.000
***** 2. row *****
    CHANNEL_NAME: 56-cluster
```

One row per replication channel
when using multi-master replication

Replication Connection Status

```
mysql> select * from performance_schema.replication_connection_status\G
***** 1. row *****
    CHANNEL_NAME:
    GROUP_NAME:
    SOURCE_UUID: 00013253-1111-1111-1111-111111111111
        THREAD_ID: 23
    SERVICE_STATE: ON
COUNT_RECEIVED_HEARTBEATS: 27
LAST_HEARTBEAT_TIMESTAMP: 2015-09-18 15:01:00
RECEIVED_TRANSACTION_SET: <GTID SET>
    LAST_ERROR_NUMBER: 0
    LAST_ERROR_MESSAGE:
LAST_ERROR_TIMESTAMP: 0000-00-00 00:00:00
***** 2. row *****
    CHANNEL_NAME: 56-cluster
    GROUP_NAME:
    SOURCE_UUID: 16d226ee-f89b-11e4-9d4f-b36d08df49cf
        THREAD_ID: 25
    SERVICE_STATE: ON
COUNT_RECEIVED_HEARTBEATS: 27
    ...

```



Replication Applier Status Config / State

```
mysql> select * from performance_schema.replication_applier_configuration;
```

CHANNEL_NAME	DESIRED_DELAY
	0
56-cluster	0

One row per replication channel
when using multi-master replication

```
mysql> select * from performance_schema.replication_applier_status;
```

CHANNEL_NAME	SERVICE_STATE	REMAINING_DELAY	COUNT_TRANSACTIONS_RETRIES
	ON	NULL	0
56-cluster	ON	NULL	0

Replication Applier Status Details

```
mysql> select * from performance_schema.replication_applier_status_by_coordinator;
```

CHANNEL_NAME	THREAD_ID	SERVICE_STATE	LAST_ERROR_NUMBER	LAST_ERROR_MESSAGE	LAST_ERROR_TIMESTAMP
	24	ON	0		0000-00-00 00:00:00
56-cluster	26	ON	0		0000-00-00 00:00:00

One row per channel use for all applier state when not using parallel workers

One row per parallel worker, across all replication channels

```
mysql> select * from performance_schema.replication_applier_status_by_worker;
```

CHANNEL_NAME	WORKER_ID	THREAD_ID	SERVICE_STATE	LAST_SEEN_TRANSACTION	LAST_ERROR_NUMBER	LAST_ERROR_MESSAGE	LAST_ERROR_TIMESTAMP
	1	27	ON	<GTID SET>	0		0000-00-00 00:00:00
	2	29	ON	<GTID SET>	0		0000-00-00 00:00:00
	3	31	ON	<GTID SET>	0		0000-00-00 00:00:00
	4	33	ON	<GTID SET>	0		0000-00-00 00:00:00
56-cluster	1	28	ON	<GTID SET>	0		0000-00-00 00:00:00
56-cluster	2	30	ON	<GTID SET>	0		0000-00-00 00:00:00
56-cluster	3	32	ON	<GTID SET>	0		0000-00-00 00:00:00
56-cluster	4	34	ON	<GTID SET>	0		0000-00-00 00:00:00



"Unfortunately, it's unlikely that I'll be able to create a reproducible test case, because there's no way to actually see what is happening. I hope that a future version of MySQL will include a more comprehensive set of tables for inspecting locks, requests, and waits at all layers of the server."

Baron Schwartz, CEO, VividCortex

WL#5879 - PERFORMANCE_SCHEMA, MDL lock instrumentation

WL#5879 - PERFORMANCE_SCHEMA, MDL lock instrumentation

- Added the `wait/lock/metadata/sql/mdl` instrument
- 2 new current instance tables
 - `metadata_locks`
 - `table_handles`

p_s.metadata_locks structure

Field	Type	Null	Key	Default	Extra
OBJECT_TYPE	varchar(64)	NO		NULL	
OBJECT_SCHEMA	varchar(64)	YES		NULL	
OBJECT_NAME	varchar(64)	YES		NULL	
OBJECT_INSTANCE_BEGIN	bigint(20) unsigned	NO		NULL	
LOCK_TYPE	varchar(32)	NO		NULL	
LOCK_DURATION	varchar(32)	NO		NULL	
LOCK_STATUS	varchar(32)	NO		NULL	
SOURCE	varchar(64)	YES		NULL	
OWNER_THREAD_ID	bigint(20) unsigned	YES		NULL	
OWNER_EVENT_ID	bigint(20) unsigned	YES		NULL	

p_s.metadata_locks sample output

```
mysql> select object_type as scope, object_schema, object_name, lock_type, lock_duration, lock_status
->   from metadata_locks
->  order by object_type = 'global' desc, object_type = 'schema' desc,
->            object_type = 'table' desc, object_type = 'commit' desc;
```

scope	object_schema	object_name	lock_type	lock_duration	lock_status
GLOBAL	NULL	NULL	SHARED	EXPLICIT	GRANTED
GLOBAL	NULL	NULL	INTENTION_EXCLUSIVE	STATEMENT	PENDING
GLOBAL	NULL	NULL	INTENTION_EXCLUSIVE	STATEMENT	PENDING
GLOBAL	NULL	NULL	INTENTION_EXCLUSIVE	STATEMENT	PENDING
GLOBAL	NULL	NULL	INTENTION_EXCLUSIVE	STATEMENT	PENDING
GLOBAL	NULL	NULL	INTENTION_EXCLUSIVE	STATEMENT	PENDING
GLOBAL	NULL	NULL	INTENTION_EXCLUSIVE	STATEMENT	PENDING
GLOBAL	NULL	NULL	INTENTION_EXCLUSIVE	STATEMENT	PENDING
TABLE	mem_events	action_logs	SHARED_READ	TRANSACTION	GRANTED
TABLE	mem_events	events	SHARED_READ	TRANSACTION	GRANTED
TABLE	performance_schema	metadata_locks	SHARED_READ	TRANSACTION	GRANTED
COMMIT	NULL	NULL	SHARED	EXPLICIT	GRANTED



sys.schema_table_lock_waits

```
mysql> select * from sys.schema_table_lock_waits\G
***** 1. row *****
    object_schema: test
    object_name: t
    waiting_thread_id: 43
    waiting_pid: 21
    waiting_account: msandbox@localhost
    waiting_lock_type: SHARED_UPGRADABLE
    waiting_lock_duration: TRANSACTION
    waiting_query: alter table test.t add foo int
    waiting_query_secs: 988
    waiting_query_rows_affected: 0
    waiting_query_rows_examined: 0
    blocking_thread_id: 42
    blocking_pid: 20
    blocking_account: msandbox@localhost
    blocking_lock_type: SHARED_NO_READ_WRITE
    blocking_lock_duration: TRANSACTION
    sql_kill_blocking_query: KILL QUERY 20
    sql_kill_blocking_connection: KILL 20
```

sys.innodb_lock_waits

```
mysql> SELECT * FROM sys.innodb_lock_waits\G
***** 1. row ****
    wait_started: 2014-11-11 13:39:20
        wait_age: 00:00:07
    wait_age_secs: 7
        locked_table: `db1`.`t1`
        locked_index: PRIMARY
        locked_type: RECORD
        waiting_trx_id: 867158
    waiting_trx_started: 2014-11-11 13:39:15
        waiting trx age: 00:00:12
        waiting_trx_rows_locked: 0
        waiting_trx_rows_modified: 0
            waiting_pid: 3
                waiting_query: UPDATE t1 SET val = val + 1 WHERE id = 2
                waiting_lock_id: 867158:2363:3:3
                waiting_lock_mode: X
                blocking_trx_id: 867157
                    blocking_pid: 4
...

```



"Suppose I accessed your database and listed the entire set of stored functions and procedures. How many of them are you not even sure are in use anymore? How many of them you think you can DROP, but are too afraid to, and keep them in *just in case*?"

Shlomi Noach

WL#5766 - PERFORMANCE_SCHEMA, Stored programs instrumentation

<http://code.openark.org/blog/mysql/why-delegating-code-to-mysql-stored-routines-is-poor-engineering-practice>



WL#5766 - Stored programs instrumentation

- Tracks **Stored Procedures, Stored Functions, Triggers and Events**
- Added 16 new **statement/sp/%** instruments
 - Expose the different work flows that stored programs use, such as cursor operations, workflow controls etc.
 - Integrated in to the normal statement instrumentation
- 1 new summary table
 - **events_statements_summary_by_program**

p_s.memory_summary_global_by_event_name

```
mysql> select * from events_statements_summary_by_program\G
***** 1. row *****
    OBJECT_TYPE: PROCEDURE
    OBJECT_SCHEMA: ps_demo
    OBJECT_NAME: ps_demo_proc
    COUNT_STAR: 1
    SUM_TIMER_WAIT: 6970931000
    MIN_TIMER_WAIT: 6970931000
    AVG_TIMER_WAIT: 6970931000
    MAX_TIMER_WAIT: 6970931000
    COUNT_STATEMENTS: 5
    SUM_STATEMENTS_WAIT: 6802181000
    MIN_STATEMENTS_WAIT: 16372000
    AVG_STATEMENTS_WAIT: 1360436000
    MAX_STATEMENTS_WAIT: 6484366000
    SUM_LOCK_TIME: 176401000000
    SUM_ERRORS: 0
    SUM_WARNINGS: 0
                SUM_ROWS_AFFECTED: 1
                SUM_ROWS_SENT: 0
                SUM_ROWS_EXAMINED: 1
    SUM_CREATED_TMP_DISK_TABLES: 0
    SUM_CREATED_TMP_TABLES: 0
    SUM_SELECT_FULL_JOIN: 0
    SUM_SELECT_FULL_RANGE_JOIN: 0
    SUM_SELECT_RANGE: 0
    SUM_SELECT_RANGE_CHECK: 0
    SUM_SELECT_SCAN: 0
    SUM_SORT_MERGE_PASSES: 0
    SUM_SORT_RANGE: 0
    SUM_SORT_ROWS: 0
    SUM_SORT_SCAN: 0
    SUM_NO_INDEX_USED: 0
    SUM_NO_GOOD_INDEX_USED: 0
```

WL#5864 - PERFORMANCE SCHEMA, instrument TRANSACTIONS

- Added the transaction instrument
- 8 new raw and summary tables
 - `events_transactions_current`
 - `events_transactions_history`
 - `events_transactions_history_long`
 - `events_transactions_summary_by_account_by_event_name`
 - `events_transactions_summary_by_host_by_event_name`
 - `events_transactions_summary_by_thread_by_event_name`
 - `events_transactions_summary_by_user_by_event_name`
 - `events_transactions_summary_global_by_event_name`

WL#5864 - PERFORMANCE SCHEMA, instrument TRANSACTIONS

- Exposes details of transactions in raw or summary views
- Show details such as
 - Transaction latency
 - Isolation levels
 - Auto commit
 - Savepoint info
 - GTID or transaction IDs
- Link transactions to the statements within the transactions

p_s.events_transactions_current

```
mysql> select * from events_transactions_current\G
***** 1. row ****
      THREAD_ID: 1
      EVENT_ID: 23733
    END_EVENT_ID: 23742
      EVENT_NAME: transaction
        STATE: COMMITTED
      TRX_ID: 281479898269256
        GTID: NULL
        XID: NULL
      XA_STATE: NULL
      SOURCE: handler.cc:1246          NUMBER_OF_SAVEPOINTS: 0
    TIMER_START: 31140612726000      NUMBER_OF_ROLLBACK_TO_SAVEPOINT: 0
    TIMER_END: 31140647445000      NUMBER_OF_RELEASE_SAVEPOINT: 0
    TIMER_WAIT: 34719000          OBJECT_INSTANCE_BEGIN: NULL
    ACCESS_MODE: READ WRITE
ISOLATION_LEVEL: REPEATABLE READ
    AUTOCOMMIT: YES
                NESTING_EVENT_ID: NULL
                NESTING_EVENT_TYPE: NULL
```



sys.ps_thread_trx_info(<thread_id>)

```
SELECT sys.ps_thread_trx_info(48) as trx_info\G
***** 1. row *****
trx_info: [
  {
    "time": "790.70 us",
    "state": "COMMITTED",
    "mode": "READ WRITE",
    "autocommitted": "NO",
    "gtid": "AUTOMATIC",
    "isolation": "REPEATABLE READ",
    "statements_executed": [
      {
        "sql_text": "INSERT INTO info VALUES (1, 'foo')",
        "time": "471.02 us",
        "schema": "trx",
        "rows_examined": 0,
        "rows_affected": 1,
        "rows_sent": 0,
        "tmp_tables": 0,
        "tmp_disk_tables": 0,
        "sort_rows": 0,
        "sort_merge_passes": 0
      },
      ...
    ]
  }
]
```



sys.ps_trace_thread procedure

- **ps_trace_thread()** monitors a specific thread for a period
- Captures as much information on the thread activity as possible
- Returns a “dot” formatted file, that can graph the event hierarchy
 - [http://en.wikipedia.org/wiki/DOT_\(graph_description_language\)](http://en.wikipedia.org/wiki/DOT_(graph_description_language))

sys.ps_trace_thread procedure example

```
mysql> call ps_trace_thread(27768, '/tmp/stack_27768.dot', 60, 0.1, true, true, true);
+-----+
| Info |
+-----+
| Data collection starting for THREAD_ID = 27768 |
+-----+
1 row in set (4.82 sec)

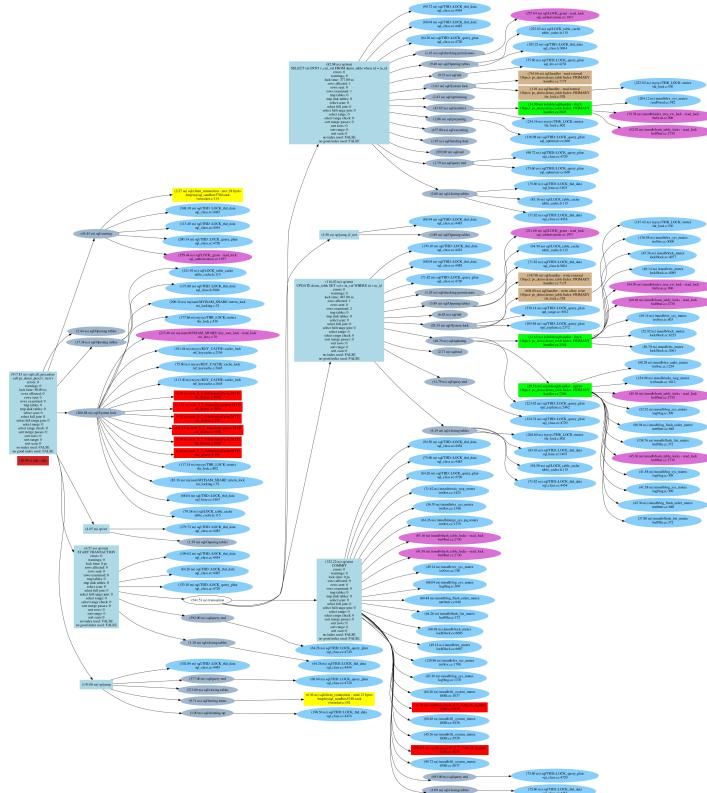
+-----+
| Info |
+-----+
| Stack trace written to /tmp/stack_27768.dot |
+-----+
1 row in set (60.90 sec)

+-----+
| Convert to PDF |
+-----+
| dot -Tpdf -o /tmp/stack_27768.pdf /tmp/stack_27768.dot |
+-----+
1 row in set (60.90 sec)

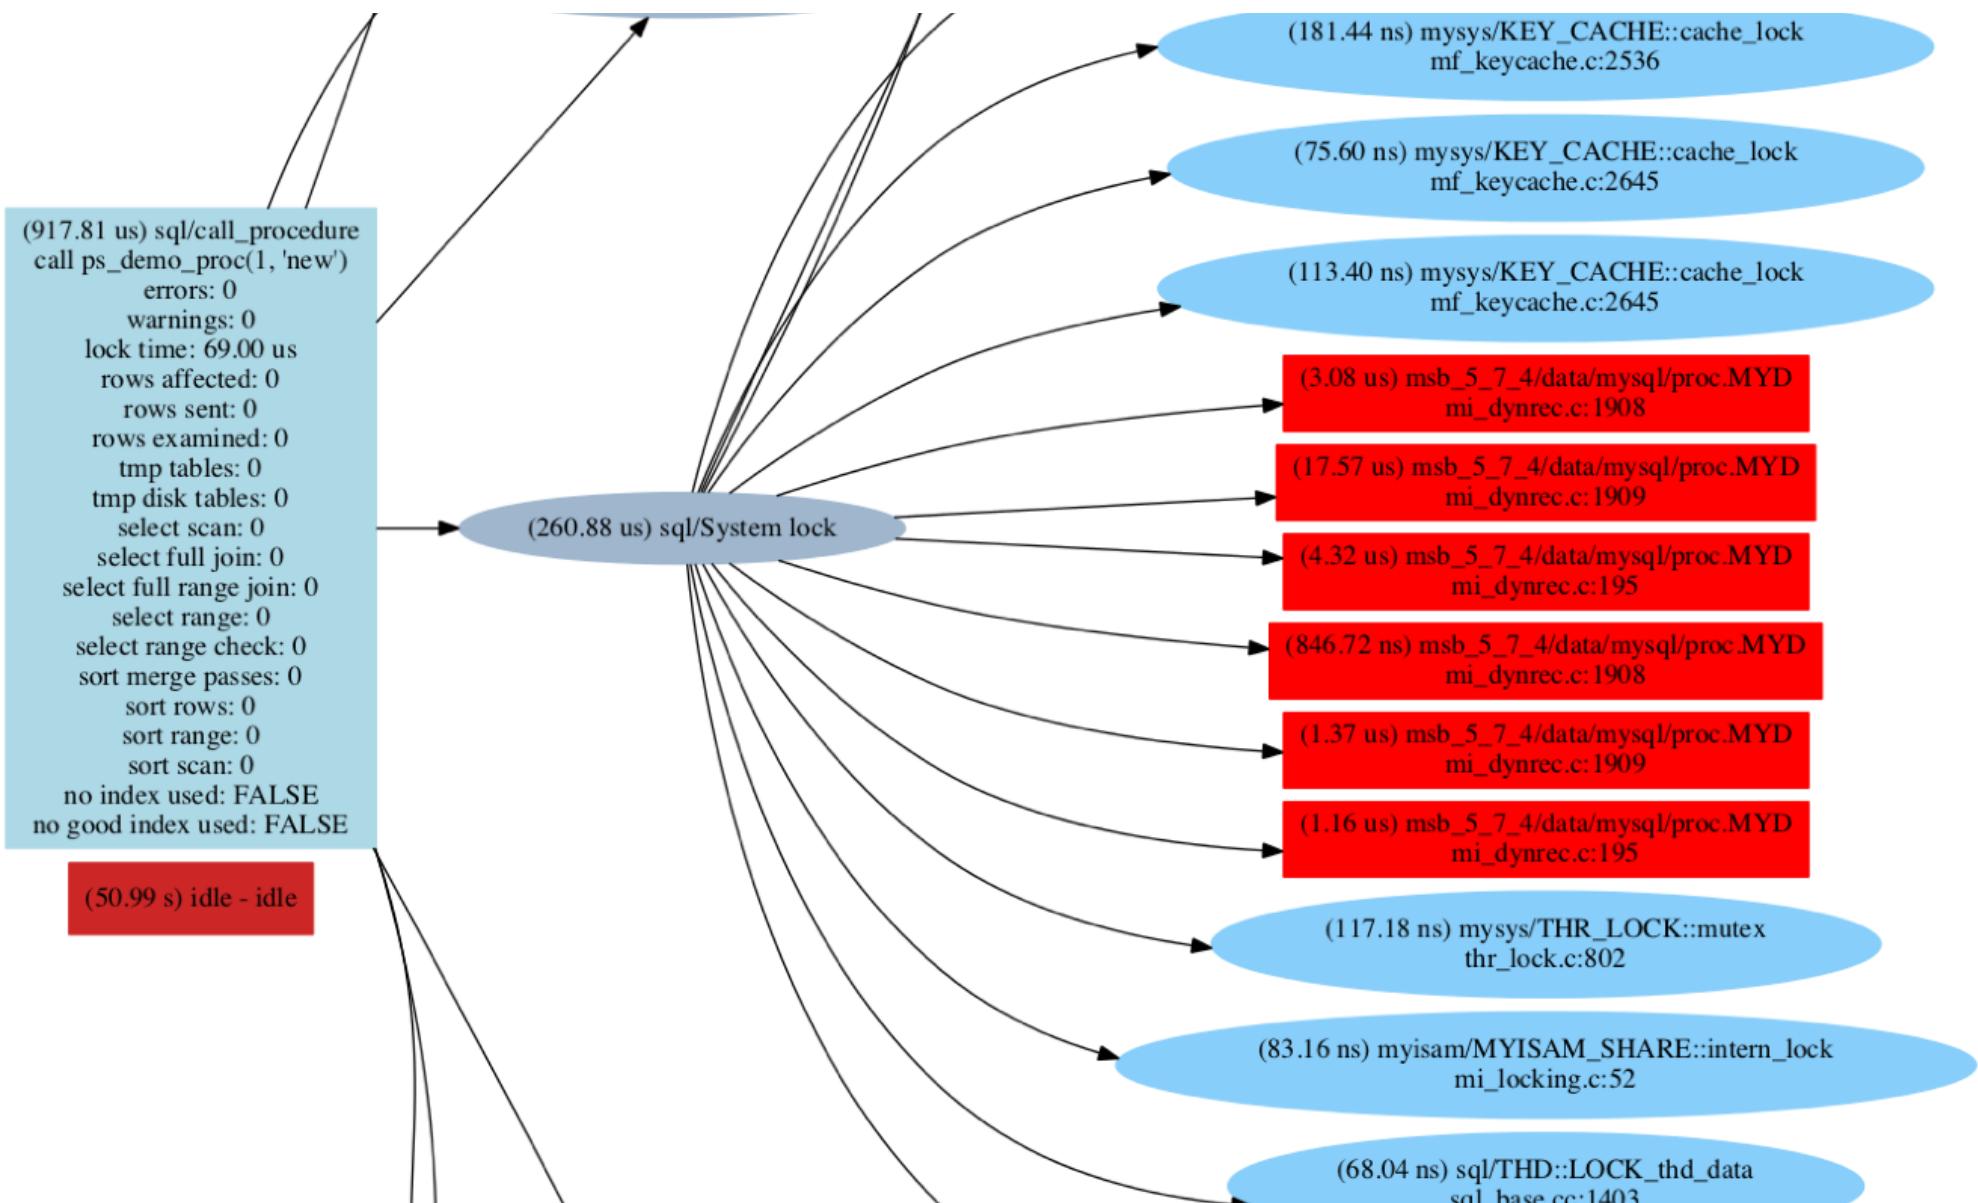
+-----+
| Convert to PNG |
+-----+
| dot -Tpng -o /tmp/stack_27768.png /tmp/stack_27768.dot |
+-----+
1 row in set (60.90 sec)+
```

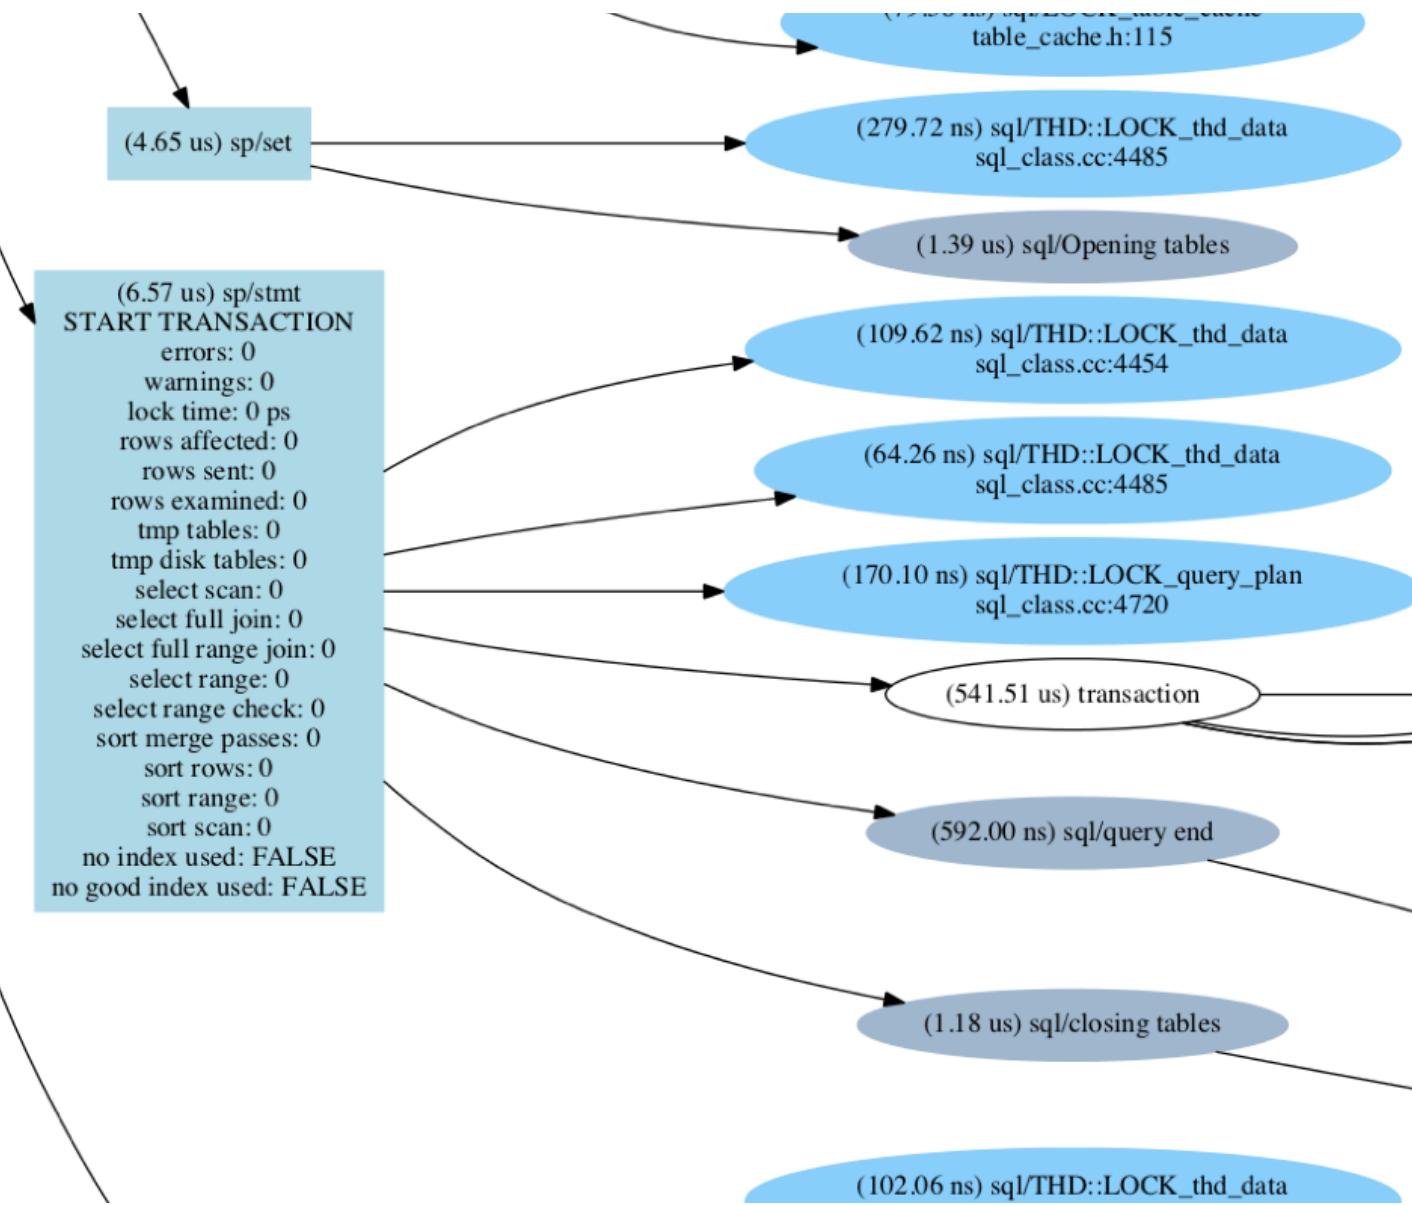


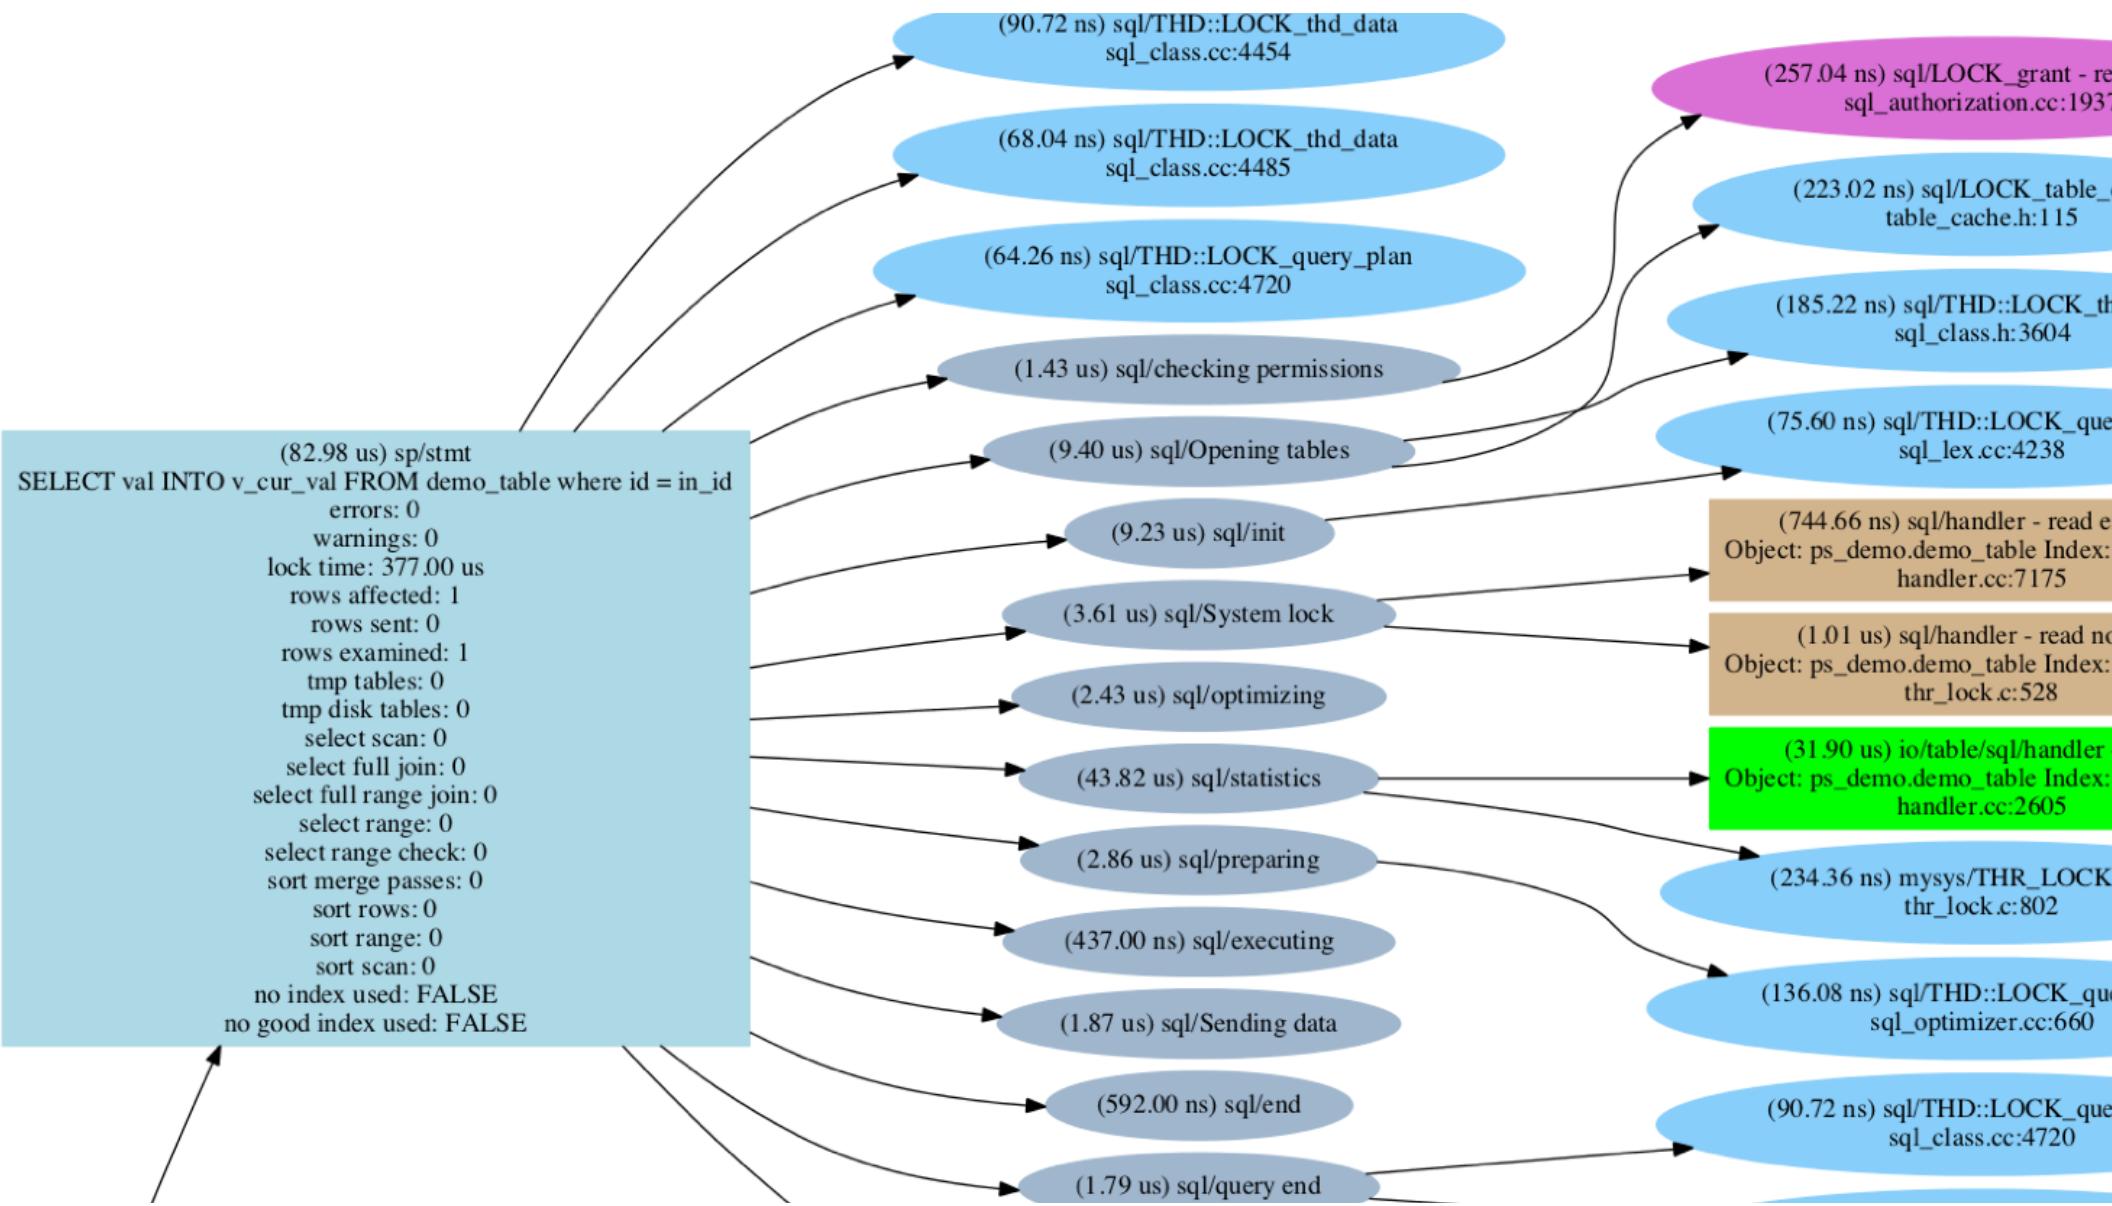
ps_trace_thread procedure example

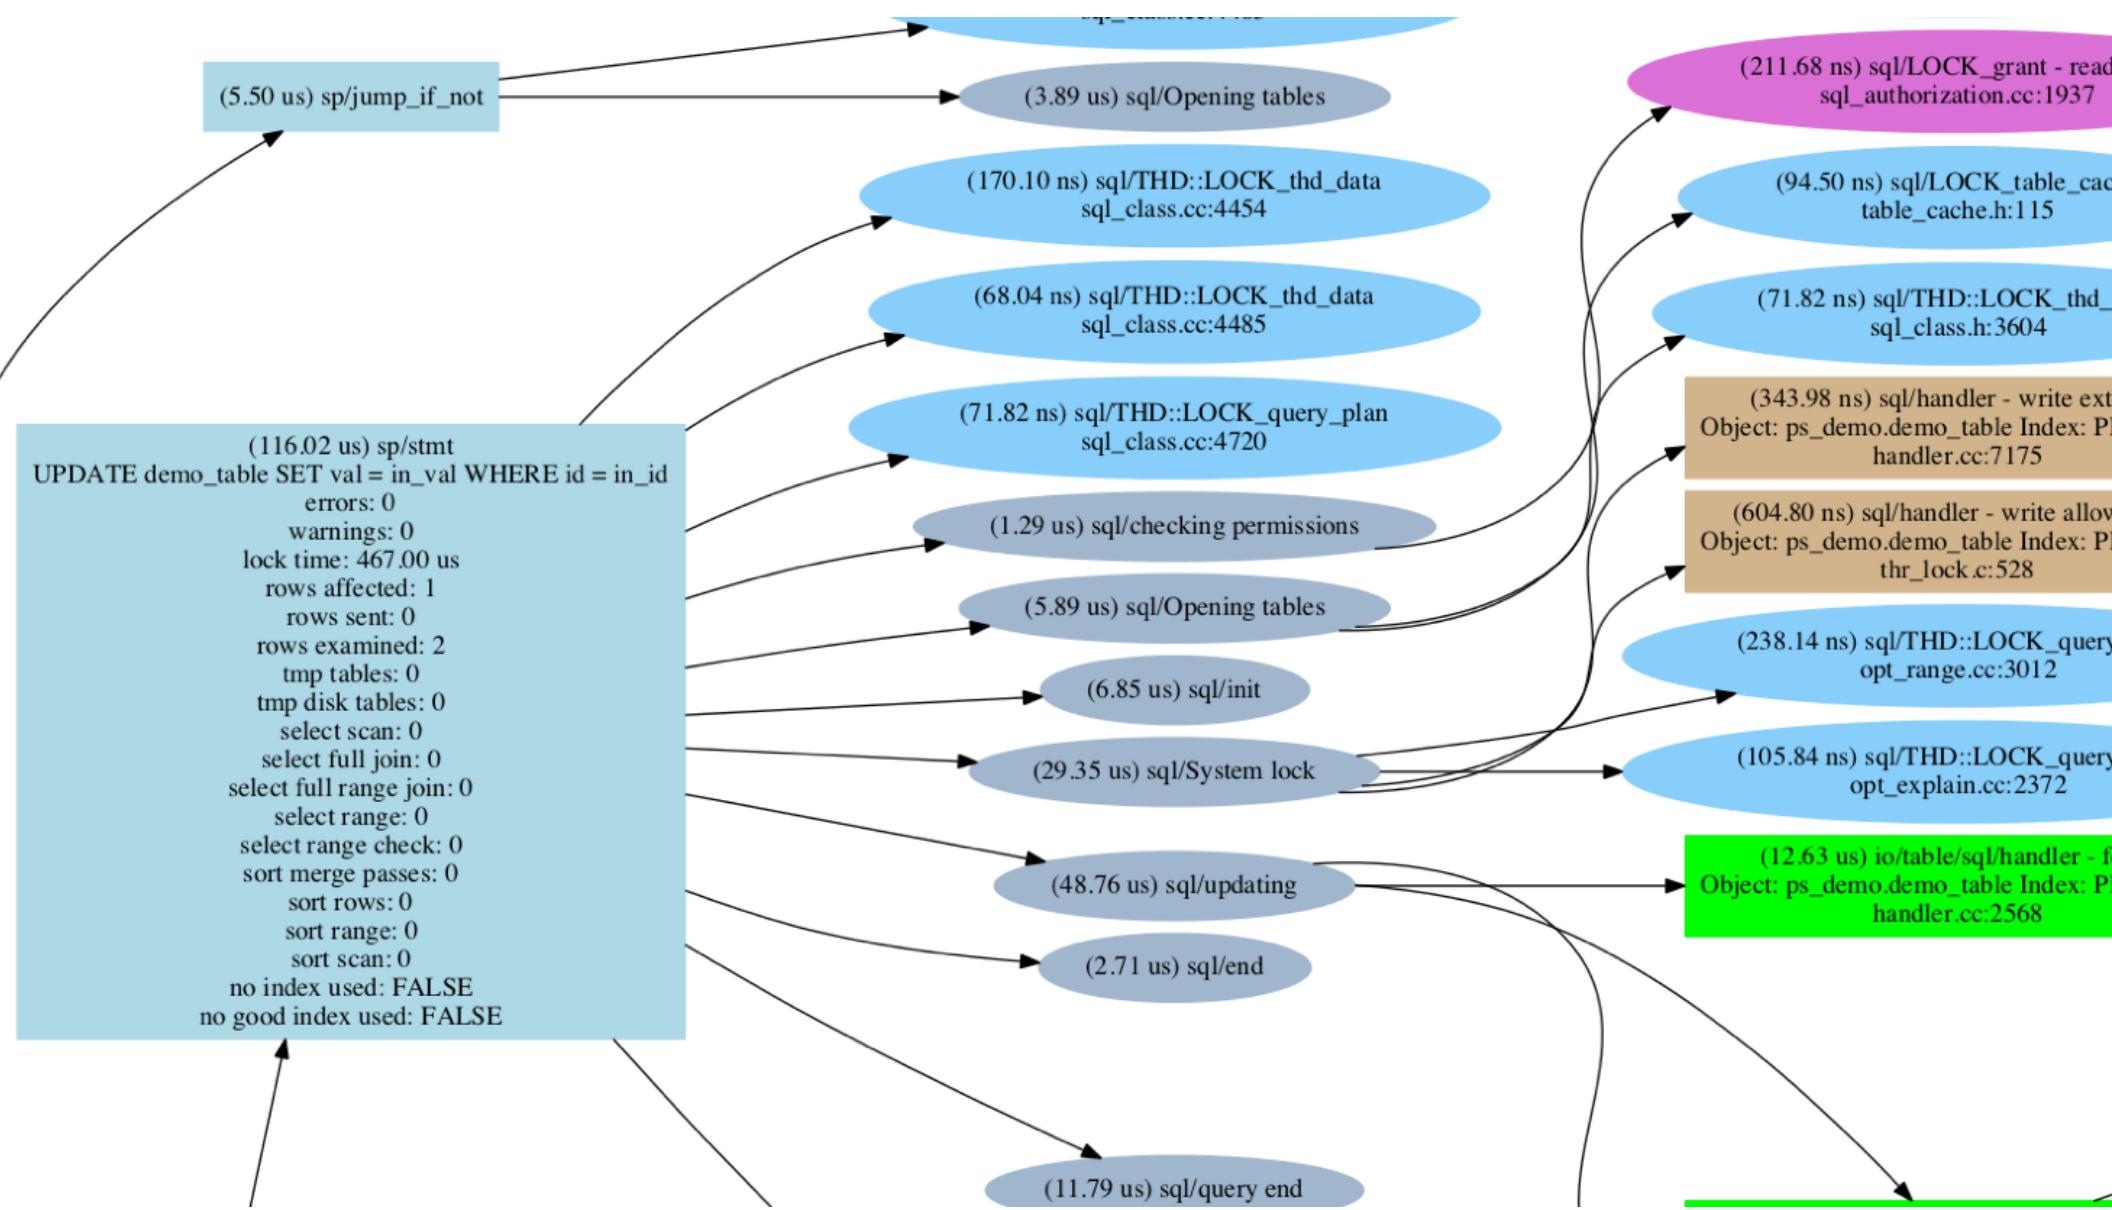


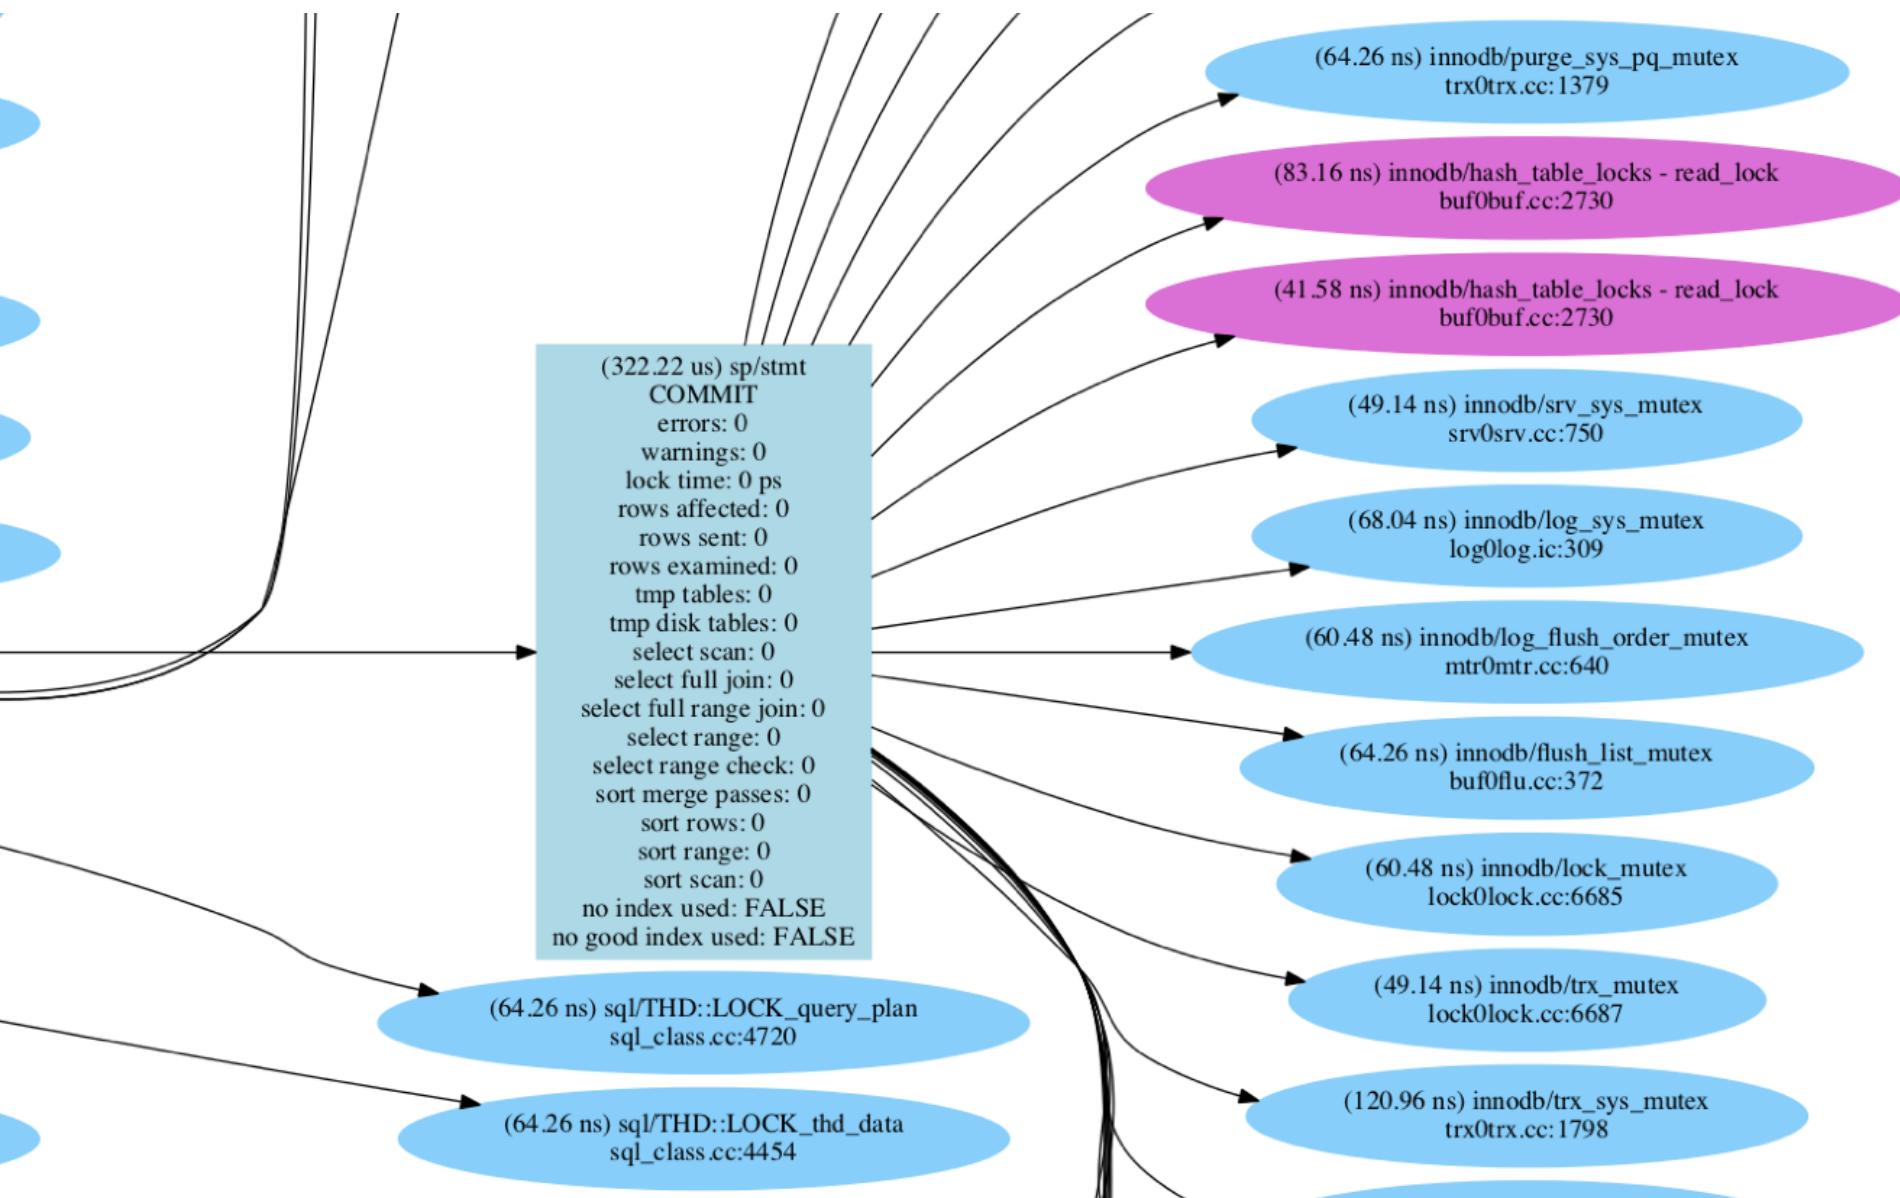
MySQL 5.7
**Procedures
Transactions
Statements
Stages
Waits**

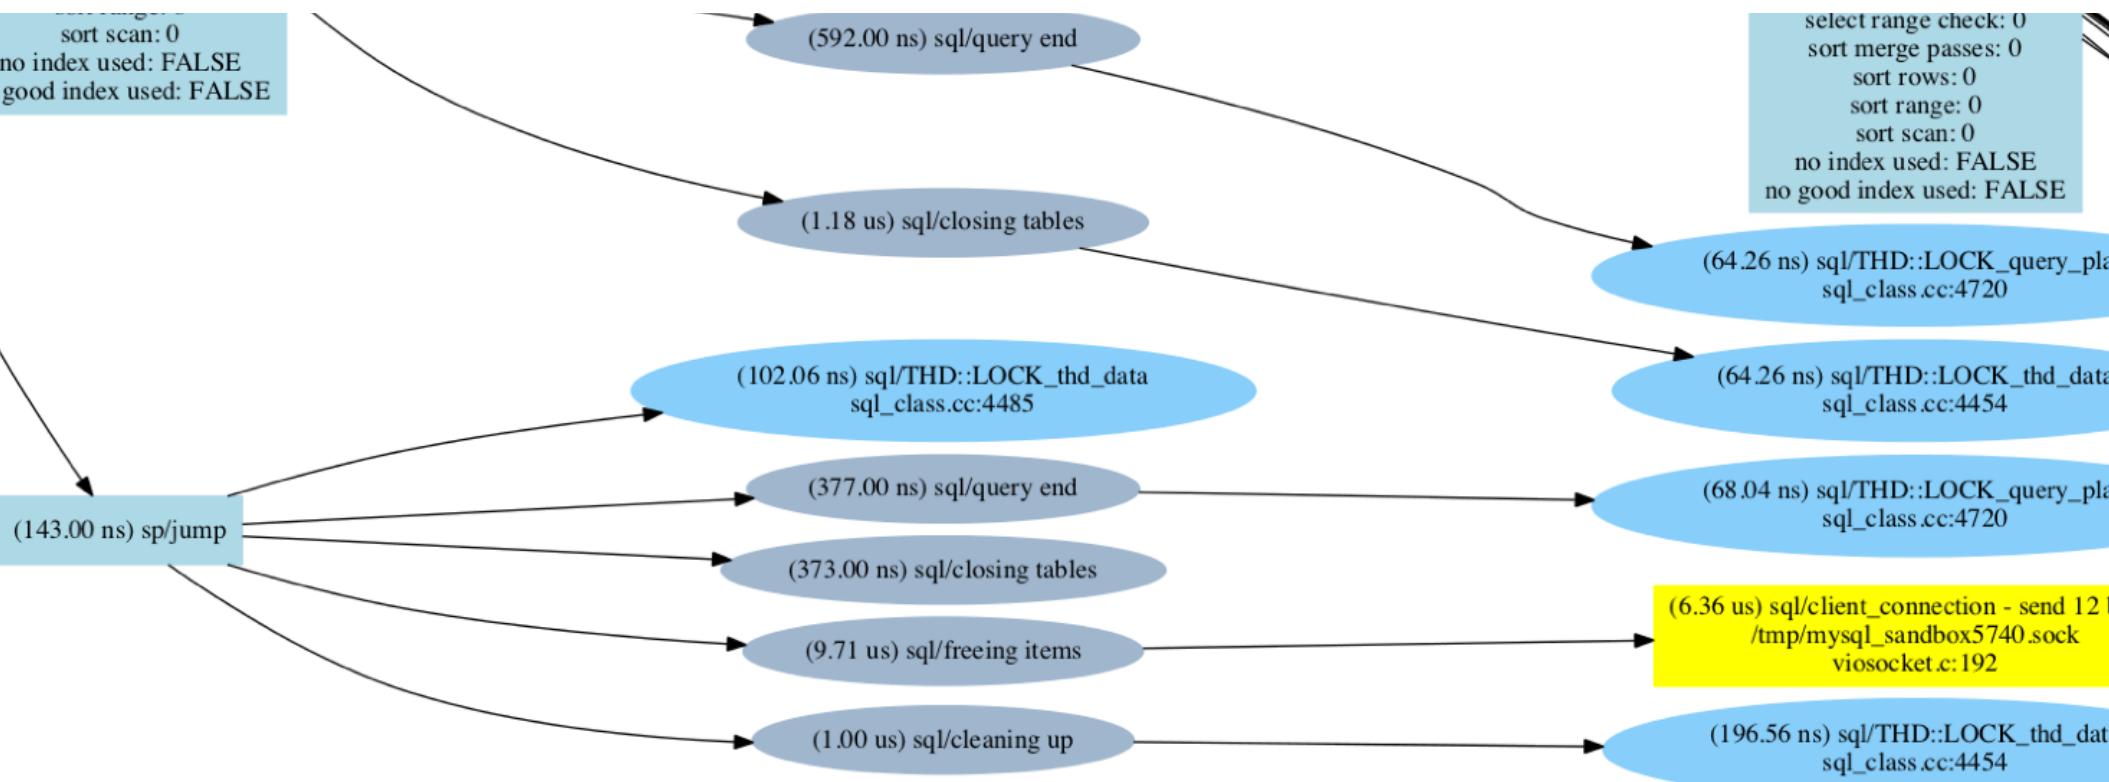












p_s.prepared_statement_instances

```
mysql> select * from performance_schema.prepared_statements_instances\G
***** 1. row *****
          OBJECT_INSTANCE_BEGIN: 140198306602144
              STATEMENT_ID: 1
      STATEMENT_NAME: stmt1
          SQL_TEXT: select * from test.t1
        OWNER_THREAD_ID: 54003
        OWNER_EVENT_ID: 935
    OWNER_OBJECT_TYPE: NULL
OWNER_OBJECT_SCHEMA: NULL
OWNER_OBJECT_NAME: NULL
      TIMER_PREPARE: 15386274000
      COUNT_REPREPARE: 0
      COUNT_EXECUTE: 1
      SUM_TIMER_EXECUTE: 217699000
      MIN_TIMER_EXECUTE: 217699000
      AVG_TIMER_EXECUTE: 217699000
      MAX_TIMER_EXECUTE: 217699000
          SUM_LOCK_TIME: 0
          SUM_ERRORS: 0
          SUM_WARNINGS: 0
          SUM_ROWS_AFFECTED: 0
          SUM_ROWS_SENT: 0
          SUM_ROWS_EXAMINED: 0
          SUM_CREATED_TMP_DISK_TABLES: 0
          SUM_CREATED_TMP_TABLES: 0
          SUM_SELECT_FULL_JOIN: 0
SUM_SELECT_FULL_RANGE_JOIN: 0
          SUM_SELECT_RANGE: 0
          SUM_SELECT_RANGE_CHECK: 0
          SUM_SELECT_SCAN: 0
          SUM_SORT_MERGE_PASSES: 0
          SUM_SORT_RANGE: 0
          SUM_SORT_ROWS: 0
          SUM_SORT_SCAN: 0
          SUM_NO_INDEX_USED: 0
          SUM_NO_GOOD_INDEX_USED: 0
```

WL#6884 - PERFORMANCE SCHEMA, USER VARIABLES

```
mysql> set @debug := 1;  
Query OK, 0 rows affected (0.10 sec)
```

```
mysql> set @totalRows := 300;  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> select * from performance_schema.user_variables_by_thread;  
+-----+-----+-----+  
| THREAD_ID | VARIABLE_NAME | VARIABLE_VALUE |  
+-----+-----+-----+  
| 33 | totalRows | 300 |  
| 33 | debug | 1 |  
| 34 | totalRows | 150 |  
+-----+-----+-----+
```

WL#7415 - PERFORMANCE SCHEMA, STAGE PROGRESS

- Added the **WORK_COMPLETED** and **WORK_ESTIMATED** columns to the **events_stages_current** table
- Work is in units, and depends on the stage what it may be, rows copied, bytes processed, etc.
- Started with adding tracking to:
 - **stage/sql/copy to tmp table**
 - *(more to come later in this presentation)*

WL#5889 - Add InnoDB events to Performance Schema's Event Stage table

- Follow up to WL#7415 - PERFORMANCE SCHEMA, STAGE PROGRESS
- Added stage progress instruments within InnoDB:
 - stage/innodb/buffer pool load
 - stage/innodb/alter table (read PK and internal sort)
 - stage/innodb/alter table (merge sort)
 - stage/innodb/alter table (insert)
 - stage/innodb/alter table (flush)
 - stage/innodb/alter table (log apply index)
 - stage/innodb/alter table (log apply table)
 - stage/innodb/alter table (end)

WL#5889 - Add InnoDB events to Performance Schema's Event Stage table

- InnoDB ALTER TABLE instruments work in tandem for the entire ALTER TABLE for example:

	work_complete	work_estimated
'..../alter table (read..)'	0 .. 100	700
'..../alter table (merge..)'	101 .. 300	700
'..../alter table (insert)'	301 .. 500	700
'..../alter table (flush)'	500 .. 650	700
'..../alter table (end)'	651 .. 700	700

sys.sessions progress estimation

```
mysql> select * from sys.sessions\G
***** 1. row *****
    thd_id: 44524
    conn_id: 44502
        user: msandbox@localhost
        db: test
    command: Query
        state: alter table (flush)
        time: 18
current_statement: alter table t1 add column g int
statement_latency: 18.45 s
        progress: 98.84
    lock_latency: 265.43 ms
    rows_examined: 0
        rows_sent: 0
    rows_affected: 0
    tmp_tables: 0
tmp_disk_tables: 0
    ...

```



WL#7800 - PERFORMANCE SCHEMA, SETUP_ACTORS ENABLED COLUMN

- The `performance_schema.setup_actors` table previously only worked for *including* users to instrument, but it was impossible to instrument all but one or two users easily
- Added `ENABLED` enum('YES','NO') to `setup_actors`
- To disable user “mark@localhost”:
 - `INSERT INTO performance_schema.setup_actors VALUES ('localhost', 'mark', '%', 'NO');`

WL#6629 - PERFORMANCE_SCHEMA, STATUS VARIABLES

- Added 9 new tables:
 - System Variables
 - **global_variables**
 - **session_variables**
 - **variables_by_thread**
 - Status Variables:
 - **global_status**
 - **session_status**
 - **status_by_thread**
 - **status_by_account**
 - **status_by_user**
 - **status_by_host**

WL#6629 - PERFORMANCE_SCHEMA, STATUS VARIABLES

```
mysql> select * from performance_schema.status_by_thread where variable_name like 'bytes%';
+-----+-----+-----+
| THREAD_ID | VARIABLE_NAME | VARIABLE_VALUE |
+-----+-----+-----+
|      37 | Bytes_received |      3901 |
|      37 | Bytes_sent     | 135557   |
+-----+-----+-----+
2 rows in set (0.01 sec)
```

```
mysql> select * from performance_schema.status_by_user where variable_name like 'bytes%';
+-----+-----+-----+
| USER      | VARIABLE_NAME | VARIABLE_VALUE |
+-----+-----+-----+
| NULL      | Bytes_received |      191  |
| NULL      | Bytes_sent     |      160  |
| msandbox  | Bytes_received |    4281  |
| msandbox  | Bytes_sent     | 136263  |
| root      | Bytes_received |      360  |
| root      | Bytes_sent     |      334  |
| mark      | Bytes_received |    3434  |
| mark      | Bytes_sent     | 101628  |
+-----+-----+-----+
```

Or by host, account

sys.metrics

```
mysql> select * from sys.metrics limit 199, 10;
```

Variable_name	Variable_value	Type	Enabled
tc_log_page_waits	0	Global Status	YES
threads_cached	7	Global Status	YES
threads_connected	3	Global Status	YES
threads_created	10	Global Status	YES
threads_running	3	Global Status	YES
uptime	275689	Global Status	YES
uptime_since_flush_status	275689	Global Status	YES
adaptive_hash_pages_added	0	InnoDB Metrics - adaptive_hash_index	NO
adaptive_hash_pages_removed	0	InnoDB Metrics - adaptive_hash_index	NO
adaptive_hash_rows_added	0	InnoDB Metrics - adaptive_hash_index	NO

performance_schema.global_status

information_schema.innodb_metrics

Program Agenda

- 1 ➤ That Introduction**
- 2 ➤ Improvements to monitoring in MySQL 5.7**
- 3 ➤ Improvements made so far in MySQL 8.0**
- 4 ➤ Questions and feedback!**

WL#6616 - PERFORMANCE SCHEMA, INDEXES

- Whilst Performance Schema tries to be low overhead from a monitoring perspective, the load of complex queries on the actual data could have significant overhead
- Most of this was due to the server having to process all event data when searching, or joining, or aggregating etc.
- Very noticeable in for instance the sys.session view with many connected sessions

WL#6616 - PERFORMANCE SCHEMA, INDEXES

- 125 “indexes” across the current 95 tables
- They are faked, we don’t maintain indexes for all instrumentation (would be too much overhead) - we fake them to the optimizer and then do lookups on the in memory data at the storage engine level
- Took the sys.session runtime, with 1,000 connected threads, from ~30 seconds to ~1 second (it’s one of the more complex joining views in sys)
- Added to most columns we *feel* will be searched upon (need feedback)

WL#8058 - PERFORMANCE SCHEMA, INSTRUMENT SERVER ERRORS

- Understanding the types of errors happening has been hard, you can get raw error data in the statement history tables in `performance_schema`, but nowhere that summarises what errors have happened over time
- Added a new **error** instrument
- 5 new summary tables (`events_errors_summary_[global | by_user | by_host | by_account | by_thread]`)
- We summarise by errors raised vs handled (i.e. within a routine HANDLER statement)

WL#8058 - PERFORMANCE SCHEMA, INSTRUMENT SERVER ERRORS

```
mysql> select * from events_errors_summary_global_by_error
->   where sum_error_raised > 0
->   order by sum_error_raised desc\G
*****
               1. row ****
ERROR_NUMBER: 1287
ERROR_NAME: ER_WARN_DEPRECATED_SYNTAX
SQL_STATE: HY000
SUM_ERROR_RAISED: 190
SUM_ERROR_HANDLED: 0
FIRST_SEEN: 2016-10-04 08:19:56
LAST_SEEN: 2016-10-04 09:10:19
*****
               2. row ****
ERROR_NUMBER: 3554
ERROR_NAME: ER_NO_SYSTEM_TABLE_ACCESS
SQL_STATE: HY000
SUM_ERROR_RAISED: 58
SUM_ERROR_HANDLED: 0
FIRST_SEEN: 2016-10-03 14:11:22
LAST_SEEN: 2016-10-04 09:10:21
```

WL#8688 - Support ability to persist SET GLOBAL settings

- Added the **variables_info** table
- Exposes metadata about system variables, such as the min/max values and where the variable was set from
- May be expanded upon (such as who set what when)

WL#8688 - Support ability to persist SET GLOBAL settings

```
mysql> select * from variables_info join global_variables using (variable_name) order by variable_source desc limit 10;
+-----+-----+-----+-----+-----+-----+
| VARIABLE_NAME | VARIABLE_SOURCE | VARIABLE_PATH | MIN_VALUE | MAX_VALUE | VARIABLE_VALUE |
+-----+-----+-----+-----+-----+-----+
| foreign_key_checks | DYNAMIC | | 0 | 0 | ON
| host_cache_size | DYNAMIC | | 0 | 65536 | 100
| show_compatibility_56 | DYNAMIC | | 0 | 0 | OFF
| basedir | COMMAND_LINE | | 0 | 0 | /Users/mark/git/bld-wl883/
| datadir | COMMAND_LINE | | 0 | 0 | /Users/mark/git/bld-wl883/data
| port | COMMAND_LINE | | 0 | 65535 | 8001
| socket | COMMAND_LINE | | 0 | 0 | /tmp/mysql-8001.sock
| auto_increment_increment | COMPILED | | 1 | 65535 | 1
| key_cache_division_limit | COMPILED | | 1 | 100 | 100
| tmp_table_size | COMPILED | | 1024 | 18446744073709551615 | 16777216
+-----+-----+-----+-----+-----+
10 rows in set (0.01 sec)
```



WL#8688 - Support ability to persist SET GLOBAL settings

- COMPILED
 - The default compiled in value
- GLOBAL
 - Set as part of global options file
- SERVER
 - Set as part of \$MYSQL_HOME/my.cnf options file
- EXPLICIT
 - Set within --defaults-file options file
- EXTRA
 - Set within --defaults-extra-file options file
- USER
 - Set within ~/.my.cnf options file
- LOGIN
 - Set within login path options file
- COMMAND LINE
 - Set as a command line parameter
- PERSISTED
 - Set within the new mysqld-auto.cnf persistent options file
- DYNAMIC
 - Set from a session dynamically

Question and feedback time!



Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |

Safe Harbor Statement

The preceding is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

ORACLE®